

Gdańsk, Poland, 26 July 2026

Dear Editors,

Please find enclosed the revised manuscript, “A Comparability Protocol for Benchmarking Relational Database Access Layers in Express.js,” submitted as a Research Paper to Information and Software Technology.

This revision addresses the scientific-status concerns raised in the July 23 review:

- It separates independent expected-result conformance from cross-implementation differential equivalence. A deterministic seed-specification oracle now checks reads without using native-driver output, while mutations are admitted only after database-state validation.
- The new state preflight exposed shared MySQL contamination that mutual equality had missed. A subsequent 662-row cross-engine seed-parity gate found and removed differing fan-out comment-author assignments. A result-blind author audit also removed Knex warning I/O and an unused MikroORM aggregation accessor. Every affected pilot was preserved and rejected; the full admission chain and measurement restarted from repetition 1.
- The treatment is now a mechanically named “policy-selected documented configuration,” not a model of practitioner behavior. The protocol requires a preregistered reconstructable policy but does not privilege documentation-first selection.
- RQ2 is limited to backend-stack transfer in the tested campaigns. DBMS, adapter, lower-level driver, protocol, and defaults may change together; close rank reversals are exploratory and no cross-host transfer is claimed.
- The protocol claim is narrowed to a concrete access-layer-specific discipline integrating established controls. CYNTHIA is recognized as direct prior cross-ORM differential-testing work, and a source-located audit applies the protocol codebook to nine external benchmarks.
- Reproduction labels now distinguish current author-run isolated reconstruction, historical author-run archive-isolated reconstruction, current measurement re-execution, and independent third-party reproduction. The current candidate check verified 60/60 JSON hashes and regenerated 35 table outputs plus four derived JSON files byte-for-byte without starting either DBMS; third-party reproduction has not occurred and is not claimed.

The manuscript also records two residual limitations without disguising them: all adapters and treatment assignments remain single-author, and both measurement campaigns use the same virtualized host. The protocol now includes independent implementation-review provenance as a recommended stage; result-blind review packets are supplied, but no unperformed human audit is counted as evidence.

The artifact includes pinned dependencies, immutable documentation snapshots and hashes, exact 42-file source archives for both accepted campaigns, the seed oracle, campaign-state validator, raw per-repetition measurements, a machine-readable protocol checklist, the external-audit dataset, and generators mapping every reported table and figure to its inputs. The complete revision artifact is archived as v1.12.15 under immutable Zenodo version DOI 10.5281/zenodo.21599104; concept DOI 10.5281/zenodo.21313858 resolves to the release series.

The article uses the required structured abstract. The separately documented conservative submission equivalent is 13,859 words, including all required declarations, leaving 1,141 words below the 15,000-word limit.

The work is original, has not been published previously, and is not under consideration elsewhere. I declare no competing interests and received no funding. The manuscript includes declarations covering generative-AI assistance in writing and research software; the author accepts responsibility for all content and results.

Sincerely,

Mateusz Miotk Faculty of Mathematics, Physics and Informatics, University of Gdańsk, Poland
mateusz.miotk@ug.edu.pl