

Highlights

- Nine idiomatic layers plus tuned native baselines on PostgreSQL and MySQL
- Deep-fetch spread of 6.5x narrows to 1.5x under a same-SQL control
- Prisma's native-throughput parity holds at 5x CPU but reverses under scaling
- Saturated tail-latency gaps largely dissolve at matched utilization
- Open harness verifies byte-identical responses across layers before measuring