

Highlights

- A reusable comparability protocol turns access-layer benchmarks into experiments
- Semantic-equivalence gate: byte-identical responses required before any timing
- Common raw-SQL contrast reduces observed deep-fetch spread from 7.15x to 1.68x
- Rankings are configuration-specific and version-sensitive, not a universal order
- Open harness + dual-engine case study (PostgreSQL, MySQL) demonstrate the protocol