

Highlights

- Nine documentation-selected layers plus tuned baselines on PostgreSQL and MySQL
- Deep-fetch spread of 7.15x narrows to 1.68x under a same-SQL control
- Rankings are version-sensitive: an earlier five-core Prisma result did not recur
- High-load tail-latency gaps largely dissolve at matched utilization
- Open harness verifies byte-identical responses across layers before measuring