

Highlights

- Nine idiomatic layers plus tuned native baselines on PostgreSQL and MySQL
- Deep-fetch spread of 6.5x collapses to 2.2x when layers run identical SQL
- Prisma matches or leads native throughput at five times the CPU of other layers
- Read rankings transfer across engines; write rankings do not under default fsync
- Open harness verifies byte-identical responses across layers before measuring