

A Comparability Protocol for Benchmarking Relational Database Access Layers in Express.js

Mateusz Miotk^{a,*}

^a*Faculty of Mathematics, Physics and Informatics, University of Gdańsk, Wita Stwosza 57, 80-308, Gdańsk, Poland*

Abstract

Context: Benchmarks of relational database access layers inform technology-selection decisions, yet can time implementations that returned different results, or report latency at unequal proximity to saturation. Whether the compared configurations are comparable is rarely established. **Objectives:** We propose an executable protocol for deciding when configured implementations are comparable for a stated estimand, and what their measurements support, then exercise it in one empirical instantiation rather than validating it independently. **Methods:** The protocol requires task-derived read checks, cross-implementation equivalence, post-write state validation, a predeclared configuration-selection policy, and separate capacity and latency operating points. An open Express.js experiment compares nine policy-selected configurations and two tuned native references across PostgreSQL, MySQL, and five access patterns in a 25-run randomized-block campaign, with a second same-host campaign testing rank stability over time and a single-rater audit of nine external benchmarks. **Results:** The widest spread is the tested deep fetch, $7.01\times$ on PostgreSQL and $5.12\times$ on MySQL; forcing common SQL through raw facilities leaves $1.71\times$ and $2.03\times$, a compound contrast isolating no single mechanism. High-load and matched-utilization p99 are different quantities, not one optimized against the other; the matched-utilization figures rest on

*Corresponding author.

Email address: `mateusz.miotk@ug.edu.pl` (Mateusz Miotk)

five runs and are secondary. Four cross-stack rank reversals recur in both campaigns, all involving the one treatment whose MySQL driver differs, and no other pair clears the same bar without it, while whole-ordering agreement stays imperfect either way. Only that treatment reverses durably at the stated bar, and the experiment does not identify the DBMS as the cause. **Conclusion:** The protocol is a candidate: a structured, auditable proposal with one worked instantiation rather than a validated methodology, and supports no universal ranking or library-intrinsic overhead estimate. Its audit has one rater, the adapters lack independent review, and both campaigns share one host, so independent applicability remains untested.

Keywords: benchmarking methodology, database access layer, object-relational mapping, Node.js, response-time tail (p99), empirical software engineering

1. Introduction

Express.js is among the most widely used web frameworks for Node.js, and almost every non-trivial service built on it reaches a relational database through some *access layer*. Access layers trade raw performance for ergonomics, type safety, and protection from pitfalls such as the N+1 query problem, where fetching a result graph issues one query per parent row instead of one for the graph [1], a pervasive anti-pattern (Section 2). Native drivers expose the wire protocol directly (`pg`, `mysql2`); query builders assemble SQL programmatically (Knex); and object-relational mappers (ORMs) map rows to objects, following patterns such as Data Mapper and Active Record [2]. Node.js back-end performance under load is a recognized empirical topic [3, 4, 5], but that literature varies the runtime and framework, not the access layer.

The most visible comparisons are vendor-authored. Peer-reviewed studies in the sources searched cover up to three ORMs in a controlled layer comparison; a broader Node-Bench study includes five access products but also changes frameworks and environments, while the dual-engine precedent holds the access

layer fixed [6, 7, 8, 9, 10, 11]. Supplement Table S51 records these coverage limits. The sole broad suite reporting a tail gives one P95 at one load point [12], and the academic studies time inside the process rather than through an HTTP framework [10, 13]. These gaps motivate the coverage of this experiment, not a priority claim. A more basic gap is methodological: the reported numbers are not established to be *comparable*: a layer returning fewer fields or erroring faster than it works can look faster, a library’s chosen query strategy is conflated with its own cost, and a single, arbitrary load point conflates capacity, tail latency under demand, and per-request latency. Establishing that comparability is the problem this paper takes up. The problem is a software-engineering one: where such benchmarks inform a technology choice, a comparison not established to be comparable can misdirect that choice rather than merely misreport a number.

This paper addresses these gaps with a *comparability protocol* (a structured, executable proposal rather than a validated methodology) and one non-vendor-authored empirical instantiation of it (Section 3.1 defines the term): an *identical* Express application served through all nine access layers (two native drivers, Knex, and six ORMs: Drizzle, Prisma, Sequelize, TypeORM, Objection, MikroORM) over five selected access patterns: point read, keyset range scan, deep fetch, per-author aggregation, and insert. The five stress distinct aspects of the access layer and are not a representative sample of Express/database workloads. Seven layers run against *both* engines under an identical schema, dataset, and pool size; two tuned native-driver baselines are disclosed reference points, not an optimization ceiling. Every (layer, engine, pattern) reports throughput and the high-load p99 under equal closed-loop demand at a fixed 50-connection point over 25 runs. That point sits near saturation, where the p99 largely tracks capacity and queueing, so we also report the tail at *matched utilization*. performance of each layer (RQ1–RQ3), and a *common-SQL raw-path sensitivity* contrast reporting the residual under identical SQL, which changes query formulation, protocol, preparation, round trips, and mapping together rather than isolating any one. The full harness (adapters, seed data, an autocannon-driven [14] load runner, a specification-derived expected-result

oracle, differential-equivalence checks, and write-state validation) is released as an open replication package.¹

The experiment answers three research questions, each scoped to the benchmarked implementations under the specified operating point:

RQ1 (*performance difference*): How large is the throughput and response-time-p99 *difference* of the benchmarked implementations relative to the native-driver baseline, and how does it vary across the five access patterns? We keep distinct three quantities a fixed operating point conflates: *capacity* (estimated saturating throughput), *high-load latency* (p99 under equal demand), and *latency at matched utilization* (equal fractions of each layer’s own capacity).

RQ2 (*backend-stack transfer*): Does the observed layer ordering transfer between the tested PostgreSQL and MySQL backend stacks, where changing the DBMS can also change its supported adapter and lower-level driver? The question concerns transfer, not causation: the design cannot attribute a failure to transfer to the DBMS rather than to the adapter or driver that the backend brings with it. We answer it at two levels: whether a specific pair reverses durably under a stated rule, and whether the whole ordering agrees, reported descriptively.

RQ3 (*pattern spread*): Which of the five tested patterns shows the largest native-relative spread in this configuration?

Supplement Table S47 states, per question, the quantity the design estimates and the quantity it does *not* identify.

This paper’s primary contribution is a **comparability protocol for access-layer benchmarking**, specified normatively in Section 3.1 (inputs, stages, cell-

¹Harness, deterministic seed, all adapters, raw per-cell measurements, and the table/figure-generating scripts: <https://github.com/mmiotk/express-db-access-performance>. The archived release and its DOIs are given in the Data availability statement.

admission criteria, interpretation, and limits): a concrete discipline for establishing and reporting comparability.

- **Semantic admission.** A seed-specification oracle checks expected read results without native-driver responses, differential checks establish mutual equivalence, and every timed insert must reach the intended database state. All precede admission of the corresponding timing, so a layer that returns less, writes wrong, or errors faster than it works cannot appear faster (Section 3).
- **Treatment definition and common-SQL sensitivity.** The treatment is fixed by a predeclared, reproducible selection rule (*policy-selected documented* here; the protocol privileges none, admitting expert-tuned, usage-frequency, or vendor-recommended policies provided the rule is fixed independently of the outcome). Reproducible selection is not representativeness: it prevents result-dependent tuning but establishes nothing about whether the configuration is optimal, common, or preferred in production. A *common-SQL raw-path sensitivity contrast* then reports what remains once SQL is held fixed, a diagnostic contrast, not a causal attribution.
- **Operating-point separation.** Capacity, high-load tail under equal demand, and tail at matched utilization are reported as *distinct* quantities, not conflated at one load point.

These constituent controls are established performance-evaluation and testing ideas [15, 16, 17]; CYNTHIA already applies differential equivalence across ORM implementations [18]. The contribution is not any individual control but their integration and operationalization for relational access-layer comparison. We *exercise* that discipline through an in-experiment ablation and a single-rater audit of nine external benchmarks; neither is independent validation, and Section 5 sets out what the exercise does and does not establish.

Two contributions should be kept apart. The *methodological* contribution is the protocol and its checklist: proposals for how access-layer comparisons should be conducted, reported, and audited. The *empirical* contribution is one

instantiation of them, scoped to the conditions and software versions of Section 3. No result is offered as a property of a library, and the study measures no non-performance quality such as maintainability, type safety, migration tooling, or ecosystem maturity.

2. Related Work

We survey prior work along four *coverage* axes (taxonomy, engine, joint throughput/tail reporting, and vendor involvement) spanned by this experiment (Supplement Table S51). We separately audit which comparability controls each source reports. A structured scoping search (not a systematic review) of the ACM Digital Library, IEEE Xplore, Scopus, and Google Scholar (2006–July 2026) retained empirical comparisons of relational access layers or engines, with search details archived in the replication package (`notes/related-work-search.md`); being scoping rather than exhaustive, we scope every gap claim to the sources searched.

2.1. *Object-relational mapping and the access-layer spectrum*

Access layers bridge the object-relational impedance mismatch between the relational model and an application’s object graphs [19], spanning native drivers, query builders, and ORMs following Data Mapper / Active Record patterns [2] (Section 1). Torres et al. [20] survey two decades of ORM patterns; their trade-off (productivity against a hard-to-predict runtime cost) motivates what we measure while holding the store family (relational rather than non-relational [21]) fixed. The treatment is the configured access-layer implementation bundle, including its selected API, SQL, protocol, and mapping path; the study does not isolate a library binary from those choices.

2.2. *ORM performance and data-access anti-patterns*

A parallel line studies ORM performance in reverse, repairing the inefficient SQL frameworks emit. The costliest is the N+1 query problem (Section 1);

it and related redundant-query anti-patterns pervade real-world web applications [22, 23, 24], their impact quantified in ORM-backed systems [25, 1] and often *removable* by query synthesis, batching, or caching [26, 27, 28] (with a JavaScript `async/await` analogue [29]). Closest to our design, Lorenz et al. [30] show the best mapping strategy depends on database technology, van Zyl et al. [31] that it is tuning-sensitive, and a recent study adds an energy dimension [32]. Overhead is thus configurable, not fixed, as our per-pattern results confirm, but this line never compares drivers, query builders, and ORMs head-to-head under load, our aim.

Cross-implementation ORM equivalence checking has a direct precedent. CYNTHIA generates an abstract query, translates it to equivalent queries for five ORMs, executes them over four relational DBMSs, and differentially compares the results to find correctness bugs [18]; matching results alone do not establish correctness against an independent task specification. CYNTHIA is a testing method, not a performance-benchmark admission protocol: it does not use equivalence to decide which timing cells may enter a benchmark, define a premeasurement rule for which implementation represents each library, or separate equal-demand capacity effects from matched-utilization latency. Our claim is restricted to using specification evidence and cross-implementation equivalence as benchmark admission, joined with predeclared treatment selection and operating-point separation.

2.3. Peer-reviewed access-layer comparisons

Peer-reviewed head-to-head access-layer measurement is scarce. Colley et al. [6] demonstrate selected Entity Framework-generated and hand-written SQL in a Contoso application; it is an ORM/non-ORM illustration, not the eight-ORM cross-language benchmark attributed to it in an earlier draft. Three recent studies partially address this: one in *Procedia Computer Science* compares Prisma against optimized raw SQL over eight query patterns on PostgreSQL [7] (differing versions, hardware, and workload, so we neither build on nor dispute its figures); a second compares Prisma, TypeORM, and Sequelize

across four relation types, also on PostgreSQL, by response time, memory, and CPU [8]. Pratama and Raharja [9] cross Node-Bench results for five access products with eleven Node.js frameworks and three virtualization/container environments. That study directly varies access products, but framework, access path, and environment also change across cells, so it does not isolate an access-layer estimand. The most recent study times internal query stages of three ORMs rather than client-observed requests [10] and reports its own concurrency-dependent ordering. A further Express study optimizes one stack rather than comparing access layers [13].

2.4. Database-engine, SQL/NoSQL, and Node.js performance

A second, disjoint line benchmarks the layers below and above the access layer without varying it. Salunke and Ouda [11] compare PostgreSQL and MySQL under standard OLTP workloads (the only dual-engine precedent found), but benchmark the *engines themselves*; SQL/NoSQL comparisons [33] likewise measure the store, not the layer. On the Node.js side, work establishes Node’s web viability [3], tracks back-end drift under sustained load [4], and compares frameworks broadly [5]. The Node-Bench comparison above shows why a product count alone is insufficient: its framework–library–environment cells do not hold a common access-layer task fixed.

2.5. Standard benchmarks and benchmarking methodology

Standard benchmarks define how systems are measured: YCSB [34] for cloud-serving stores and TPC-C/TPC-E for OLTP. These target the *engine*, not the driver, query builder, or ORM this study varies, but set the bar for controlled, repeatable load. A methodological literature explains why single-metric reporting fails: Fruth et al. [16] and Raasveldt et al. [15] catalog benchmarking pitfalls, Dean and Barroso [35] show the tail drives user-perceived latency under fan-out, and Li et al. [36] trace its sources. Such measurement is often ad hoc [37, 38] and hard to reproduce [39], motivating a harness built for relevance, repeatability, and fairness [40, 41]. Following that guidance, we report throughput *and*

p99 for the same run: a pairing the general literature recommends, which we did not find in the access-layer benchmarks surveyed here.

2.6. Practitioner and vendor benchmarks

A larger body comes from practitioners and vendors, each authored by a party with a stake in a compared product. Drizzle publishes a Northwind micro-suite (nine targets) and a k6 load test reporting req/s, average latency, P95, and CPU [12]. Prisma’s site compares Prisma, Drizzle, and TypeORM by median query latency over 500 iterations, reporting neither throughput nor any percentile [42]. IMDBench adds Sequelize and node-postgres, reporting throughput and latency on three CRUD operations [43]. Together these span more products than the peer-reviewed literature, but none of these three combines dual-backend-stack coverage, response-time p99, capacity characterization, and frozen source-located treatment provenance.

2.7. Positioning

Audit corpus.. The nine audited sources are those retained by the same scoping search, under stated criteria rather than convenience. Eligible were empirical performance comparisons of relational access layers, or direct PostgreSQL-versus-MySQL comparisons, bearing on the four positioning properties; peer-reviewed studies, vendor suites, and practitioner benchmarks were all eligible, since the audit concerns what a benchmark reports, not who published it. Excluded were non-empirical ORM papers, non-relational-only studies, framework benchmarks holding the access layer fixed, and pages with unverifiable versions and no reproducible method. CYNTHIA is retained for methodological positioning but is not among the nine, being a differential-testing tool rather than a performance benchmark. Coding was performed once, by the author, from a stable locator for every judgement; `not_reported` records what the inspected source does not state, never a finding that the authors omitted the practice. The corpus is not exhaustive, so every gap statement below is scoped to it: we

do not claim that no other benchmark implements these controls. Full criteria, codebook, screening record, and per-judgement evidence are in the online supplement and the replication package.

In the sources searched, no access-layer benchmark reports all seven protocol dimensions, unsurprisingly, since the seven are this paper’s own construction: the audit measures distance from a new standard, not failure to meet a known one. The source-located audit in Supplement Table S37 records partial controls rather than treating non-reporting as proof of absence, and we found no study combining broad product coverage, both engines, joint throughput/tail reporting, and non-vendor authorship, each covers at most two or three (Supplement Table S51), the strongest claim the scoping search supports, not one of priority.

But coverage is not the contribution: the protocol operationalizes established experimental principles for this genre, not a new one. Kounev et al.’s systems-benchmarking criteria [44], Hasselbring’s software-engineering standard [45], and Raasveldt and Mühleisen’s fair-benchmarking guidance [15] supply its principles at a deliberately general level. They leave an access-layer study to operationalize *what code represents a library*, how competing implementations are checked for semantic equivalence, and how strategy-sensitive results are contrasted. Our protocol supplies one concrete operationalization: an arbitrary but predeclared treatment policy; a common-SQL raw-path sensitivity contrast; and a per-cell gate requiring specification-conformant, mutually equivalent responses and verified write state *before* a timing is admitted. Correctness-before-timing has precedents (SPEC validation, TPC ACID, and SQLancer [46]); cross-implementation equivalence has the direct precedent discussed above. Our narrower contribution is stated in Section 1. Together these constructs provide a concrete discipline for establishing and reporting comparability among competing implementations of one task. Its evidence is the in-experiment ablation and a single-rater retrospective audit of nine external benchmarks (Supplement Table S37); this demonstrates utility and applicability, not independent inter-rater validation or invention of the underlying principles.

3. Study Design

In Goal–Question–Metric terms, the study analyzes configured relational-access-layer implementations for the purpose of comparing throughput and response-time p99 under PostgreSQL and MySQL back ends [47]. The view-point is the benchmark analyst evaluating policy-defined treatments; it is not a model of observed developer behavior. This is a *controlled benchmarking experiment* on one synthetic service, dataset, and environment, not inference about ORMs as a population. We use *experiment* in the sense of controlled software-engineering guidance [48, 49]: the treatments are assigned and executed automatically rather than observed in situ, the factors are fixed by design, the run order is randomized in blocks, and each cell is measured repeatedly. That design is what licenses the paired within-campaign analysis; it also bounds external validity to the single instantiation measured here. The experiment instantiates the protocol with three decisions: semantic admission before a cell’s timing is used, a predeclared treatment-selection policy, and explicit operating points. Admission gates the *use* of a timing, not always its execution: several admission checks were added after the campaigns they screen (Supplement Table S48).

The primary RQ1–RQ3 treatment is the *policy-selected documented configuration*: the path selected mechanically from frozen official documentation under the rule below. The protocol itself does not privilege documentation; another study may preregister an expert-tuned, usage-frequency, or vendor-recommended policy. A *common-SQL raw-path sensitivity contrast* then has every portable layer execute the same SQL through its raw facility. SQL, API level, protocol, preparation, round trips, and mapping still change together, so the contrast measures sensitivity to standardizing the whole bundle; it decomposes nothing and bounds nothing (Supplement Table S42).

A fixed concurrency imposes equal *demand* on layers of unequal throughput and so drives them to unequal *utilization*. RQ1 is therefore characterized under three operating conditions, not one privileged point, the protocol’s *operating-point separation*. *Equal external demand* is the fixed 50-connection *high-load*

point covering the full five-pattern, two-engine matrix; a per-pattern concurrency sweep (ladder 1–200, both engines; Supplement Table S35) places every pattern’s throughput knee at or below 50 connections, so this point sits at high utilization of each pattern’s own capacity (median 98–100%), a capacity-and-queueing regime rather than a sub-saturation latency comparison. The two utilization-dependent conditions need a per-layer capacity target and are shown on the deep fetch: *equal utilization* is the coordinated-omission-corrected open-loop sweep offering each layer a fixed fraction (50/70/85/95%) of its *own* estimated capacity, and an *equal compute budget* is the exploratory equal-CPU check. These answer different questions and need not agree, so all three are reported. The capacity proxy $\hat{C}_{50}$ is the median throughput of the 25-run, 50-connection primary cell; it remains an estimate, evaluated against the bootstrap interval and the separately measured sweep maximum in Supplement Table S45. The 50% and most 70% cells remain below saturation; 85% and 95% are interpreted only as near-saturation trends (Supplement Tables S21–S22).

3.1. The comparability protocol

We state the protocol independently of this experiment, then instantiate it.

Comparability, operationally.. Configured treatments are *comparable for a stated estimand* when each satisfies the predeclared semantic-admission and treatment-definition criteria that estimand requires, and the differences that remain between them are the differences the estimand is about. Comparability is therefore a property of a claim, not of a pair of systems: treatments comparable for one estimand need not be comparable for another, as this study’s selected-configuration and common-SQL estimands illustrate. The definition deliberately does *not* require identical SQL, round-trip counts, internal algorithms, or framework responsibilities; requiring any of these would suppress the design choices an access-layer comparison exists to measure. What admission establishes is narrower: every timed treatment performed the declared task, returned the declared result, and left the declared state, so a

difference in timing is not a difference in work done. Admission is a floor, not a certificate of equivalence: admitted treatments stay unequal in ways the protocol records rather than removes, and no result below should be read as comparing otherwise-identical systems.

The seven stages are cited throughout by stable identifiers: *M1* semantic admission and *M2* treatment definition are mandatory; *R3* common-SQL raw-path sensitivity, *R4* capacity characterization, *R5* operating-point separation, *R6* resource accounting, and *R7* implementation-waste evidence are recommended. They are not interchangeable: each addresses a distinct threat to interpretability, and a study may satisfy some and not others, which narrows the admissible claims rather than invalidating the reported numbers. The full normative specification (stage definitions, the cell-admission rule, applicability limits of the output comparator, the six-level compliance scheme, and the level each workload cell attains here) is in the online supplement, with a machine-readable encoding in `protocol-checklist.yaml`; Figure 1 states the inputs, stages, cell-admission rule, and output interpretation. The stages cover the admission-through-interpretation dimensions addressed here; they are not asserted to be a complete taxonomy of every threat to access-layer benchmarking. Here M1–M2 cover the full matrix, R4 every pattern, R3/R5/R6 the deep fetch, and R7 the measured source.

Protocol development versus protocol validation. The protocol was not specified in full before the measurements it governs. We separate *development evidence* (the protocol finding defects in this study’s own harness) from *internal demonstration* (the worked instantiation reported here) and *external validation* (independent analysts applying the protocol to their own comparisons), which has not been performed. Supplement Table S48 dates each element against the campaigns; four postdate the accepted campaigns and are therefore retrospective, applied to archived state rather than gating the timed runs. That the protocol exposed real defects shows the checks do work, but the defects were found by the author who designed them, so this is not independent validation.




Figure 1: The comparability protocol. Semantic admission (M1) and treatment definition (M2) form the mandatory core. The recommended stages R3–R7 license common-SQL, capacity, utilization, resource, and implementation-provenance interpretations, and may be scoped to a representative workload point rather than the full matrix. A finite gate supplies task-conformance evidence; it does not prove arbitrary-input correctness. For reporting compliance, M1 and M2 together form one mandatory validity-core level and R3–R7 define five claim-specific extensions, giving the six-level scheme of Supplement Table S40.

3.2. Factors and treatments

We vary three factors. The *access layer* has eleven levels across three tiers. Two native drivers (`pg`, `mysql2`) and their two engine-specific tuned baselines (`pg-tuned` reuses named prepared statements, `mysql2-tuned` the binary protocol via `execute()`) each run only on their own engine; the query builder Knex and the six ORMs (Drizzle, Prisma, Sequelize, TypeORM, Objection, MikroORM) run on both. The tuned baselines are reference points, not selected treatments, leaving *nine* policy-selected layers. The 7 portable layers run on both engines and the 4 native/tuned on one, so $4 + 7 \times 2 = 18$ layer $\times$ engine cells $\times$ 5 patterns = 90 measured configurations. We classify a library by its dominant interface, not its self-description (Section 1). The nine policy-selected layers are a deliberately broad product cross-section, not balanced sampling of the three tiers and not an exhaustive catalogue: each is maintained, used, and, for the portable tiers, runs on both engines. It excludes non-portable PostgreSQL-only Slonik and pg-promise and query-builder Kysely (covered by Knex; `METHODOLOGY.md` gives the criteria). *Non-vendor authorship* means only that no evaluated library’s maintainers were involved; the consequential author choices (selection, schema, pool size, tuning) remain and are disclosed. All adapter code and treatment assignments were produced by the single author, a limitation analysed in Section 6.

The *backend-stack* factor has two nominal DBMS levels (PostgreSQL 18.4 and MySQL 9.7.1), but switching it can also switch the supported adapter, driver, wire protocol, and defaults, so RQ2 estimates transfer between stacks, not a pure DBMS effect. The *access pattern* has five levels, realized as HTTP endpoints (Supplement Table S39). Supplement Table S29 records which factors are held constant, vary, or are uncontrolled on this single virtualized host, the reason we report this configuration rather than general library performance.

3.3. Objects: schema, workload, and data

The workload is a minimal blog domain (`authors 1-* posts 1-* comments *-1 authors`). The five endpoints implement five selected access patterns [34]

(Supplement Table S39), selected as in the Introduction: a keyset page (primary-key cursor, not a large `OFFSET`), a single-row read, a one-to-many list, the deep fetch (a post with its author and every comment, each with its own comment-author), and a per-author aggregation. A deterministic seed (2,000 authors, 100,000 posts, 1,000,000 comments) is loaded by the native driver, independently of any layer under test, so both engines receive byte-identical row values; on-disk representation, index layout, and collation naturally differ. The working set fits in memory, so measurements emphasize the access layer’s per-request instruction, protocol, and mapping cost, disk-read latency largely removed, while engine-side execution, buffering, locking, and logging remain in every measurement [50]. Each request’s target identifier is drawn uniformly from the seeded range (posts 1–100,000, authors 1–2,000), spreading load across the working set; a skewed, production-like distribution is a planned variant (Section 6).

3.4. The adapter contract and $N+1$ control

Every access layer implements the same factory interface (five query methods plus a teardown), so the server and load runner are layer-agnostic and each returns an *identical result shape*. Each avoids the deep-fetch $N+1$ problem by its selected mechanism, a hand-written join for the native drivers, Knex, and query-builder-style Drizzle (tuned baselines reuse it), and relation eager loading (`include, populate, relations, ...`) for the data-mapper ORMs. The comparison is thus of each layer’s *selected configuration* under a fixed adapter-selection rule, not a fixed query plan: a measured difference reflects the SQL generated, the round-trip count, and result materialization. Server-side logging captures the exact SQL, round-trip count, and `EXPLAIN` plans; per-endpoint counts differ by design (Supplement Table S2).

The API was fixed by a rule recorded before the accepted campaigns (2026-07-22; Supplement Table S48) and applied to them unchanged: the eager-loading API each layer’s version-compatible official documentation presents first, ties broken (as for Drizzle’s two builders) toward the lower-level API that the library documents as its base layer. Such a tie shows why the rule must be fixed

in advance; it is not evidence that the rule picks the better builder. This is an intentionally artificial, predeclared *policy*, not evidence that the chosen API is performance-optimal, the most common production choice, or what an expert would deploy, one instance of the protocol’s predeclared treatment rule (Section 3.1), which privileges none; a study targeting expert-tuned usage would substitute its own (corroboration and limits: Table S33, Section 6). Reproducible selection is also not yet independently *reproduced*. The result-blind treatment-selection packets are prepared but unreturned: all four review slots, covering treatment selection, the adapter audit, and the external protocol audit, stand pending (`notes/reviewer-packets/review-register.csv`). Whether a second analyst applying this rule to the same frozen pages would select the same APIs is therefore untested, and establishing it is future work.

Supplement Table S16, `METHODOLOGY.md`, and the rubric (`notes/documentation-selection.md`) record, per layer, the selected API with its documentation page, version, and URL, its round-trip count, the documented alternative not used, that logging, hooks, and validation are off, and the freeze date. The artifact ships the exact pre-freeze page captures with timestamps and hashes, so reconstruction does not depend on today’s website; a capture proves the page state at its recorded timestamp, not that it was unchanged through the freeze, and the record therefore separates the freeze date from the evidence date (mechanics in the online supplement). If official pages conflict, the version-matched relations/eager-loading page controls over a quick-start or marketing page; remaining equal prominence invokes the declared documented-base-layer tie-break, and an unresolved tie must be reported as ambiguous.

For the protocol’s *common-SQL raw-path sensitivity contrast*, every adapter also executes the same two-statement deep-fetch SQL and row mapping through its raw-SQL facility, with the same `EXPLAIN` plan per engine. Server-side capture confirms the statements are byte-identical across every layer on both engines; only the wire protocol differs (Supplement Table S14). A secondary loading-strategy sensitivity switches only between documented, byte-identical deep-fetch alternatives; it does not redefine the primary policy-selected treatment or esti-

mate a library-intrinsic effect (Supplement Table S18).

Read admission has an expected-result layer independent of any timed reference: `bench/verify-spec.mjs` replays the documented seed PRNG and fan-out fixture in memory, without querying a database or importing an access layer, and derives exact fields, values, ordering, graph membership, and aggregates. All nine engine-compatible adapters pass 36,540 expected-result checks per engine, 73,080 in total (Supplement Table S38).

A separate differential layer asks whether implementations agree. Shared canonical constructors fix fields, key order, and scalar types, copying already materialized rows without joining, regrouping, sorting, or caching; their standalone cost is bounded in Section 6. Fixed probes and a randomized edge-input sweep compare serialized outputs across implementations, with no divergence. The specification and differential layers provide finite evidence for the declared task and tested inputs; neither is presented as proof for arbitrary inputs. A companion check validates database *state* rather than HTTP status: an out-of-band native-driver connection confirms, for every adapter on both engines, that the primary single-row insert wrote exactly the requested values and identifier, and for the adapters implementing the exploratory transactional method (five on PostgreSQL, four on MySQL) that the post and five comments commit atomically while an invalid final comment rolls back to no partial state; all declared methods pass. Per-adapter scope, the fan-out and timestamp conventions, and the fixture details are in `experiments/METHODOLOGY.md` and `write-admission.json`. The read cross-check caught two defects during development (a fan-out aggregation bug and a driver result-shape mismatch; Section 5).

3.5. Controls

Four controls hold the shared conditions fixed (connection pool, logical task, physical dataset state, and observer) so the *configured access-layer implementation bundle* is the only *deliberately varied* factor. They do not isolate the library alone: the bundled differences in SQL, query count, protocol, loading

strategy, and hydration are the treatment, not confounds to remove (Supplement Table S42). (i) *Pool size* is fixed at ten connections for every adapter, so the comparison reflects the access-layer choice, not pool tuning [51, 52]. Because the generator opens 50 concurrent connections against that pool, roughly forty in-flight requests always wait in each library’s acquisition queue; that queueing is deliberately measured and held identical across layers. (ii) *Common task semantics*: each endpoint implements the same declared input–output task, while SQL text, query count, plan, protocol, loading strategy, and hydration remain measured properties of the treatment. Aggregation is the exception, with a common correlated-subquery formulation; the separate common-SQL contrast standardizes SQL only for the deep fetch.

(iii) *Write isolation*: the declared reset rebuilds PostgreSQL from its seed template and uses a disclosed in-place delete/OPTIMIZE TABLE/allocator reset on MySQL, removing generated rows before warm-up, measurement, and the next cell. A post-review state oracle found that the historical MySQL rebuild could restore AUTO_INCREMENT below the cleanup floor, so generated rows could persist; a cross-engine gate exposed a second shared-fixture defect, an advancing PRNG that gave different comment-author assignments across engines despite identical comment counts. The corrected harness pins the allocator with a sentinel, verifies exact base and fan-out counts by idempotent preflight, and materializes the fixture once so the joined representation is byte-identical across engines. The affected pilot was rejected and a corrected-state campaign remeasured all five patterns for all 18 compatible treatment–stack pairs under one version of the instrument; the historical MySQL range-scan, aggregation, and insert cells are not used as clean evidence. Reset mechanics, sentinel values, and fixture hashes are in `experiments/METHODOLOGY.md`. (iv) *Observer separation*: the HTTP load generator runs in a process separate from the server.

3.6. Measurement procedure

Load was generated with `autocannon` [14], a closed-loop (constant-concurrency) HTTP/1.1 generator [53] recording a latency histogram. The

accepted corrected-state primary matrix contains 90 configurations with 25 repeated runs each; the superseded campaign remains provenance only. With INDEP=1, each replicate randomizes adapter-engine blocks and boots one fresh server process per block, so no application process crosses an adapter, engine, or replicate boundary; read endpoints within a block share that process, and database buffers plus host state are not reset. Each endpoint receives a discarded 15-second warm-up followed by one 12-second measurement at 50 connections. A forward/reverse-order write-cell check found no detectable order effect. The campaign records its ID, deterministic order seed, environment fingerprint, and exact source manifest; process, teardown, and rebuild mechanics are in `experiments/METHODOLOGY.md`.

Within a replicate, every layer and engine draws target identifiers from one generator seeded on endpoint and replicate, so all face the same identifier distribution, though not an identical served sequence, since the generator is shared with the warm-up and layers of unequal throughput enter the measured window at different offsets. The paired analysis is licensed by the block design, not by sequence identity. The inferential unit is the endpoint-level run; individual requests are subsamples. We report the median of the 25 replicate values of throughput and of the run-level percentiles p50, p90, p97.5, and p99.

Throughout, *p99* denotes this *median run-level p99*, not the p99 of the pooled request distribution, which the design does not estimate. Each run-level p99 is the generator’s histogram percentile over one 12-second window, supported by between about 5,700 and 91,000 requests depending on the cell’s throughput, and the reported value is the median of the 25 such estimates. Ten 60 s re-measurements of the deep fetch give 12-versus-60-second rank correlations of 1.00 on both PostgreSQL and MySQL; the largest cell-median shifts are 5.6% and 3.8%, respectively (Supplement Table S23). This is endpoint- and campaign-specific sensitivity evidence, not a general claim that 12 s suffices for p99 estimation. We sampled server CPU and peak RSS over the process tree, with database and load-generator CPU separately. Inserts were measured under each engine’s *default* durability (PostgreSQL `fsync`

and `synchronous_commit` on; MySQL `innodb_flush_log_at_trx_commit=1`, `sync_binlog=1`), the regime a deployment runs unless relaxed; a labelled secondary run relaxed all of these to isolate the durability mechanism (Supplement Table S3). To characterize behavior under load we swept the deep fetch across 1–256 connections for every layer and engine [54]; and a native `undici`-based constant-arrival generator drove the deep fetch under a fixed offered rate (250–4,000 req/s) rather than fixed concurrency, measuring from intended send time and clipping timeouts at their cutoff, a coordinated-omission-corrected design [55, 56] checking the closed-loop tail’s sensitivity to a different load model (Section 4).

3.7. Analysis

We separate *primary* outcomes, which carry the research-question claims, from *secondary* and *exploratory* outcomes reported descriptively as controls and sensitivity analyses (Supplement Table S34): the primary outcomes are throughput and the closed-loop p99 on the five patterns against both engines at the fixed operating point.

Because absolute throughput is hardware-specific, we analyze *relative* differences within an engine; a pattern’s *spread* is the untuned native driver’s median throughput divided by that of the slowest layer on that pattern (native-relative, so comparable across patterns). We report medians over the 25 runs, the coefficient of variation as a stability signal [17, 57], and a fixed-seed percentile bootstrap 95% CI for the median; these intervals capture within-campaign, run-to-run variability on this host and configuration, not variation across machines, versions, or deployments (Section 6). The design is a *randomized block* (within each replicate every layer is measured once against the same seeded identifier distribution) so adjacently ranked layers are compared with *paired* tests on their per-replicate throughput ratios (a two-sided sign-flip permutation test cross-checked with Wilcoxon signed-rank), reporting the geometric-mean paired ratio, its bootstrap CI, the paired dominance, and a TOST against a $\pm 5\%$ margin (author-selected and application-dependent, not universal; supplement; the

margin was fixed on 2026-07-13, before the accepted campaigns but after pilot data existed, so the TOST is reported as a sensitivity analysis and not as confirmatory equivalence evidence) for the closest pair; we foreground these effect sizes and interval estimates over p -values [58, 59]. Because the seven adjacent pairs follow the *observed* ranking, their significances are a post-hoc, descriptive robustness check, not a confirmatory family.

All tests operate on the 25 paired run-level values; no interval is computed over pooled requests. The sign-flip test is valid under the assumption that, when the null of no paired difference holds, the per-replicate log ratios are symmetric about zero and their signs exchangeable across replicates. Blocking makes this credible rather than guaranteed: replicates share a schedule and an environment, and randomized order within a replicate protects against drift and position effects. Randomization here is over *measurement order*, so it supports temporal balance, not causal identification; the treatments are pre-existing implementations, not interventions assigned to units, and no result should be read as a causal effect of a library on throughput. Pairing by replicate therefore controls the temporal and environmental block and the seeded workload design, not the individual request stream. The Wilcoxon signed-rank result is reported as a robustness check on the same 25 paired values, not as independent confirmation, since it reuses the data and design of the sign-flip test.

Rank correlations are descriptive: the seven layers are the benchmarked population of configurations, not a random sample from a population of software systems, so a coefficient summarizes agreement in this set and supports no inference to layers outside it. Adjacent-layer p99 gaps of one or two milliseconds sit at the millisecond measurement floor and are not read as real; conclusions rest on the large-ratio patterns, chiefly the deep fetch. The full construction (procedure and seed, tie handling, the layer $\times$ engine interaction test, rank correlations, equivalence margin) is in the statistical-methods notes ([notes/supplement-methods](#)); paired p99 comparisons are Supplement Table S30.

3.8. Experimental setup

All measurements were taken on a single host, a virtualized Intel Xeon Gold 6140 instance (16 vCPUs, single NUMA node, 126 GB RAM; Linux 6.17; Node.js 24.18.0; Express 5), PostgreSQL 18.4 and MySQL 9.7.1 local. The fourteen labelled secondary experiments (order-invariance, durability, same-SQL, open-loop, fan-out, equal-CPU, concurrency, utilization, pool-size, cluster, mixed-workload, wait-event, post-restart, longer-window tail) ran 2026-07-16 to 2026-07-18. The accepted corrected-state primary matrix was measured overnight on 2026-07-24 and the same-host validation campaign on 2026-07-25, each dated by its archived environment capture; the superseded pilot dates from 2026-07-13 (Supplement Table S48). The secondary experiments therefore predate the accepted primary matrix rather than sharing its window, and the clean-evidence scoping below applies to them as it does to the pilot. Each library is pinned to the *latest stable release compatible with the harness adapter contract* as of the freeze date (2026-07-15), recorded from the committed lockfile and an automated environment capture (Supplement Table S11); only Sequelize invokes the earlier-release exception (its version-7 line is a prerelease, so the current 6.x stable line is used). Because access-layer performance is version-sensitive, this pins a dated snapshot rather than a permanent ranking (Section 4).

3.9. AI-assisted research tooling

AI assistance covered research software and data analysis, not language editing alone, so it is documented here as part of the method; the publisher’s separate manuscript-preparation declaration appears before the references.

Tools, models, and provenance. The sole assistant was Claude (Anthropic), via the Claude Code command-line interface, in three model versions: Claude Opus 4.8, Claude Fable 5, and Claude Opus 5. No other vendor’s assistant contributed software, analysis, or measurements. OpenAI Codex was used once, after the manuscript was otherwise complete, for an independent critical pre-submission

review; the author evaluated its comments and is responsible for which were acted on. Assistance is recorded per commit: 215 of 223 commits carry a **Co-Authored-By** trailer naming the model (159 Opus 4.8, 27 Fable 5, 29 Opus 5), counted from the trailers themselves rather than message text. The other 8 predate the convention and include two squash-merges of assistant-authored branches and the largest revision commit, so the count understates assistance.

Affected components. Assistance covered the eleven adapters and harness scaffolding, the table and figure generators, the statistical estimators (**bench/stats.mjs**), and manuscript prose. Only the estimators carry analytical rather than scaffolding logic: the generators deterministically transform archived records into table cells, and the adapters satisfy a fixed contract (Section 3.4).

Review and independent validation. Author review is not the safeguard; each category carries a check independent of the generated code. The estimators are unit-tested against hand-computed values and closed-form properties: Wilcoxon signed-rank against a hand-computed z , bootstrap ratio intervals on known-answer inputs, and permutation p -values checked for symmetry and the $1/(B+1)$ floor (**bench/stats.test.mjs**). Every *measured* table cell traces through a committed generator to an archived per-cell record; **MANIFEST.md** marks the authored analytical items, which report no measurement and lie outside that guarantee. The adapters are gated by the specification-derived oracle, which replays the deterministic seed without importing an adapter or reading timed native output, and by differential and write-state checks (Supplement Table S38).

Extent of assistance in analysis. The candidate procedures, the paired sign-flip permutation test and its Wilcoxon cross-check, the geometric-mean paired ratio and its bootstrap interval, Cliff’s delta, the equivalence test and its $\pm 5\%$ margin, and the blocked interaction test, were proposed with assistant support, as was prose describing what they compute. Each is a named, published method, none generated or bespoke. The author chose among the candidates against

performance-evaluation guidance [17, 58, 59], verified each against the blocked design, wrote every interpretive claim, and checked every reference against its source, taking full responsibility for all of them.

Author independence therefore does not secure this analysis; the mechanical checks do, and a reader who discounts our review can still run them. Their limit is exact: they establish that the estimators compute what they claim and that each measured number came from an archived measurement. *No* check bears on whether the chosen procedures suit the question, or whether the claims drawn from them are sound; those rest on author judgement alone.

3.10. Replication package

The harness (the Express application, the adapter contract, all eleven adapters, the deterministic seed, the `autocannon` matrix runner, the semantic-admission checks, and the scripts regenerating every table in Section 4) is released so the full matrix re-runs with one command against a containerized or local PostgreSQL and MySQL.

Two reproduction statuses are kept distinct throughout, and the distinction is not cosmetic. An author-run *reconstruction* regenerates the archived tables and analysis files from the archived raw records and verifies their hashes; it establishes that what is archived is intact and regenerable, and it can be repeated by a reader without a database. *Independent reproduction* would require another team re-executing the measurements on its own hardware, and none has been performed. No reconstruction reported here re-executes a measurement, and none is offered as evidence that the measurements would recur elsewhere; the dates and counts of the reconstructions run are in the Data availability statement.

4. Results

Tables report throughput (req/s, higher better) and tail latency (p99 ms, lower better) per access layer and engine. The two consequential patterns are in

the main text (deep fetch, Table 1; insert, Supplement Table S52); point read, keyset range scan, and aggregation are in Supplement Tables S12, S13, and S15. Each native driver has an untuned policy-selected row (`pg`, `mysql2`) and a tuned one (`pg-tuned`, `mysql2-tuned`) marking a hand-tuned reference point, not one attained without code changes. We compare *relative* differences within an engine. Two estimands recur: the performance of the selected configured treatment (RQ1–RQ3), which is not library-only overhead, and the common-SQL raw-path sensitivity contrast (Supplement Table S14), the residual spread when every layer runs common SQL through its raw facility (Section 6).

Table 1: Deep fetch: throughput (req/s, higher is better; median with a within-campaign 95% bootstrap interval over the 25 repeated runs — run-to-run variability within the source campaign of each cell, not across hardware or days) and response-time p99 (ms, lower is better; median run-level p99 with the same within-campaign 95% bootstrap interval) by access layer and engine. p99 is reported at 1 ms resolution, so cells differing by ≤ 1 ms are practically unresolved (see the paired significance analysis in the supplement).

Layer	Category	PostgreSQL		MySQL	
		req/s [95% CI]	p99 [95% CI]	req/s [95% CI]	p99 [95% CI]
pg	native-driver	3638 [3602–3662]	18 [17–20]	–	–
mysql2	native-driver	–	–	2416 [2413–2425]	28 [27–28]
pg-tuned	native-tuned	3700 [3656–3773]	18 [16–19]	–	–
mysql2-tuned	native-tuned	–	–	2398 [2386–2419]	25 [24–26]
knex	query-builder	2431 [2415–2459]	27 [25–28]	1982 [1963–2006]	30 [29–31]
drizzle	orm	1829 [1815–1846]	34 [33–35]	1301 [1290–1308]	48 [46–49]
prisma	orm	1175 [1160–1179]	52 [51–53]	630 [623–638]	91 [89–94]
sequelize	orm	978 [973–983]	60 [60–61]	880 [871–886]	68 [67–68]
typeorm	orm	1219 [1205–1233]	49 [48–50]	1017 [1008–1030]	57 [56–58]
objection	orm	1017 [1011–1029]	56 [54–58]	836 [824–849]	70 [68–70]
mikroorm	orm	519 [512–528]	108 [107–112]	472 [464–475]	121 [118–126]

4.1. RQ1: performance difference relative to the native baseline

The native implementation leads nine of the ten engine–pattern comparisons: `pg` or `pg-tuned` is fastest on all five PostgreSQL patterns, and `mysql2` on four of five MySQL patterns. The exception is the MySQL insert, where the four fastest

treatments span 1,766–1,816 req/s with broadly overlapping intervals and the native driver ranks fourth (Supplement Table S52); that ordering is practically unresolved rather than a defeat, and we do not read a leader from it. The current-generation ORMs sit mid-to-low on the read patterns; Prisma is near the bottom of point reads and aggregation on both engines (Table 1; Supplement Tables S12 and S15), retiring an earlier version’s headline of Prisma tying the native driver at a large CPU premium (Section 5).

`pg-tuned` and `mysql2-tuned` track their untuned counterparts within about 2% on the deep fetch (`pg-tuned` 3,700 vs. `pg` 3,638; `mysql2-tuned` 2,398 vs. `mysql2` 2,416). The one material gain is PostgreSQL aggregation, where named prepared statements add 11% (7,583 vs. 6,807). On MySQL the binary-protocol baseline is *slower* than `mysql2` on every pattern (by 0.4–4.7%), so `execute()` buys nothing at this operating point; both tuned baselines are therefore disclosed reference points, not optimization ceilings.

On the **deep fetch** (a post, its author, and all comments with their authors) the native-relative spread is the widest measured, $7.01\times$ on PostgreSQL (`pg` 3,638, 95% CI [3,602–3,662], down to MikroORM’s 519 [512–528]) and $5.12\times$ on MySQL (`mysql2` 2,416 [2,413–2,425] down to MikroORM’s 472 [464–475]). The native driver leads on both engines, the data-mapper ORMs (especially MikroORM) trailing furthest (Table 1); the comparison is paired (Section 3), every adjacent MySQL step large relative to its bootstrap interval (Supplement Table S36). Aggregation is the least layer-sensitive pattern: every layer forwards a single query for it, and it avoids the deep fetch’s nested-graph materialization, though the design isolates neither feature as the cause. The remaining per-pattern spreads are collected under RQ3 below.

4.2. Common-SQL raw-path sensitivity contrast

Server-side capture confirms common statement text and database plans across raw paths on both engines; API level, wire protocol, preparation, round trips, and mapping can still differ (Supplement Table S42). The deep-fetch native-relative spread is then far smaller, $1.71\times$ (PostgreSQL) and $2.03\times$

(MySQL) against the selected-configuration $7.01\times$ and $5.12\times$ (Supplement Table S14). The selected paths therefore differ more than these common-SQL raw paths, but the compound contrast assigns no share to any mechanism. A loading-strategy sensitivity check (Supplement Table S18) argues against a mere selected-strategy artifact: forcing the join-selected ORMs (Sequelize, MikroORM) to select-in makes them slower, while Objection’s batched select-in runs about $1.6\times$ slower than a single join on MySQL. Objection’s result is thus sensitive to the documented loading option; this comparison does not isolate a library-only cost.

4.3. RQ2: backend-stack transfer of the ranking

Switching the backend changes a bundle that can include the DBMS, supported adapter, lower-level driver, wire protocol, and defaults. RQ2 therefore concerns *backend-stack transfer*, not a causal DBMS effect. In the primary campaign, the seven portable layers have stack-ratio heterogeneity on every pattern (blocked interaction $F = 127.1\text{--}458.5$, permutation $p < 5\times 10^{-5}$, the $1/(B+1)$ resolution floor at $B = 20,000$; Supplement Table S28), but this within-campaign test does not establish durable rank reversals.

A second same-host campaign re-runs 25 repetitions of the seven portable layers for deep fetch, aggregation, and insert under an independently drawn block order (Supplement Table S46). It re-randomizes execution order only: the identifier generator is seeded on endpoint and replicate, so the campaign replays the same draws and cannot detect sensitivity to the workload sample. Deep-fetch ranks reproduce exactly on both stacks ($\rho = 1.00$); aggregation reproduces 5/7 PostgreSQL and 7/7 MySQL positions. MySQL insert is less stable: 3/7 positions, $\rho = 0.82$, and the leader changes from Drizzle to Knex.

We therefore promote a cross-stack reversal only when it occurs in both campaigns, exceeds the 5% practical margin on both stacks, and paired bootstrap intervals exclude equality in opposite directions in both. The margin was fixed before these campaigns (Section 3), but the composite rule was specified after both completed: it is a post-hoc screen applied uniformly to all 21 layer pairs

on each of the three patterns. No familywise correction is applied; recurrence across both campaigns is the screen’s error control. The reversals it promotes are the stable subset of an exploratory comparison, not a prespecified confirmatory test. Four pairs meet all criteria: on deep fetch, Prisma exceeds Objection and Sequelize on PostgreSQL but trails both on MySQL; on aggregation, Prisma exceeds Objection on PostgreSQL but trails it on MySQL; on insert, Prisma exceeds MikroORM on PostgreSQL but trails it on MySQL. All four therefore involve one treatment, and that treatment is the one whose lower-level driver differs between the stacks: Prisma runs the MariaDB adapter on MySQL because it ships no `mysql2` adapter, while every other layer uses `mysql2` (Section 6). The recurrence is consistent with backend-stack sensitivity as RQ2 defines it, but it cannot be separated from that driver substitution, and it is not evidence about ORMs as a group. Other changes in close ranks, including the identity of the MySQL insert leader, remain exploratory.

Because all four involve that one treatment, we repeated the promotion procedure unchanged on the six remaining portable layers (Supplement Table S49). No pair is promoted on any of the three patterns: the sole recurring non-Prisma deep-fetch reversal (Sequelize versus Objection, $0.96\times$ against $1.05\times$) falls below the 5% margin, and the recurring aggregation and insert reversals fail that margin or the interval criterion. The promoted-reversal evidence therefore rests entirely on the treatment carrying the driver substitution. Whole-ordering agreement is a separate and purely descriptive question, and it moves in both directions when that treatment is removed: ρ rises on deep fetch (0.86 to 0.94) and on insert (0.68 to 0.83 primary, 0.71 to 0.94 validation), and falls on aggregation (0.68 to 0.54; 0.64 to 0.49). It is imperfect on both stacks either way. We do not read that as a further reversal claim: the promotion rule excludes sub-threshold pairs by design, and we do not readmit them to argue a broader conclusion. What the leave-one-out contrast licenses, and what it cannot identify, is set out in Section 5.

The magnitude of stack transfer is also treatment-specific. Primary-campaign PostgreSQL $\div$ MySQL ratios range from 1.10 to $1.84\times$ on deep

fetch, 1.27 to 1.82 $\times$ on aggregation, and 1.55 to 3.06 $\times$ on insert (Supplement Table S28). These ratios and the recurring reversals show that the observed ordering cannot be assumed to transfer unchanged between the two tested stacks. They do not allocate the difference to the DBMS, adapter, driver, protocol, or defaults.

The primary MySQL insert medians illustrate why fine ordering is not a headline: Drizzle, Objection, Knex, and the native driver lie between 1,766 and 1,816 req/s, while several replicate distributions are visibly multimodal or heavy-tailed and reach 13.6% within-campaign CV (Supplement Figure S1). Prisma is the slowest policy-selected treatment at 894 req/s, giving a 1.98 $\times$ native-relative spread on a broad interval [1.65–2.09]. Under relaxed durability the native MySQL and Knex paths gain 2.42 $\times$ and 2.07 $\times$, with gains differing by treatment; concurrent wait-event sums are consistent with a shared commit-path contribution but are not a causal time fraction (Supplement Tables S3 and S19). PostgreSQL insert spans 3.18 $\times$ (6,018 to 1,894 req/s) and changes by at most 14% under relaxed durability. These compound interventions do not decompose backend-stack cost.

4.4. RQ3: which pattern spreads widest

The primary native-relative spreads on PostgreSQL are deep fetch (7.01 $\times$), keyset range scan (4.22 $\times$), insert (3.18 $\times$), point read (2.83 $\times$), and aggregation (1.82 $\times$). On MySQL they are deep fetch (5.12 $\times$), range scan (3.88 $\times$), point read (2.14 $\times$), insert (1.98 $\times$), and aggregation (1.87 $\times$). Thus the tested deep fetch has the largest spread and aggregation the smallest on both stacks. A full concurrency sweep places the median knee at 8–32 connections and the median utilization at 50 connections at 98–100% for every pattern–stack combination; the lowest individual values are 87–88% for two PostgreSQL patterns (Supplement Table S35). The fixed point is therefore a high-utilization comparison throughout, although it is not an identical utilization fraction for every layer. An exploratory fan-out sweep from 0 to 500 comments finds a roughly fixed multiplier rather than a spread that grows with breadth; it does not vary graph

depth, schema, or query mix.

4.5. Tail latency

At equal closed-loop demand, deep-fetch p99 spans 18–108 ms on PostgreSQL and 28–121 ms among the policy-selected layers on MySQL (Table 1). The run is the analysis unit; a 12 s slow-layer run supplies only about 60 observations to its top percentile, so we use the median of 25 run-level p99 values and show their raw spread. Ten 60-second repetitions preserve the ordering (rank correlation 1.00 on both engines; largest median shift 5.6% and 3.8%; Supplement Table S23), which is endpoint-specific evidence, not a general sufficiency claim.

At the stable nominal 50% capacity target, coordinated-omission-corrected p99 is 2.3–4.9 ms on PostgreSQL and 1.9–5.5 ms on MySQL, a different operating point and estimand, not an optimization applied to the systems; most 70% cells remain similarly small (Supplement Tables S21–S22 and S43). Re-expressed against the full-sweep maximum, nominal 50% falls at 46.6–51.0% and nominal 70% at 65.2–71.5% (S45). The 85% and 95% results are unstable near-saturation observations and do not support a convergence claim. The stable 50% result shows the large equal-demand p99 spread contains a material capacity-and-queueing component; it does not identify its share or imply a deployment-general sub-saturation bound.

The two operating points also *order the layers differently*, which is the sharpest consequence of separating them. Between the equal-demand and 50%-utilization p99 orderings, Spearman ρ is 0.64 on PostgreSQL and 0.29 on MySQL (Supplement Table S43). The native driver moves furthest: fastest on the tail at equal demand on both engines, it falls to third of eight on PostgreSQL and seventh of eight on MySQL once each layer is held at half its own capacity. On MySQL the reversal is unambiguous, `knex` 1.9 ms [1.8–2.2] against `mysql2` 4.7 ms [4.0–5.6], non-overlapping over five runs, while on PostgreSQL the native and leading portable ranges overlap (`pg` 2.6 [2.2–3.0], `knex` 2.3 [2.3–2.7]), so only the MySQL inversion is claimed. Reporting a tail at one arbitrary load point

would have recorded the native driver as the tail leader on both engines. This is direct evidence that the operating-point separation stage changes a conclusion rather than only its caveats, and it is why capacity and tail are reported as distinct quantities.

4.6. Measurement stability and effect sizes

Deep-fetch throughput CV reaches 4.5%, much smaller than its multi-fold spreads. The largest primary throughput CV is 13.6% on MySQL insert, whose raw replicates are shown because a median and interval can hide regime mixtures. All 90 cells contain exactly 25 run-level samples and zero timed errors, timeouts, or non-2xx responses. One pre-warm-up startup failure was retried once and recorded; no timed sample preceded that retry. The fail-closed roster validation requires 90×25 observations before accepting the file.

Predeclared native-driver contrasts are foregrounded; ranking-selected adjacent comparisons remain descriptive. Paired bootstrap intervals characterize only within-campaign variation, while the same-host validation campaign is reported separately. Differences of about 1 ms in p99 and throughput ratios inside the predeclared $\pm 5\%$ practical margin are not promoted as substantive ordering.

5. Discussion

Version-sensitivity. An earlier campaign found Prisma (then 5.22, on an in-process Rust query engine) tying the native driver on the deep fetch while drawing about five application cores. The *identical* harness on the re-frozen toolchain did not reproduce it: application CPU now stays near one core and Prisma sits at mid-to-low throughput (Section 4). The non-reproduction is a fact, its cause a hypothesis, since the re-freeze moved the whole toolchain at once, so the headline is version-sensitive [31, 60].

Implications for practice. RQ1 and RQ3 agree the selected-configuration difference is specific to the access *pattern*: the native-relative spread is widest on

the deep fetch, where a layer must hydrate rows and resolve associations rather than merely issue SQL (the object-relational impedance mismatch [19] made concrete), and narrowest on aggregation. Round-trip count alone does not explain throughput: on the PostgreSQL deep fetch, single-query MikroORM is slowest while two-query Drizzle outruns four-query Prisma (Supplement Table S2; Table 1). Result-graph materialization and the other bundled implementation differences are plausible contributors, but the study does not isolate them.

Single-row inserts and backend-stack transfer. Single-row insert behavior is both stack- and treatment-specific. In the primary campaign PostgreSQL spans $3.18\times$, whereas MySQL spans $1.98\times$ and several MySQL replicate distributions are multimodal or heavy-tailed. Relaxing durability raises the native MySQL and Knex paths by $2.42\times$ and $2.07\times$, but treatment ratios range from 1.02 to $2.44\times$; aggregate concurrent wait-event time can exceed wall time and is not a causal fraction. Together these controls are consistent with a shared commit-path contribution, not with attributing a performance floor or a fixed share to the DBMS (Supplement Tables S3 and S19).

The second same-host campaign preserves PostgreSQL insert ranks but few MySQL positions (Section 4), so we claim no durable fine ordering for MySQL insert. The recurring, material insert reversal is narrower: Prisma exceeds MikroORM on PostgreSQL and trails it on MySQL in both campaigns, with paired intervals excluding equality in opposite directions (Supplement Table S46).

Every promoted reversal involves Prisma, whose MySQL path uses the MariaDB adapter rather than `mysql2`, so backend-stack sensitivity here is inseparable from that driver substitution and describes one treatment, not ORMs as a class. Repeating the promotion procedure without that treatment promotes no pair at all (Supplement Table S49). *Observed:* every reversal clearing the promotion bar involves the one layer whose MySQL driver differs, and whole-ordering agreement without it rises on deep fetch and insert but falls on aggregation, to $\rho = 0.54$ and 0.49 . *Interpretation:* at the stated bar only the asymmetric treat-

ment reverses durably; agreement of the whole ordering stays imperfect on both stacks with or without that treatment, which we report descriptively rather than as a further reversal claim. *Not identified:* the design cannot attribute the promoted reversals to the DBMS rather than to the adapter, the driver, or Prisma’s own implementation, because holding out the treatment holds out the substitution with it. The hold-out followed observing that all four promoted reversals involved one treatment; its criterion is nevertheless a design fact, the sole asymmetric driver, fixed independently of the outcomes. This supports backend-stack sensitivity in the tested implementations while leaving cross-host transfer unresolved. A five-replicate, one-layer-per-tier six-statement transaction remains exploratory and does not extend the single-row-insert finding to writes as a class (Supplement Table S8).

Methodological lessons: a checklist for access-layer benchmarking. Several failures changed or could have changed the comparison; we distil them into a reusable checklist ([notes/supplement-methods](#)): an **aggregation fan-out** (a join before SUM inflating the aggregate) and a driver result-shape mismatch, both caught by specification/differential admission; **OFFSET pagination** collapsing on MySQL, easily misread as an access-layer weakness; **write-state contamination**, where MySQL allocated inserts below the cleanup floor; and a **query-formulation confound** where a mis-scoped pre-aggregation re-scanned the whole table, a query-plan artifact masquerading as an access-layer effect [61]. The write path needs its own gate: MikroORM 7’s dropped `persistAndFlush` made every insert throw a 500 faster than a real one, briefly *leading* MySQL until the non-2xx counter flagged it. Because a 2xx alone would miss a *silently* incorrect write, the gate also validates post-write state: field values, identifiers, row-count changes, and rollback. None is schema-specific; each extends the database-benchmarking pitfalls of Fruth et al. [16] and Raasveldt et al. [15] to the access-layer sub-genre, and at least one plausibly explains part of the gap between some vendor-reported and controlled benchmark results.

Evidence for the protocol contribution. Each stage had a concrete consequence: an independent seed specification caught shared state and fixture defects that mutual equality would miss; treatment provenance made the loading-policy sensitivity auditable; capacity characterization changed the tail interpretation; and the common-SQL path narrowed, but did not decompose, the selected-treatment spread (Supplement Table S50). Applied to the nine external sources, the codebook yields claim licenses rather than a binary quality score (Supplement Table S37).

This demonstrates that the protocol is *executable*: it admits and rejects concrete cells, it changed the reported interpretation at every stage exercised here (R7 excepted, being unsatisfied), and it codes an external corpus to a reproducible per-stage result. It does not demonstrate *reusability by another person*. Two limits sharpen this: no source was coded *satisfied* on any stage, so the codebook’s upper range is untested, and agreement against this rater would measure agreement with the protocol’s own designer, where two blind raters would be the stronger design. The codebook, audit record, agreement scorer, and result-blind forms are released so the question can be settled against this study. The claim is therefore a *candidate* access-layer-specific discipline and reusable artifact, demonstrated once: not invention of the constituent controls, not a complete theory of benchmarking validity, and not a validated methodology. The demotion follows from the circularity stated in Section 3: until an analyst outside this study applies the protocol, it is a checklist worth testing rather than one already tested.

Practical guidance. The synthesis follows from RQ1–RQ3 and describes the benchmarked implementations in the tested configuration (Section 6), not access-layer categories in general. Because the fan-out sweep varies breadth (0–500 children) but not depth, schema, or query mix, the deep-fetch peak locates a difference among five probes on one synthetic schema: an observation to re-measure on the target workload, not a law. Where a hot path resembles the tested deep fetch the access layer is consequential and worth benchmarking:

the thin paths (native driver, Knex) led here *for throughput and response-time p99*. Where the tested workload is read-simple, or for the tested MySQL single-row insert where shared commit behavior contributes materially under default durability, the observed layer spread is smaller and may be met by adding application processes until another bottleneck (contention, connection limits). These are performance findings only: the choice also turns on correctness, maintainability, security, type safety, and developer productivity, which this study does not measure and may outweigh them; and the common-SQL results are a compound sensitivity contrast on the selected-configuration cost, not a licence to use raw SQL.

6. Threats to Validity

We organize the threats along the four standard validity categories [62, 48].

Construct validity of the measures. Our dependent variables are HTTP-level throughput (req/s) and tail latency (p99) at the Express endpoint, not query-execution time at the driver boundary. This is deliberate (it captures the full request round trip) though overhead attributed to the *access layer* may partly reflect Express routing and serialization. We hold those near-constant: every adapter satisfies a documented *adapter contract* and funnels rows through shared canonical constructors (fixed field set, key order, typing, under TZ=UTC). These constructors copy and type-normalize already-grouped rows or entity graphs without querying, joining, regrouping, sorting, or caching. Their standalone median peaks at 17.6 μ s, 0.146% of the fastest reported HTTP p50, too small to explain the multi-fold gap though near-ties stay practically unresolved.

The specification-derived oracle passes 73,080 expected-result checks independently of all adapters (Section 3), and differential checks then establish finite mutual equivalence; these found, among other defects, a Drizzle timestamp conversion that shifted PostgreSQL instants by the host UTC offset, fixed before measurement. A fixed-payload endpoint measures the shared Express-and-serialization floor at 11,238 and 10,893 req/s

(`results/sameplan-corrected.json`), at least three times the fastest deep-fetch cell, so the framework compresses rather than manufactures between-layer differences, attributable to how each layer reaches an identical response, not to what is returned. The reported single-row-insert finding uses each engine’s *default* vendor-standard durability (Section 3), not one we relaxed, so those insert spreads (Supplement Table S52) describe the tested default deployment, not a tuning choice; a stack-specific relaxed-durability sensitivity quantifies how each configured treatment changes under that compound intervention without decomposing its cause (Section 5, Supplement Table S3).

Construct validity of the treatment. RQ1’s treatment is each layer’s *policy-selected documented* implementation and eager-loading strategy: an operational rule (the API the official documentation presents first) not validated against surveyed, typical, or performance-tuned usage. It is an intentionally artificial, predeclared selection policy, not a behavioral practitioner model. The frozen pages, conflict rule, and assignments make the decision auditable, but the single author both defined and applied it; two blind selection reviews are prepared and still pending. Where the policy differs from a performance-tuned configuration, RQ1’s spread mixes the library’s execution machinery with that strategy choice, made material by the one-to-four statement range on the deep fetch (Supplement Table S2).

Three sensitivity checks qualify it. First, every selected path and alternative is preserved with the source-page hash (Supplement Table S33 and `METHODOLOGY.md`); this is provenance, not evidence of frequency or idiomatity. Second, the *common-SQL raw-path sensitivity contrast* runs common SQL through every layer’s raw facility, so the raw-path spread (deep fetch $1.71\times$, against $7.01\times$ for the selected configuration; Section 4) is far smaller, a compound sensitivity contrast, not a decomposition or bound on how much comes from the loading strategy versus the library machinery. Third, where the documented alternative is a byte-identical drop-in, a loading-strategy check (Supplement Table S18) shows the selected path is the *faster* strategy for the

join-selected ORMs, so their deficit is not an artifact of a poor selected strategy; the lone exception, Objection’s slower selected select-in path on MySQL, is disclosed rather than corrected. RQ1 therefore answers what this policy yields, not a library-only overhead.

Internal validity. Each adapter follows the registered policy and passes specification-derived and differential checks, but all were implemented by one author. A semantically correct yet needlessly allocating, converting, or serializing adapter would pass; no independent human implementation review is claimed. A result-blind author self-audit found two avoidable timed-path operations before campaign acceptance: Knex passed its PostgreSQL-style `returning` option on MySQL (44,907 console warnings within two pilot repetitions), and MikroORM’s aggregation obtained and discarded an unused Knex handle per request. Both were removed without changing the declared treatment or SQL, the admission chain passed again, and measurement restarted from repetition 1 (`notes/adapter-self-audit.md`); because the implementer performed it, this is author review.

Both are now detected mechanically instead, which is why R7 requires evidence rather than a reviewer: a static check reports values acquired and never read, clean on all 13 modules of the measured source, and a runtime gate rejects any cell writing to `stderr` or `stdout` during a measured run. The runtime gate postdates the reported campaign and did not gate it, so absence of output there rests on author observation. Neither covers waste that is silent *and* invisible in the source (an unnecessarily expensive documented configuration, say) so the residual narrows without closing: a correct, policy-conformant adapter uniformly slower than a differently written one would pass every check. We report it as open, not discharged. The same caution applies more widely: four admission elements (the specification-derived oracle, seed-parity and campaign-state verification, and the dead-work check) postdate the accepted campaigns (Supplement Table S48). Because they replay a deterministic seed and inspect archived state, they can invalidate archived results, and did so for the super-

seeded pilot; but they did not gate the timed runs as those runs executed, so a defect observable only while a run was in progress would have escaped them.

Per-cell warm-up phases exceeding the slowest layer’s cold-start stabilization absorb JIT, pool-fill, and plan-cache effects [63]; two further controls address non-layer confounds: one correlated-subquery aggregation plan, and the corrected-state campaign’s verified reset before each write cell. The historical MySQL allocator defect and the corrected reset are described in Section 3; its consequence here is that the affected historical cells are not used as clean evidence, and all five patterns were remeasured in one campaign. The seeded identifier design, and why served sequences are not identical, are set out in Section 3; a forward-versus-reversed write-cell-order check found no detectable order effect, so order does not confound layer identity with run position. A residual risk is coordinated omission, which can under-report the tail at saturation [16, 55]; the corrected constant-arrival sweeps show small tails below measured capacity and sharp growth near it (Supplement Tables S6 and S17).

External validity. All measurements come from a single schema, dataset size, and 16-vCPU virtualized host, a pool fixed at ten, and specific engine and library versions (PostgreSQL 18.4, MySQL 9.7.1; latest compatible stable versions as of the 15 July 2026 freeze, Supplement Table S11). The fixed pool does not drive the ranking: the pool-size frontier (Supplement Tables S7, S25) preserves the ordering from 1 to 50 connections, and a mixed read/write workload (S26) preserves the read ordering. Two choices are disclosed rather than resolved: Sequelize on its stable 6.x line (7.x is prerelease), and Prisma’s MySQL path on the MariaDB adapter, as Prisma ships no `mysql2` adapter, an implementation×engine asymmetry. The memory-resident working set is a hot-cache regime that makes access-layer overhead more visible; as a workload becomes I/O-bound the spreads may compress toward 1×.

Same-host checks do not bound an independent substrate: a post-restart re-run of six representative cells moved absolute throughput by up to 16% against the campaign prevailing when it ran, which is the superseded pilot rather

than the accepted campaign, since the restart check preceded the latter. The corrected-state 25-run campaign repeats all five patterns with a new randomized order and environment fingerprint on the same host. Cross-campaign agreement or disagreement is reported directly; it reduces dependence on one overnight ordering but does not address host, hypervisor-neighbor, CPU-steal, or frequency transfer. No empirical ordering is intended to travel beyond the tested campaigns, and close cross-backend rank reversals are exploratory unless their direction repeats. Every promoted RQ2 reversal involves the single treatment whose MySQL driver differs, and a leave-one-out repetition of the promotion procedure without it promotes none (Supplement Table S49); the promoted-reversal evidence is therefore conditional on that driver substitution, which the design cannot separate from the treatment or from the DBMS. The protocol evidence is also bounded: its external audit has one rater, so general independent applicability remains to be tested.

Conclusion validity. Our analysis (Section 3) reports medians with the coefficient of variation and within-campaign bootstrap 95% intervals (repeatability, not population-general) and, because the design is a randomized block, compares adjacently ranked layers with *paired* tests on their per-replicate throughput ratios (paired sign-flip permutation, cross-checked with Wilcoxon; Supplement Table S36). Two bounds reinforce the ordering: the widest spreads dwarf the deep-fetch run-to-run CV ($\leq 4.5\%$); insert CV reaches 6.4% on PostgreSQL and 13.6% on MySQL (Supplement Tables S1, S9), which is why no fine insert ordering is claimed, and it holds across the concurrency sweep at ≥ 32 connections (Supplement Figure S2). Rather than assert a specific statistical power (absent a prespecified effect-size model), we rest the conclusions on effect sizes and interval estimates, foremost the *predefined* native-driver contrasts (Supplement Table S41), keeping the ranking-selected adjacent-pair permutation tests only as a descriptive post-hoc check. Application processes do not cross adapter, engine, or replicate blocks, although read endpoints within a block share one process [64]. All blocks share one host; primary intervals remain within-campaign,

while the corrected RQ2 subset has a second same-host campaign.

The secondary experiments covering only a layer subset or few replicates (Supplement Table S32) are read as *sensitivity evidence* on the tested subset, not as claims about all eleven implementations. The common-SQL contrast rests on ten replicates and the matched-utilization tails on five, so both are reported with their observed spread (Supplement Tables S14, S21–S22) and neither carries an interval estimate. The matched-utilization tails carry a second, unpropagated uncertainty: each offered rate is a fraction of an *estimated* capacity, whose spread is evaluated separately (Supplement Table S45) rather than carried into the reported p99. They are secondary evidence for the operating-point distinction, not a precise tail estimate.

7. Conclusion and Future Work

This paper proposes a concrete discipline for access-layer comparisons: derive expected read results from the declared task, check mutual equivalence and post-write state before a timing is admitted, define each configured treatment by a predeclared policy, and report fixed-demand latency separately from capacity and matched utilization. The predeclared selection policy takes documented configurations from frozen pages; it is reproducible evidence about those treatments, not observed practitioner behavior. The common-SQL raw-path experiment is a compound sensitivity contrast, not a library-intrinsic overhead estimate.

In the pinned Express.js experiment, the largest observed gap occurs on the tested deep fetch: $7.01\times$ on PostgreSQL and $5.12\times$ on MySQL, narrowing to $1.71\times$ and $2.03\times$ when the layers use common SQL through raw facilities. At stable 50% and most 70% capacity targets, p99 values form a much smaller band than at the high-load point. At the stated evidential bar only one treatment reverses durably across the two tested backend stacks, and it is the one whose MySQL driver differs; no other pair clears the same bar without it, while whole-ordering agreement stays imperfect on both stacks either way. The experiment

therefore records limited, treatment-specific non-transfer without identifying the DBMS as its source. These are treatment- and campaign-specific observations; handwritten native SQL and ORM relation APIs are deliberately asymmetric practical implementations.

The evidence record separates completed checks from unperformed validation: the oracle, preflight, second same-host campaign, checklist, and retrospective audit make the method inspectable, while the audit has one rater, every adapter is single-author, and both campaigns share one virtualized host. Result-blind packets are provided instead of a claim of independent validation.

Future work should estimate inter-rater agreement for the external audit, repeat the reduced RQ2 matrix on a physically independent host, and extend the workload across datasets, topologies, and cold-cache regimes. The first is prepared rather than proposed: a rater-extensible audit record and a committed agreement scorer ship with the artifact, so an independent rater can settle it without re-running a measurement. The released harness and explicit compliance levels make these extensions additive rather than changes to an implied universal ranking.

CRedit authorship contribution statement

Mateusz Miotk: Conceptualization, Methodology, Software, Investigation, Data curation, Formal analysis, Visualization, Writing – original draft, Writing – review & editing.

Declaration of competing interest

The author declares no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper. The author is not affiliated with, nor funded by, any of the database libraries or vendors evaluated in this study.

Data availability

The public repository is available at <https://github.com/mmioth/express-db-access-performance>. The complete revision artifact is archived as release v1.13.1, which the concept DOI <https://doi.org/10.5281/zenodo.21313858> resolves to [65]; the Zenodo record also carries an immutable per-version DOI for that release. A reproducibility summary — versions, requirements, run commands, resources needed, and which results the harness regenerates automatically from the archived raw data — is in Supplement Table S31 and REPRODUCE.md. **Exact reproduction requires the reference engines** (PostgreSQL 18.4 and MySQL 9.7.1) via the user-space conda path documented there; the convenience Docker path pins the same engine versions but uses relaxed durability in a containerized topology, so it reproduces the workflow, not the headline default-durability insert numbers. A machine-readable encoding of the protocol ships as `protocol-checklist.yaml`. Reproduction status is separated explicitly: author-run isolated reconstructions verified the archived JSON hashes and regenerated the committed tables byte-for-byte without starting either DBMS (dates and counts in Supplement Table S31). Neither is independent reproduction nor a re-execution of the current measurements, and neither is described as such.

Funding

This research did not receive any specific grant from funding agencies in the public, commercial, or not-for-profit sectors.

Declaration of generative AI and AI-assisted technologies in the manuscript preparation process

During the preparation of this work the author used Claude (Anthropic) to draft and edit manuscript text and improve its readability and language, and OpenAI Codex once to obtain an independent critical review of the complete manuscript. After using this tool, the author reviewed and edited the

content as needed and takes full responsibility for the content of the published article. AI assistance that formed part of the research itself is separate from the preparation use declared here and is documented in Section 3.9: it covered implementation of the benchmark harness and analysis code, and extended to proposing candidate analysis procedures, which the author selected among, verified, and interpreted. That section states the models, the affected components, the checks that independently validate the generated code, and the limits of those checks.

References

- [1] T.-H. Chen, W. Shang, Z. M. Jiang, A. E. Hassan, M. Nasser, P. Flora, Detecting performance anti-patterns for applications developed using object-relational mapping, in: Proceedings of the 36th International Conference on Software Engineering (ICSE), ACM, 2014, pp. 1001–1012. doi:10.1145/2568225.2568259.
- [2] M. Fowler, Patterns of Enterprise Application Architecture, Addison-Wesley, Boston, MA, 2002.
- [3] I. K. Chaniotis, K.-I. D. Kyriakou, N. D. Tselikas, Is Node.js a viable option for building modern web applications? a performance evaluation study, Computing 97 (10) (2015) 1023–1044. doi:10.1007/s00607-014-0394-9.
- [4] P. Kuffel, B. Walter, Changes in Node.js server-side application performance characteristics over time under load, in: Advances in Information Systems Development, Vol. 77 of Lecture Notes in Information Systems and Organisation, Springer, 2025, pp. 275–294. doi:10.1007/978-3-031-87880-0_14.
- [5] K. X. Carmo, F. Ferreira, E. Figueiredo, Performance evaluation of back-end frameworks: A comparative study, in: Proceedings of the 20th Brazilian Symposium on Information Systems (SBSI), ACM, 2024, pp. 1–9. doi:10.1145/3658271.3658314.

- [6] D. Colley, C. Stanier, M. Asaduzzaman, The impact of object-relational mapping frameworks on relational query performance, in: 2018 International Conference on Computing, Electronics & Communications Engineering (iCCECE), IEEE, 2018, pp. 47–52, java, C#, Python, PHP — deliberately excludes JavaScript/Node. doi:10.1109/iccecome.2018.8659222.
- [7] J. C. Yusmita, R. Arya, J. M. Wijaya, K. M. Suryaningrum, R. R. Siswanto, Optimizing database access strategy: A performance analysis comparison of raw sql and prisma orm, *Procedia Comput. Sci.* 269 (2025) 1201–1210. doi:10.1016/j.procs.2025.09.061.
- [8] J. Attala, C. Khemapatpapan, Comparing the performance of ORMs, *J. Comput. Creat. Technol.* 3 (2) (2025) 201–213. doi:10.14456/jcct.2025.16. URL <https://so13.tci-thaijo.org/index.php/jcct/article/view/2330>
- [9] I. P. A. E. Pratama, I. M. S. Raharja, Node.js performance benchmarking and analysis at VirtualBox, Docker, and Podman environment using Node-Bench method, *Int. J. Inform. Vis.* 7 (4) (2023) 2240–2246. doi:10.62527/joiv.7.4.1762.
- [10] S. Zhadko-Bazilevych, Analysis of ORM framework approaches for Node.js, *J. Comput. Sci. Inst.* 37 (2025) 426–430. doi:10.35784/jcsi.7951.
- [11] S. V. Salunke, A. Ouda, A performance benchmark for the PostgreSQL and MySQL databases, *Future Internet* 16 (10) (2024) 382. doi:10.3390/fi16100382.
- [12] Drizzle Team, Drizzle orm benchmarks, <https://orm.drizzle.team/benchmarks>, k6, 1M prepared requests, two-machine setup; reports req/s, avg + P95 latency, CPU (accessed 10 July 2026). (2024).
- [13] M. D. Imam Sakti, D. Dirgantara, A. K. Ramadhan, M. A. Ibrahim, Optimizing asynchronous performance in Node.js with Ex-

- press and PostgreSQL, *Procedia Comput. Sci.* 269 (2025) 172–181. doi:10.1016/j.procs.2025.08.270.
- [14] M. Collina, et al., autocannon: fast http/1.1 benchmarking tool for node.js, <https://github.com/mcollina/autocannon>, version 7.15.0; accessed 10 July 2026. (2023).
 - [15] M. Raasveldt, P. Holanda, T. Gubner, H. Mühleisen, Fair benchmarking considered difficult: Common pitfalls in database performance testing, in: *Proceedings of the Workshop on Testing Database Systems (DBTest)*, ACM, 2018, pp. 1–6. doi:10.1145/3209950.3209955.
 - [16] M. Fruth, S. Scherzinger, W. Maurer, R. Ramsauer, Tell-tale tail latencies: Pitfalls and perils in database benchmarking, in: *Performance Evaluation and Benchmarking (TPCTC 2021)*, Vol. 13169 of *Lecture Notes in Computer Science*, Springer, 2022, pp. 119–134. doi:10.1007/978-3-030-94437-7_8.
 - [17] A. Georges, D. Buytaert, L. Eeckhout, Statistically rigorous Java performance evaluation, in: *Proceedings of the 22nd ACM SIGPLAN Conference on Object-Oriented Programming Systems, Languages and Applications (OOPSLA)*, ACM, 2007, pp. 57–76. doi:10.1145/1297027.1297033.
 - [18] T. Sotiropoulos, S. Chaliasos, V. Atlidakis, D. Mitropoulos, D. Spinellis, Data-oriented differential testing of object-relational mapping systems, in: *Proceedings of the 43rd IEEE/ACM International Conference on Software Engineering (ICSE)*, IEEE, 2021, pp. 1535–1547. doi:10.1109/ICSE43902.2021.00137.
 - [19] C. Ireland, D. Bowers, M. Newton, K. Waugh, A classification of object-relational impedance mismatch, in: *2009 First International Conference on Advances in Databases, Knowledge, and Data Applications (DBKDA)*, IEEE, 2009, pp. 36–43. doi:10.1109/DBKDA.2009.11.

- [20] A. Torres, R. Galante, M. S. Pimenta, A. J. B. Martins, Twenty years of object-relational mapping: A survey on patterns, solutions, and their implications on application design, *Inf. Softw. Technol.* 82 (2017) 1–18. doi:10.1016/j.infsof.2016.09.009.
- [21] P. J. Sadalage, M. Fowler, *NoSQL Distilled: A Brief Guide to the Emerging World of Polyglot Persistence*, Addison-Wesley, Upper Saddle River, NJ, 2012.
- [22] J. Yang, P. Subramaniam, S. Lu, C. Yan, A. Cheung, How not to structure your database-backed web applications: A study of performance bugs in the wild, in: *Proceedings of the 40th International Conference on Software Engineering (ICSE)*, ACM, 2018, pp. 800–810. doi:10.1145/3180155.3180194.
- [23] C. Yan, A. Cheung, J. Yang, S. Lu, Understanding database performance inefficiencies in real-world web applications, in: *Proceedings of the 2017 ACM Conference on Information and Knowledge Management (CIKM)*, ACM, 2017, pp. 1299–1308. doi:10.1145/3132847.3132954.
- [24] S. Shao, Z. Qiu, X. Yu, W. Yang, G. Jin, T. Xie, X. Wu, Database-access performance antipatterns in database-backed web applications, in: *2020 IEEE International Conference on Software Maintenance and Evolution (ICSME)*, IEEE, 2020, pp. 58–69. doi:10.1109/ICSME46990.2020.00016.
- [25] T.-H. Chen, W. Shang, Z. M. Jiang, A. E. Hassan, M. Nasser, P. Flora, Finding and evaluating the performance impact of redundant data access for applications that are developed using object-relational mapping frameworks, *IEEE Trans. Softw. Eng.* 42 (12) (2016) 1148–1161. doi:10.1109/TSE.2016.2553039.
- [26] A. Cheung, A. Solar-Lezama, S. Madden, Optimizing database-backed applications with query synthesis, in: *Proceedings of the 34th ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI)*, ACM, 2013, pp. 3–14. doi:10.1145/2491956.2462180.

- [27] A. Cheung, S. Madden, A. Solar-Lezama, Sloth: Being lazy is a virtue (when issuing database queries), in: Proceedings of the 2014 ACM SIGMOD International Conference on Management of Data, ACM, 2014, pp. 931–942. doi:10.1145/2588555.2593672.
- [28] T.-H. Chen, W. Shang, A. E. Hassan, M. Nasser, P. Flora, CacheOptimizer: Helping developers configure caching frameworks for Hibernate-based database-centric web applications, in: Proceedings of the 2016 24th ACM SIGSOFT International Symposium on Foundations of Software Engineering (FSE), ACM, 2016, pp. 666–677. doi:10.1145/2950290.2950303.
- [29] A. Turcotte, M. D. Shah, M. W. Aldrich, F. Tip, DrAsync: Identifying and visualizing anti-patterns in asynchronous JavaScript, in: Proceedings of the 44th International Conference on Software Engineering (ICSE), ACM, 2022, pp. 774–785. doi:10.1145/3510003.3510097.
- [30] M. Lorenz, J.-P. Rudolph, G. Hesse, M. Uflacker, H. Plattner, Object-relational mapping revisited—a quantitative study on the impact of database technology on O/R mapping strategies, in: Proceedings of the 50th Hawaii International Conference on System Sciences (HICSS), 2017. doi:10.24251/HICSS.2017.592.
- [31] P. van Zyl, D. G. Kourie, L. Coetzee, A. Boake, The influence of optimisations on the performance of an object relational mapping tool, in: Proceedings of the 2009 Annual Research Conference of the South African Institute of Computer Scientists and Information Technologists (SAICSIT), ACM, 2009, pp. 150–159. doi:10.1145/1632149.1632169.
- [32] A. Bonvoisin, C. Quinton, R. Rouvoy, Understanding the performance-energy tradeoffs of object-relational mapping frameworks, in: 2024 IEEE International Conference on Software Analysis, Evolution and Reengineering (SANER), IEEE, 2024, pp. 626–636. doi:10.1109/SANER60148.2024.00069.
- [33] Y. Li, S. Manoharan, A performance comparison of SQL and NoSQL databases, in: 2013 IEEE Pacific Rim Conference on Communications,

- Computers and Signal Processing (PACRIM), IEEE, 2013, pp. 15–19. doi:10.1109/PACRIM.2013.6625441.
- [34] B. F. Cooper, A. Silberstein, E. Tam, R. Ramakrishnan, R. Sears, Benchmarking cloud serving systems with YCSB, in: Proceedings of the 1st ACM Symposium on Cloud Computing (SoCC), ACM, 2010, pp. 143–154. doi:10.1145/1807128.1807152.
 - [35] J. Dean, L. A. Barroso, The tail at scale, *Commun. ACM* 56 (2) (2013) 74–80. doi:10.1145/2408776.2408794.
 - [36] J. Li, N. K. Sharma, D. R. K. Ports, S. D. Gribble, Tales of the tail: Hardware, OS, and application-level sources of tail latency, in: Proceedings of the ACM Symposium on Cloud Computing (SoCC), ACM, 2014, pp. 1–14. doi:10.1145/2670979.2670988.
 - [37] Z. M. Jiang, A. E. Hassan, A survey on load testing of large-scale software systems, *IEEE Trans. Softw. Eng.* 41 (11) (2015) 1091–1118. doi:10.1109/TSE.2015.2445340.
 - [38] P. Leitner, C.-P. Bezemer, An exploratory study of the state of practice of performance testing in Java-based open source projects, in: Proceedings of the 8th ACM/SPEC International Conference on Performance Engineering (ICPE), ACM, 2017, pp. 373–384. doi:10.1145/3030207.3030213.
 - [39] C. Collberg, T. A. Proebsting, Repeatability in computer systems research, *Commun. ACM* 59 (3) (2016) 62–69. doi:10.1145/2812803.
 - [40] K. Huppler, The art of building a good benchmark, in: Performance Evaluation and Benchmarking (TPCTC 2009), Vol. 5895 of Lecture Notes in Computer Science, Springer, 2009, pp. 18–30. doi:10.1007/978-3-642-10424-4_3.
 - [41] S. E. Sim, S. Easterbrook, R. C. Holt, Using benchmarking to advance research: A challenge to software engineering, in: Proceedings of the 25th

- International Conference on Software Engineering (ICSE), IEEE, 2003, pp. 74–83. doi:10.1109/ICSE.2003.1201189.
- [42] Prisma, Query performance benchmarks, <https://benchmarks.prisma.io/>, prisma/Drizzle/TypeORM; median of 500 iterations; no throughput, no p95/p99 (accessed 10 July 2026). (2024).
 - [43] Gel Data (EdgeDB), Imdbench: Orm benchmarks, <https://github.com/geldata/imdbench>, accessed 10 July 2026. (2024).
 - [44] S. Kounev, K.-D. Lange, J. von Kistowski, Systems Benchmarking: For Scientists and Engineers, Springer, 2020. doi:10.1007/978-3-030-41705-5.
 - [45] W. Hasselbring, Benchmarking as empirical standard in software engineering research, in: Proceedings of the Evaluation and Assessment in Software Engineering (EASE), ACM, 2021, pp. 457–462. doi:10.1145/3463274.3463361.
 - [46] M. Rigger, Z. Su, Testing database engines via pivoted query synthesis, in: Proceedings of the 14th USENIX Symposium on Operating Systems Design and Implementation (OSDI), USENIX Association, 2020, pp. 667–682.
 - [47] V. R. Basili, H. D. Rombach, The TAME project: Towards improvement-oriented software environments, IEEE Trans. Softw. Eng. 14 (6) (1988) 758–773. doi:10.1109/32.6156.
 - [48] C. Wohlin, P. Runeson, M. Höst, M. C. Ohlsson, B. Regnell, A. Wesslén, Experimentation in Software Engineering, 2nd Edition, Springer, Berlin, Heidelberg, 2012. doi:10.1007/978-3-642-29044-2.
 - [49] B. A. Kitchenham, S. L. Pfleeger, L. M. Pickard, P. W. Jones, D. C. Hoaglin, K. El Emam, J. Rosenberg, Preliminary guidelines for empirical research in software engineering, IEEE Trans. Softw. Eng. 28 (8) (2002) 721–734. doi:10.1109/TSE.2002.1027796.

- [50] S. Harizopoulos, D. J. Abadi, S. Madden, M. Stonebraker, OLTP through the looking glass, and what we found there, in: Proceedings of the 2008 ACM SIGMOD International Conference on Management of Data, ACM, 2008, pp. 981–992. doi:10.1145/1376616.1376713.
- [51] K. Gupta, M. Mathuria, Improving performance of web application approaches using connection pooling, in: 2017 International Conference of Electronics, Communication and Aerospace Technology (ICECA), IEEE, 2017, pp. 355–358. doi:10.1109/ICECA.2017.8212833.
- [52] R. Laigner, Y. Zhou, M. A. Vaz Salles, Y. Liu, M. Kalinowski, Data management in microservices: State of the practice, challenges, and research directions, Proc. VLDB Endow. 14 (13) (2021) 3348–3361. doi:10.14778/3484224.3484232.
- [53] B. Schroeder, A. Wierman, M. Harchol-Balter, Open versus closed: A cautionary tale, in: Proceedings of the 3rd USENIX Symposium on Networked Systems Design and Implementation (NSDI), USENIX Association, 2006. URL <https://www.usenix.org/legacy/event/nsdi06/tech/schroeder.html>
- [54] X. Yu, G. Bezerra, A. Pavlo, S. Devadas, M. Stonebraker, Staring into the abyss: An evaluation of concurrency control with one thousand cores, Proc. VLDB Endow. 8 (3) (2014) 209–220. doi:10.14778/2735508.2735511.
- [55] G. Tene, How NOT to measure latency, <https://www.infoq.com/presentations/latency-response-time/>, conference presentation; origin of the term “coordinated omission” (accessed 10 July 2026). (2015).
- [56] B. Krishnamachari, How to do statistical evaluations in ECE/CS papers: A practical playbook for defensible results, <https://arxiv.org/abs/2605.00428> (2026).
- [57] C. Laaber, J. Scheuner, P. Leitner, Software microbenchmarking in the

- cloud. how bad is it really?, *Empir. Softw. Eng.* 24 (4) (2019) 2469–2508. doi:10.1007/s10664-019-09681-1.
- [58] A. Arcuri, L. Briand, A hitchhiker’s guide to statistical tests for assessing randomized algorithms in software engineering, *Softw. Test. Verif. Reliab.* 24 (3) (2014) 219–250. doi:10.1002/stvr.1486.
- [59] N. Cliff, Dominance statistics: Ordinal analyses to answer ordinal questions, *Psychol. Bull.* 114 (3) (1993) 494–509. doi:10.1037/0033-2909.114.3.494.
- [60] A. V. Papadopoulos, L. Versluis, A. Bauer, N. Herbst, J. von Kistowski, A. Ali-Eldin, C. L. Abad, J. N. Amaral, P. Tûma, A. Io-sup, Methodological principles for reproducible performance evaluation in cloud computing, *IEEE Trans. Softw. Eng.* 47 (8) (2021) 1528–1543. doi:10.1109/TSE.2019.2927908.
- [61] V. Leis, A. Gubichev, A. Mirchev, P. Boncz, A. Kemper, T. Neumann, How good are query optimizers, really?, *Proc. VLDB Endow.* 9 (3) (2015) 204–215. doi:10.14778/2850583.2850594.
- [62] T. D. Cook, D. T. Campbell, *Quasi-Experimentation: Design and Analysis Issues for Field Settings*, Houghton Mifflin, Boston, MA, 1979.
- [63] T. Mytkowicz, A. Diwan, M. Hauswirth, P. F. Sweeney, Producing wrong data without doing anything obviously wrong!, in: *Proceedings of the 14th International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*, ACM, 2009, pp. 265–276. doi:10.1145/1508244.1508275.
- [64] T. Kalibera, R. Jones, Rigorous benchmarking in reasonable time, in: *Proceedings of the 2013 International Symposium on Memory Management (ISMM)*, ACM, 2013, pp. 63–74. doi:10.1145/2464157.2464160.
- [65] M. Miotk, Replication package for A Comparability Protocol for Benchmarking Relational Database Access Layers in Express.js, Zenodo, version 1.13.0, [software and dataset] (2026). doi:10.5281/zenodo.21313858.