

Response to Reviewer (Revision Round 4)

We thank the reviewer for a rigorous and constructive report. The review's central observation — that several interpretations exceeded what a single-host, saturated, single-implementation design identifies, and that this concern has recurred — is correct, and we have addressed it structurally rather than cosmetically. The governing change this round is a **reposition**: the paper is now presented as a *reproducible benchmark instrument and a configuration-specific comparison*, not as a generalized ranking of database libraries. Single host, idiomatic-choice, and the saturated operating point are now **declared scope**, and the harness, its controls, and a benchmarking-pitfalls checklist are the primary contributions. We also ran the feasible experiments the review pointed toward. Section and table references are to the revised submission build; effective length is recorded in `word-count.pdf`.

The reposition (addresses the recurring "claims exceed design")

- **Title, Objective, RQs, contributions recast.** The title now reads *"A*

Reproducible Benchmark . . . A Configuration-Specific Comparison on PostgreSQL and MySQL." *The Objective leads with the harness ("the harness, not a generalized ranking, is the contribution"). Each RQ is scoped to "the benchmarked implementations under the specified operating point." The contributions list the instrument, the design, and the checklist as primary**; the empirical findings are labelled configuration-specific.

- **We propagated the fixes uniformly.** The recurrence was partly mechanical: prior rounds fixed a phrase in one place but not others, so the paper contradicted itself. We swept every occurrence. There are now **zero** instances of "independent replicate(s)" (the runs are "repeated randomized-block runs"; the paper already conceded non-independence, so this removes a self-contradiction), the primary p99 is called "saturated client-observed response-time p99" throughout, and the same-SQL control is described only as a compound intervention.

Major concerns

§6.1 — saturated p99 is a saturation outcome, not intrinsic tail. Agreed. It is now consistently named a saturated client-observed response-time p99 at the stated operating point, and we added a **utilization-controlled open-loop experiment**: each layer is offered 50/70/85/95 % of its own saturating throughput, replicated, on both engines. At 50 % of each layer's own capacity every layer holds p99 near 2–4 ms on both engines, and the large saturated gap (native driver 20 ms versus the slowest ORM 106 ms) dissolves — confirming the saturated ladder was tracking capacity/queueing, not intrinsic tail latency; the tail rises only as a layer nears its own capacity. The utilization results are now at least as prominent as the saturated p99.

§6.2 — same-SQL does not identify a mechanism. Agreed and corrected everywhere. The control is now described strictly as a compound intervention that changes query formulation, API, protocol, statement preparation, round-trip structure, and mapping *together* and isolates no single factor; the residual is not called "raw-execution cost" unqualified. A new **controlled/uncontrolled factor table** makes explicit what the control holds constant and what it changes together.

§6.3 — "idiomatic" is subjective. We added a **pre-specified selection protocol**, fixed before measurement: the eager-loading API each library's official documentation presents first for the pinned version, with logging/hooks/validation off. The exact API, **documentation page, and pinned version** for each layer are tabulated (`METHODOLOGY.md`, Supplement Table S16). We did not invite maintainers to review the adapters; that is now disclosed as an explicit

limitation. The alternative loading strategy is measured where it is a byte-identical drop-in; the native, query-builder, and Prisma layers have no second loading strategy to offer (their SQL is fixed), so "strongest alternative for every layer" is satisfied where feasible.

§6.4 — equal pool $\neq$ equal resource. We renamed the estimand "equal configured connection limit" and added a **per-layer pool-size frontier** on both engines (throughput versus pool size, 1–50) so the fixed-pool choice is shown not to drive the ranking.

§6.5 — single-host interference. We do not have a second machine, so a separate-host / remote-database topology is named as a scoped future use of the instrument rather than claimed. The feasible test of the reviewer's specific worry — that Prisma's multi-core engine looks competitive only because the one-core layers leave cores idle — is a **multi-worker (node cluster) experiment**: each layer is scaled to 1/2/4 workers so every layer can use the spare cores. The result confirms the reviewer's intuition: at one worker the native driver and Prisma are near parity, but at four workers the native driver overtakes Prisma by 1.6 $\times$ on PostgreSQL and 1.4 $\times$ on MySQL, because the single-pool layers scale on the cores Prisma's engine already consumed. Prisma's throughput parity is therefore specific to a deployment that leaves cores idle. The single-host topology, and its consequences for the CPU comparison, are disclosed prominently.

§6.6 — "25 independent replicates." Corrected to "25 repeated runs" throughout — the body, the main-text and supplement table captions, the generic build, and the table-generating scripts (so a re-run of the pipeline is also clean). The manuscript's own non-independence paragraphs are retained.

§6.7 — narrow workload / broad recommendations. Recommendations are now scoped to "the benchmarked implementations in this configuration, workload, and host," and the unmeasured productivity/type-safety trade-off is flagged as unmeasured. We added a **mixed read/write workload** (interleaved deep-fetch reads and inserts at documented mixes) and a **60–120 s longer-run sensitivity** subset. Both preserve the read ordering, and the 90 s runs reproduce the 12 s medians within 1 %.

§6.8 — selective/post-hoc statistics. Every primary layer $\times$ engine $\times$ pattern cell now carries a **95 % bootstrap confidence interval on throughput** in the tables (not only medians and selected adjacent tests). The single-run open-loop and three-run fan-out are **replicated** (open-loop to five, fan-out to eight runs) and the fan-out, durability, loading-strategy, and open-loop analyses are explicitly labelled exploratory. The adjacent-rank tests were already labelled post-hoc.

§6.9 — the interaction p is uninformative. We now read the interaction through its **effect size** (per-layer cross-engine throughput ratio, narrow 1.0–1.6 on reads, wide 2.4–7.7 on writes) and note that the floor p conveys little about magnitude; a **per-layer PG-vs-MySQL rank table** now backs the correlations (Prisma flips from rank 1 on PostgreSQL to rank 6 on MySQL for inserts).

§6.10 — novelty. We call the literature review a *structured related-work search* (not systematic), scope the novelty to the searched databases and dates, and archive the search strings.

What was already correct (verified, not changed)

We note, in case it saves the reviewer time, that several items the review asked us to check are already implemented as recommended:

- **p99 test unit** is the run-level statistic (25 per cell), never the individual

request; the paper and code (`stats2.mjs`) both use run-level values.

- **Bootstrap** is a **percentile** bootstrap resampling **replicate indices**

(preserving pairing), from a **fixed seed** (`mulberry32(0x57a75)`); this is now stated in the methodology.

- **Docker images** are pinned by immutable `@sha256` digest; **dependencies** by a committed lockfile.

- All six **declarations** (structured abstract, CRediT, competing interest, funding, data availability with the Zenodo DOI, generative-AI) are present.

- The **author name** is consistent ("Miotk") throughout; there is no "Miotka".

Point-by-point answers to the reviewer's questions

1. **Idiomatic selection rule?** Pre-specified before measurement: the first eager-loading API in each library's official guide for the pinned version; documented with page and version (Supplement Table S16, `METHODOLOGY.md`).

1. **Maintainer/expert review?** No; disclosed as a limitation.
2. **Equivalent tracking/hooks/validation/lifecycle?** Library defaults, disabled or request-scoped on the read path; documented per adapter.

1. **Same-SQL mappers identical machine code?** Source-level equivalent, not identical code paths; this is exactly why the control is a compound intervention.

1. **Prepared-statement caches over fresh processes?** A fresh process per run means a cold statement cache each run; the tuned baselines reuse within a run.

1. **Bootstrap resampling unit?** Replicate index (25 run-level values), seeded.
2. **p99 tests on 25 run-level values?** Yes.
3. **Why saturated p99 primary vs replicated open-loop at controlled utilization?**

Both are now reported; the utilization-controlled open-loop is the near-intrinsic latency evidence.

1. **Processes CPU-pinned?** DB/server/generator share the host; the equal-CPU control pins the server; single-host is disclosed.

1. **Steal time / frequency / neighbours?** Uncontrolled on the virtualized host; disclosed; the drift, post-restart, and longer-run checks bound their effect.

1. **Load generator co-located every run?** Yes (single host); disclosed.
2. **Buffer pools / OS caches between cells?** Warm by design; randomized order and fresh boot; the post-restart check confirms the cold-start ranking.

1. **Exact IST word count?** Recomputed (`word-count.pdf`).
2. **Zenodo archive = submitted version?** Yes; the v1.4.0 tarball is 'git archive HEAD' of the tagged release.

1. **Fan-out "robust" with three runs?** The reviewer was right to doubt it. On

replication (all layers, eight runs, both engines) the spread is a **roughly constant** multiplier across 0–500 comments, not the growing one the three-run sweep suggested; we removed the "growing with graph size" claim from the abstract, intro, results, discussion, conclusion, and threats, and report the fan-out as exploratory.

1. **Uncertainty for the MySQL commit-path percentages?** Reported from the wait-event replicates.

1. **Stable public versions + immutable ids?** Digest-pinned images + lockfile.

2. **Miotk or Miotka?** Miotk everywhere (verified).

3. **AI generate/modify adapters or stats code?** Assisted; verified by the byte-identical cross-check and unit tests; disclosed.

1. **Raw logs reconstruct failed/excluded runs (knex n=24)?** Yes; raw JSON, partial files, and status logs are archived.

We believe the repositioned paper claims exactly what a single-host benchmark identifies, backed by new experiments and a fully documented, provenance-tracked instrument, and we thank the reviewer for pressing it to this point.