

Summary of Revisions

Manuscript: A Comparability Protocol for Benchmarking Relational Database Access Layers in Express.js **Journal:** Information and Software Technology **Revision date:** 26 July 2026

Semantic admission

- Added an expected-result oracle derived from the deterministic seed and fan-out specification, independent of all adapters and both DBMSs.
- Kept differential comparison as a separate equivalence layer and renamed the protocol stage accordingly.
- Added exact campaign-state and allocator preflight checks.
- Added a 662-row cross-engine seed-parity gate; it caught differing fan-out comment-author assignments hidden by equal counts, and the affected pilot was rejected before the campaign restarted from repetition 1.
- Corrected a MySQL reset defect discovered by the new oracle and remeasured all affected patterns.

Treatment and implementation validity

- Renamed the case-study estimand “policy-selected documented configuration” and removed the behavioral developer persona.
- Made the selection policy substitutable at the protocol level.
- Preserved exact documentation responses, timestamps, hashes, conflict resolution, and alternatives.
- Added implementation-review provenance as recommended stage R7 and result-blind review packets.
- Marked independent human selection and adapter review as pending rather than claiming agreement.
- Recorded a result-blind author self-audit covering every adapter. It caught and removed per-insert Knex/MySQL warning I/O and an unused per-request MikroORM aggregation accessor; every affected pilot was preserved and rejected, the full admission chain was re-run, and measurement restarted from repetition 1. This is not counted as independent R7 evidence.

RQ2 and new measurements

- Ran one coherent corrected-state primary campaign covering all 18 compatible treatment-backend pairs for all five access patterns: 90 complete cells x 25 runs, zero timed request failures, and pre/post state checks.
- Ran a second independently ordered same-host campaign on 7 portable layers x 2 stacks x 3 prioritized patterns: 42 complete cells x 25 runs, zero timed failures. Deep-fetch ranks reproduce exactly; MySQL insert reproduces only 3/7 ranks and is reported as unstable.
- Predeclared persistence criteria retain four material cross-stack reversals and classify all close/nonrecurring changes as exploratory.
- Defined RQ2 as backend-stack transfer within the tested same-host campaigns.
- Removed causal DBMS attribution and made close rank reversals exploratory.
- Retained the explicit limitation that the second campaign is same-host, not cross-host replication.

Protocol contribution

- Narrowed novelty to a concrete domain-specific operational discipline built from established controls.
- Cited and distinguished CYNTHIA and retained the direct prior-art boundary.
- Added the 2023 Node-Bench comparison as a direct coverage precedent and a ninth source in the protocol audit.
- Added a seven-stage, nine-source retrospective protocol audit with 63 source-located judgments and a machine validator.
- Described that audit as a one-rater bounded applicability demonstration, not independent validation.
- Removed universal completeness and necessity language; each level now states the interpretation it licenses.

Fresh secondary analyses

- Reran common-SQL (16 cells x 10), all-pattern capacity (720 points), matched-utilization (64 cells x 5), longer-tail (18 cells x 10 x 60 s), relaxed durability (18 cells x 5), wait-event, and canonical-constructor analyses with zero accepted request failures.
- Updated all numerical claims: selected deep-fetch spread 7.01x/5.12x; common-SQL 1.71x/2.03x; stable 50% p99 bands 2.3–4.9/1.9–5.5 ms.
- Reframed concurrent wait sums and durability changes as compound sensitivity evidence, not a DBMS floor or causal share.

Reproducibility

- Reconciled the main paper with Supplement S31.
- Completed an isolated current-candidate reconstruction: 60/60 JSON hashes verified, and 35 table outputs plus four derived analysis JSON files regenerated byte-for-byte without starting either DBMS.
- Preserved and verified exact 42-file source archives for each accepted campaign, so later harness maintenance cannot overwrite the code provenance of completed measurements.
- Fixed the standalone write-admission cleanup to restore post and comment allocators even after failure; the corrected verifier passes all 18 adapter-engine admissions and leaves both campaign states exact.
- Named the five presentation-only differences in the historical 45/50 reconstruction.
- Separated author-run offline reconstruction, historical archive-isolated reconstruction, current measurement re-execution, and independent reproduction.
- Marked the Docker path as workflow reproduction because its durability and topology differ from the reference numerical path.
- Updated the manifest so every revised output maps to its exact data and generator.

Presentation

- Replaced “strategy standardization” with “common-SQL raw-path sensitivity contrast.”
- Restricted matched-utilization conclusions to stable 50% and most 70% cells.
- Regenerated the MySQL insert distribution from corrected state.
- Standardized “deep fetch” and “single-row insert.”
- Consolidated repeated qualifications and moved exploratory interpretation out of the main argument.