

Online supplement:
A Comparability Protocol for Benchmarking
Relational Database
Access Layers in Express.js

Mateusz Miotk

Faculty of Mathematics, Physics and Informatics, University of Gdańsk
`mateusz.miotk@ug.edu.pl`

This supplement carries measurement detail referenced from the main text. All tables are generated by the replication package (<https://github.com/mmiothk/express-db-access-performance>; archived at Zenodo as release v1.12.8, <https://doi.org/10.5281/zenodo.21497051>) from the same raw data as the main-text tables.

1 Dispersion on MySQL

The main text reports the coefficient-of-variation table for PostgreSQL and the per-engine maxima. Table S1 gives the full per-layer, per-pattern CV on MySQL across the 25 repeated runs (maximum 15.9%, `drizzle`/insert; the insert cells are the noisiest in the study). Figure S1 plots the raw per-replicate insert throughput behind these coefficients: several cells are visibly bimodal or heavy-tailed within a cell, structure a single CV cannot convey, which is why the insert rank correlation and dispersion are reported conservatively.

2 Documentation-selected round-trip counts

Table S2 reports, for each layer, the number of SQL statements the documentation-selected deep fetch issues per request, captured from server-side statement logs. The counts motivate the main text’s observation that round-trip count alone does not determine throughput: the two layers issuing four statements (Prisma, Objection) sit at opposite ends of the throughput ladder.

Table S1: Coefficient of variation (%) of throughput across the 25 repeated runs, per access layer and pattern on MySQL. Low values indicate stable measurements.

Layer	point_read	range_scan	deep_fetch	aggregation	write
mysql2	2.4	2.6	2.4	2.5	10.5
knex	3.7	2.7	2.7	1.9	13.2
drizzle	2.9	2.2	2.0	3.5	15.9
prisma	1.9	2.6	3.3	1.9	7.1
sequelize	2.1	1.8	2.5	1.6	12.5
typeorm	3.6	4.0	3.2	1.8	8.5
objection	2.5	1.9	3.8	2.9	11.7
mikroorm	3.0	4.0	2.2	2.1	5.6




Figure S1: Raw per-replicate insert throughput on MySQL for the eight documentation-selected layers (25 repeated runs each, red bar = median). The insert cells are the noisiest in the study; several are visibly bimodal or heavy-tailed within a cell, structure a coefficient of variation cannot convey, which is why the insert rank correlation and dispersion are reported conservatively.

Table S2: Number of SQL round trips each access layer issues per request on PostgreSQL, captured from server-side statement logging. The native drivers, Knex, and Drizzle issue a two-query deep fetch; the data-mapper ORMs’ eager-loading strategies choose their own counts. Full generated SQL and EXPLAIN plans are in the replication package.

Layer	Point read	Range scan	Deep fetch	Aggregation	Insert
pg	1	1	2	1	1
knex	1	1	2	1	1
drizzle	1	1	2	1	1
prisma	1	1	4	1	1
sequelize	1	1	1	1	1
typeorm	1	1	2	1	2
objection	1	1	4	1	1
mikroorm	1	1	1	1	1

3 Write throughput under default versus relaxed durability

Table S3 gives the per-layer insert throughput under each engine’s default durability configuration (the paper’s primary write regime) and under the symmetric relaxed regime (secondary, 10 replicates), with the relaxed÷default ratio quoted in the main text.

4 Equal-CPU-budget sensitivity check

Table S4 gives the deep-fetch throughput with the server confined to 1, 2, or 4 cores (database and load generator pinned to disjoint cores); all nine values are quoted in the main text.

5 Application-tier versus end-to-end CPU efficiency

Table S5 reports, per layer on the deep/nested fetch, the median application-process-tree CPU and database-server CPU, the combined core-count, and throughput per application-CPU-second and per combined (application+database) CPU-second. The application-only figure understates the compute of layers that delegate more work to the database, but on the current versions the application-only and combined figures largely agree, because no layer draws more than about one application core: `mysql2` leads application-tier efficiency by about $3.8\times$ over Prisma (2,269 vs. 604 req per app-CPU-s) and by $3.3\times$ end-to-end (1,254 vs. 384 req per combined-CPU-s). There is no CPU-for-throughput trade of the kind an earlier version of this study reported.

Table S3: Insert throughput (req/s) under the default durability configuration (PostgreSQL `fsync=on`, `synchronous_commit=on`; MySQL `innodb_flush_log_at_trx_commit=1`, `sync_binlog=1`; median of 25 replicates) and under the relaxed regime (all of the above off; median of 10), with the relaxed÷default ratio. At 50 concurrent connections, group commit amortizes the per-commit flushes: relaxing durability buys at most 14% on PostgreSQL and 1.1–2.4× on MySQL, largest for the flush-bound `mysql2` and `Knex` and smallest for the overhead-bound ORMs.

Layer	PostgreSQL			MySQL		
	default	relaxed	ratio	default	relaxed	ratio
<code>pg</code>	6402	7092	1.11	–	–	–
<code>mysql2</code>	–	–	–	1753	4130	2.36
<code>pg-tuned</code>	6413	7340	1.14	–	–	–
<code>mysql2-tuned</code>	–	–	–	1750	4039	2.31
<code>knex</code>	4841	5380	1.11	1670	3505	2.10
<code>drizzle</code>	4085	4554	1.11	1730	3280	1.90
<code>prisma</code>	2774	2741	0.99	886	988	1.12
<code>sequelize</code>	2732	2714	0.99	1478	2267	1.53
<code>typeorm</code>	2648	2764	1.04	1345	1647	1.22
<code>objection</code>	3907	4329	1.11	1728	3233	1.87
<code>mikroorm</code>	1979	2040	1.03	1257	1458	1.16

Table S4: Exploratory equal-CPU sensitivity check (deep-fetch throughput, req/s, PostgreSQL, median of 3 runs, three representative layers) with the server process confined by `taskset` to 1, 2, or 4 cores (database and load generator pinned to disjoint cores). `pg` holds full throughput on a single core and gains nothing from more; Prisma and MikroORM stay far below it and do not rise monotonically with extra cores (their per-core values wobble by 10–17% run-to-run at this replicate count), so the deep-fetch ranking does not appear to be an artifact of unequal core budgets — a layer’s low per-process throughput is recovered by more processes over the tested range (Table S24), not more cores.

Layer	1 core	2 cores	4 cores
<code>pg</code>	3689	3688	3687
<code>prisma</code>	1079	1052	1219
<code>mikroorm</code>	446	495	522

Table S5: Application-tier versus end-to-end CPU efficiency on the deep/nested fetch (MySQL): median CPU (% of one logical core) of the Node.js process tree and of the database server, the combined core-count, and throughput per application-CPU-second and per combined (app+database) CPU-second. The application-only figure understates the compute of layers that delegate more work to the database; the combined figure is the end-to-end view. Load-generator CPU is excluded.

Layer	app CPU%	db CPU%	comb. cores	req/app-CPU-s	req/comb-CPU-s
mysql2	105	85	1.90	2269	1254
knex	100	75	1.75	2007	1147
drizzle	105	50	1.55	1255	850
prisma	105	60	1.65	604	384
sequelize	100	30	1.30	886	682
typeorm	100	70	1.70	1017	598
objection	100	45	1.45	834	575
mikroorm	110	25	1.35	436	356

6 Coordinated-omission-corrected open-loop sweep

Table S6 gives the full constant-arrival sweep (PostgreSQL, five layers, offered rates 250–4,000 requests/s) behind the open-loop paragraph of the main text: achieved throughput, corrected latency percentiles measured from each request’s *intended* send time, and timeout counts (10 s cap, clipped into the distribution).

Table S6: Open-loop tail latency, corrected for coordinated omission: p99 (ms) on the deep/nested fetch (PostgreSQL) under a constant offered request rate, with every latency measured from the request’s intended start time and timed-out requests included at their clipped value. [†]marks cells that did not sustain the offered rate (achieved<95% or any timeouts).

Layer	250/s	500/s	1000/s	2000/s	4000/s
pg	2	2	2	3	1483 [†]
prisma	3	3	13	10721 [†]	10947 [†]
knex	2	2	2	89	8812 [†]
sequelize	3	3	183	10745 [†]	10940 [†]
mikroorm	4	346	11283 [†]	11677 [†]	11935 [†]

7 Pool-size sensitivity

The primary campaign fixes every layer’s connection pool at 10 to isolate the access layer from pool tuning. Table S7 varies the pool from 1 to 50 for

a representative layer of each tier on the deep fetch (PostgreSQL, 50 connections). The deep-fetch ordering (`pg` > `Knex` > `Prisma` > `MikroORM`) is preserved at every pool size, and each layer’s throughput saturates at a pool well below 50, so the fixed-pool choice does not appear to drive the ranking; a very small pool (1–2) is the one regime that compresses the differences, because every layer then queues on connection acquisition.

Table S7: Deep-fetch throughput (req/s, PostgreSQL, 50 connections, median of 3) as the connection pool varies from 1 to 50, per access layer. The primary campaign fixes the pool at 10; each layer saturates at a pool well below 50 and the ordering is preserved at every pool size, so the fixed-pool choice does not appear to drive the ranking.

Layer	1	2	5	10	20	50
<code>pg</code>	905	1762	3224	3667	3798	3768
<code>knex</code>	741	1413	2392	2532	2683	2636
<code>prisma</code>	870	1072	1183	1234	1207	1105
<code>mikroorm</code>	494	487	507	527	527	522

8 Transactional multi-statement write

Beyond the single-row insert of the primary matrix, Table S8 reports a transactional write, `POST /threads`, which inserts a post and five comments in one transaction through each layer’s transaction facility, for a representative layer of each tier on PostgreSQL under default durability (median of five runs, exploratory). The full layer $\times$ engine matrix for this pattern is left to future work. Because the transaction commits once, its per-commit fsync dominates and the layers cluster more tightly than the widest read spread ($3.7\times$, `pg` 2,718 down to `Prisma` 733), though the ordering does not simply track the single-row insert: `Prisma`, mid-pack on the single-row PostgreSQL insert, falls to the bottom of the transactional subset.

9 Dispersion on PostgreSQL

Table S9 gives the full per-layer, per-pattern coefficient of variation of throughput on PostgreSQL across the 25 repeated runs (maximum 7.3%, `knex`/insert), the companion to the MySQL table in Section S1.

10 Resource footprint

Table S10 gives the per-layer application-process CPU (median % of one logical core) and peak resident memory on the deep/nested fetch (MySQL),

Table S8: Transactional write throughput (req/s) and tail latency (p99, ms) for `POST /threads`, which inserts a post and five comments in a single transaction, on PostgreSQL under default durability (50 connections, median of 5). A representative layer of each tier is shown; the full matrix is left to future work. Between-layer spread is $3.7\times$ (pg 2718 down to Prisma 733, now the slowest), well below the $7.15\times$ read spread, and the ordering broadly tracks the single-row insert apart from Prisma (mid-pack single-row, last here), so the write conclusions largely extend to a multi-statement transactional write.

Layer	req/s	p99 (ms)
pg	2718	22
knex	2322	27
prisma	733	80
typeorm	1935	32
mikroorm	856	68

Table S9: Coefficient of variation (%) of throughput across the 25 repeated runs, per access layer and pattern on PostgreSQL. Low values indicate stable measurements.

Layer	point_read	range_scan	deep_fetch	aggregation	write
pg	4.0	2.3	3.4	2.0	5.3
knex	4.1	3.0	2.7	3.6	7.3
drizzle	2.8	2.5	2.0	3.4	4.0
prisma	3.0	2.5	3.0	2.7	3.8
sequelize	2.6	3.0	2.2	2.2	2.6
typeorm	2.3	2.1	2.6	3.9	3.6
objection	2.8	2.4	2.6	3.3	5.0
mikroorm	2.4	3.8	4.0	4.2	2.5

the detail behind the CPU trade-off discussed in the main text (all layers $\approx 100\text{--}110\%$ of one core; peak RSS 215–356 MB).

Table S10: Server-process resource footprint on the deep/nested fetch (MySQL): median CPU utilization (% of one core) and peak resident memory (MB) of the Node.js process under load. The load generator and database run in separate processes.

Layer	Category	CPU (%)	RSS (MB)
mysql2	native-driver	105	294
knex	query-builder	100	215
drizzle	orm	105	297
prisma	orm	105	334
sequelize	orm	100	262
typeorm	orm	100	225
objection	orm	100	221
mikroorm	orm	110	356

11 Pinned versions

Table S11 lists the pinned versions of every evaluated access layer and tool, recorded from the logfile on the measurement date.

Table S11: Pinned versions of the evaluated access layers and tooling.

Layer / tool	Version	Layer / tool	Version
pg	8.22.0	sequelize	6.37.8
mysql2	3.23.0	typeorm	1.1.0
knex	3.3.0	objection	3.1.5
drizzle-orm	0.45.2	@mikro-orm/core	7.1.6
@prisma/client	7.8.0	express	5.2.1
autocannon	8.0.0	Node.js	24.18.0
PostgreSQL	18.4	MySQL	9.7.1

Pinned to the latest stable release compatible with the harness adapter contract at the freeze date (2026-07-15); Sequelize’s version-7 line is a prerelease, so the 6.x stable line is used. Prisma 7 is Rust-free and connects through an official JS driver adapter (`@prisma/adapters-pg` for PostgreSQL, `@prisma/adapters-mariadb` with the `mariadb` 3.5.3 driver for MySQL, as Prisma ships no `mysql2` adapter).

12 Per-pattern detail: the read patterns

The two cheapest read patterns show the smallest between-layer spread and are summarized in the main text; their full throughput and p99 tables are given here. Table S12 reports the point read (primary-key lookup) and Table S13 the keyset range scan, each by access layer and engine. The consequential patterns (deep/nested fetch, aggregation, insert) are tabulated in the main text.

Table S12: Point read: throughput (req/s, higher is better; median with a within-campaign 95% bootstrap interval over the 25 repeated runs — run-to-run variability within one host and campaign, not across hardware or days) and response-time p99 (ms, lower is better; median run-level p99 with the same within-campaign 95% bootstrap interval) by access layer and engine. p99 is reported at 1 ms resolution, so cells differing by ≤ 1 ms are practically unresolved (see the paired significance analysis in the supplement).

Layer	Category	PostgreSQL		MySQL	
		req/s [95% CI]	p99 [95% CI]	req/s [95% CI]	p99 [95% CI]
pg	native-driver	6835 [6597–6892]	9 [9–10]	–	–
mysql2	native-driver	–	–	4368 [4314–4420]	16 [16–16]
pg-tuned	native-tuned	6806 [6622–6964]	9 [9–9]	–	–
mysql2-tuned	native-tuned	–	–	4268 [4212–4331]	14 [13–14]
knex	query-builder	4942 [4794–5131]	13 [12–13]	3814 [3771–3870]	15 [15–16]
drizzle	orm	4199 [4161–4254]	14 [14–15]	2928 [2866–2977]	21 [21–22]
prisma	orm	3255 [3222–3310]	20 [20–21]	2074 [2050–2095]	29 [28–29]
sequelize	orm	3518 [3483–3591]	17 [16–17]	2764 [2726–2798]	20 [20–21]
typeorm	orm	4185 [4171–4242]	15 [15–15]	3059 [3032–3097]	19 [19–19]
objection	orm	3977 [3943–4040]	15 [15–16]	3275 [3221–3310]	17 [17–18]
mikroorm	orm	2455 [2425–2484]	24 [23–25]	2067 [2044–2086]	27 [26–28]

13 Same-SQL control: full table

Table S14 gives the full same-SQL (raw-path) control on the deep fetch for both engines: throughput when every layer executes the identical control statement (verified byte-identical by server-side capture). The main text reports the resulting collapse toward parity and reads it as a standardized-SQL contrast, not a bound, on the residual same-SQL difference — a compound intervention that isolates no single factor.

14 Per-pattern detail: aggregation

Table S15 gives the full aggregation-pattern throughput and p99 by access layer and engine. RQ3 finds aggregation the least layer-sensitive pattern

Table S13: Keyset range scan: throughput (req/s, higher is better; median with a within-campaign 95% bootstrap interval over the 25 repeated runs — run-to-run variability within one host and campaign, not across hardware or days) and response-time p99 (ms, lower is better; median run-level p99 with the same within-campaign 95% bootstrap interval) by access layer and engine. p99 is reported at 1 ms resolution, so cells differing by ≤ 1 ms are practically unresolved (see the paired significance analysis in the supplement).

Layer	Category	PostgreSQL		MySQL	
		req/s [95% CI]	p99 [95% CI]	req/s [95% CI]	p99 [95% CI]
pg	native-driver	3811 [3747–3865]	18 [17–18]	–	–
mysql2	native-driver	–	–	3295 [3274–3341]	21 [20–21]
pg-tuned	native-tuned	3830 [3766–3872]	18 [17–19]	–	–
mysql2-tuned	native-tuned	–	–	3252 [3233–3271]	19 [17–19]
knex	query-builder	3220 [3198–3240]	18 [17–19]	2995 [2942–3022]	20 [19–21]
drizzle	orm	2780 [2749–2822]	20 [20–21]	2165 [2145–2192]	29 [29–30]
prisma	orm	2443 [2428–2473]	26 [26–27]	1538 [1517–1556]	39 [38–40]
sequelize	orm	2478 [2424–2498]	23 [23–24]	2105 [2090–2117]	27 [27–28]
typeorm	orm	2564 [2537–2590]	24 [24–25]	2262 [2150–2298]	26 [25–27]
objection	orm	2686 [2659–2697]	22 [21–23]	2582 [2551–2597]	23 [22–23]
mikroorm	orm	905 [891–915]	64 [63–68]	856 [846–865]	66 [66–71]

Table S14: Standardizing SQL — the residual same-SQL spread on the deep fetch: throughput (req/s) of the documentation-selected deep fetch versus the *same-SQL* control (median of 10 runs), in which every layer executes the identical two-statement SQL with identical row mapping through its raw-SQL facility. The ratio (default $\div$ same-SQL) measures the gap between each layer’s documentation-selected relational-fetch path and the common raw-SQL path, a composite of query strategy, query count, protocol, the raw versus ORM API, and result marshalling and hydration changed together; the residual native-relative same-SQL spread ($1.68\times$ on PostgreSQL, $2.01\times$ on MySQL) is a standardized-SQL contrast on that compound difference, not a bound, and isolates no single factor.

Layer	PostgreSQL			MySQL		
	default	same-SQL	ratio	default	same-SQL	ratio
pg	3687	3625	1.02	–	–	–
mysql2	–	–	–	2382	2427	0.98
knex	2487	2926	0.85	2007	2245	0.89
drizzle	1865	3237	0.58	1318	2008	0.66
prisma	1186	2155	0.55	634	1206	0.53
sequelize	991	2436	0.41	886	1903	0.47
typeorm	1204	3170	0.38	1017	2267	0.45
objection	1028	2831	0.36	834	2310	0.36
mikroorm	516	3239	0.16	480	2453	0.20

once every layer runs the same raw correlated-subquery plan; the main text reports the ranking and the leading figures.

Table S15: Aggregation: throughput (req/s, higher is better; median with a within-campaign 95% bootstrap interval over the 25 repeated runs — run-to-run variability within one host and campaign, not across hardware or days) and response-time p99 (ms, lower is better; median run-level p99 with the same within-campaign 95% bootstrap interval) by access layer and engine. p99 is reported at 1 ms resolution, so cells differing by ≤ 1 ms are practically unresolved (see the paired significance analysis in the supplement).

Layer	Category	PostgreSQL		MySQL	
		req/s [95% CI]	p99 [95% CI]	req/s [95% CI]	p99 [95% CI]
pg	native-driver	6978 [6865–7038]	11 [10–11]	–	–
mysql2	native-driver	–	–	3994 [3941–4018]	18 [17–19]
pg-tuned	native-tuned	7538 [7493–7705]	10 [9–10]	–	–
mysql2-tuned	native-tuned	–	–	3958 [3933–3985]	15 [15–17]
knex	query-builder	5666 [5608–5710]	13 [12–14]	3842 [3815–3883]	19 [18–20]
drizzle	orm	6280 [6156–6427]	12 [12–12]	3520 [3447–3541]	21 [19–21]
prisma	orm	4408 [4368–4476]	16 [15–16]	2328 [2301–2349]	30 [28–31]
sequelize	orm	4813 [4758–4870]	14 [14–15]	3216 [3203–3235]	20 [18–22]
typeorm	orm	6098 [5985–6179]	13 [12–13]	3767 [3731–3812]	18 [15–20]
objection	orm	3812 [3724–3851]	17 [17–18]	2763 [2743–2785]	25 [23–26]
mikroorm	orm	5745 [5637–5805]	13 [12–14]	3741 [3711–3804]	19 [18–21]

15 Documentation-selected deep-fetch choices per access layer

Table S16 documents, for each access layer, the exact eager-loading API used on the deep/nested fetch, the number of SQL round trips it issues, the documented alternative loading strategy we did *not* use, and the query-logging state. Every layer uses its documented default path; query logging is disabled and verified on all eleven adapters, and no lifecycle hooks or validation run on the read path. The six data-mapper and ORM layers split both ways on the join-versus-select-in default: Sequelize, TypeORM, and MikroORM default to a single join, while Prisma and Objection default to separate batched queries. The loading-strategy sensitivity check reported in the main text (Table S18) switches the join-default layers to select-in where that alternative is a byte-identical drop-in. Full per-adapter detail is in the replication package’s `METHODOLOGY.md`.

16 Open-loop tail latency on MySQL

Table S17 is the MySQL companion to the PostgreSQL open-loop sweep (Supplement Table S6): coordinated-omission-corrected p99 on the deep/nested

Table S16: Documentation-selected deep-fetch implementation per access layer: the eager-loading API actually used, the SQL round trips it issues (measured from server-side statement logging for the portable layers; by construction for the native variants), the documented alternative loading strategy we did *not* use, and the query-logging state. Query logging is disabled everywhere; no lifecycle hooks or validation run on the read path (details in the replication package’s `METHODOLOGY.md`).

Layer	Deep-fetch API	Round-trips	Alternative not used	Logging
pg	hand-written JOIN $\times 2$	2	n/a (hand-written SQL)	off
mysql2	hand-written JOIN $\times 2$	2	n/a (hand-written SQL)	off
pg-tuned	named-prepared JOIN $\times 2$	2	n/a (hand-written SQL)	off
mysql2-tuned	execute() JOIN $\times 2$	2	n/a (hand-written SQL)	off
knex	builder .join() $\times 2$	2	n/a (hand-written SQL)	off
drizzle	builder .innerJoin() $\times 2$	2	relational query API (<code>db.query</code>)	off
prisma	include:	4	<code>relationLoadStrategy: 'join'</code>	off
sequelize	include: (<code>separate:false</code>)	1	<code>separate:true</code> (select-in)	off
typeorm	relations:	2	<code>relationLoadStrategy: 'query'</code>	off
objection	.withGraphFetched()	4	.withGraphJoined()	off
mikroorm	populate:	1	<code>strategy: 'select-in'</code>	off

fetch at a constant offered rate. The sub-saturation tail is flat and small on both engines, and each layer collapses once the offered rate crosses its capacity, so the high-load-tail ordering holds on MySQL as on PostgreSQL.

Table S17: Open-loop tail latency on **MySQL**, corrected for coordinated omission: p99 (ms) on the deep/nested fetch under a constant offered request rate, every latency measured from the intended start time and timed-out requests clipped into the distribution. [†]marks cells that did not sustain the offered rate. As on PostgreSQL (Supplement Table S6), the sub-saturation tail is flat and small and each layer collapses once the offered rate crosses its capacity, so the high-load-tail ordering is engine-robust.

Layer	250/s	500/s	1000/s	2000/s	4000/s
mysql2	2	2	3	15	10856 [†]
prisma	5	7	9518 [†]	11723 [†]	11900 [†]
knex	2	2	2	243	10904 [†]
sequelize	3	3	1628 [†]	10797 [†]	10962 [†]
mikroorm	4	953 [†]	11314 [†]	11788 [†]	11912 [†]

17 Alternative eager-loading sensitivity

Table S18 reports the loading-strategy sensitivity check: for the data-mapper ORMs whose alternative strategy is a byte-identical drop-in, documentation-selected deep-fetch throughput versus the alternative, on both engines. Both

endpoints were verified to return byte-identical responses before timing.

Table S18: Alternative eager-loading sensitivity on the deep/nested fetch: documentation-selected throughput (req/s) versus the alternative loading strategy, both engines, for the layers whose alternative is a byte-identical drop-in. Sequelize and MikroORM switch from their documentation-selected single join to select-in; in both cases the documentation-selected default is the *faster* strategy, so their deep-fetch deficit does not appear to be an artifact of a poor default. TypeORM is omitted, its alternative strategy not a byte-identical drop-in in the pinned versions (its query strategy errors on the ordered nested `findOne`, and Objection’s `withGraphJoined` changes row handling), which is itself a portability caveat.

Layer	Switch	PostgreSQL		MySQL	
		default	alt.	default	alt.
sequelize	join→select-in	915	690	830	602
mikroorm	join→select-in	447	357	420	324
objection	select-in→join	—	—	820	1301

18 MySQL insert commit-flush mechanism

Table S19 gives the direct `performance_schema` evidence behind the MySQL insert bottleneck: sampling the redo-log and binary-log flush wait instruments around a fixed insert workload at default durability, the flush is a shared engine floor—similar per-insert across the native driver, query builder, and ORM despite their very different read performance—and relaxing durability removes it.

19 Post-restart environmental-robustness check

Table S20 reports the environmental-robustness check: as the campaign’s final stage, after a full restart of both database engines (cold buffer pools, fresh connections), a representative deep-fetch subset was re-measured against the primary median. The layer ranking reproduces on both engines (**pg** > Prisma ≫ MikroORM on PostgreSQL; **mysql2** ≫ Prisma > MikroORM on MySQL); absolute throughput is up to 16% lower, a cold-buffer and run-to-run drift the paper’s relative analysis is designed to absorb. Because PostgreSQL and MySQL cells are not measured simultaneously, this drift is non-uniform across engines (0% on **mysql2**, 16% on **pg**, 16% on MikroORM/PostgreSQL), so the cross-engine interaction ratios (Table S28) could in principle carry some residual drift; the engine effects there are far larger than the 16% drift, so the engine-dependence conclusion holds.

Table S19: MySQL insert commit-flush mechanism (`performance_schema` wait instruments, default durability `innodb_flush_log_at_trx_commit=1`, `sync_binlog=1`): per-insert wall time, the redo-log+binlog flush wait per insert, and the flush share of wall time, over 4000 inserts per layer. The flush wait is a shared engine floor – similar across layers despite their very different read performance – so it caps insert throughput regardless of the access layer. Relaxing durability (`mysql2`, `trx_commit=0`, `sync_binlog=0`) removes the flush and lifts throughput to 3238 req/s, confirming the floor.

Layer (MySQL)	req/s	per-insert (ms)	flush/insert (ms)	flush share
<code>mysql2</code>	1576	0.634	1.001	158%
<code>knex</code>	1722	0.581	1.009	174%
<code>mikroorm</code>	2584	0.387	0	0%
<i>Relaxed durability (control):</i>				
<code>mysql2*</code>	3238	0.309	0.01	3%

Table S20: Environmental-robustness check (review 6.7): as the campaign’s final stage, after a full restart of both database engines (cold buffer pools, fresh connections), deep-fetch throughput of a representative subset was re-measured against the primary median. The layer ranking reproduces on both engines; absolute throughput drifts with host conditions (larger on PostgreSQL), a cold-buffer and run-to-run variation the paper’s *relative* analysis is designed to absorb. Single host, so this checks restart robustness, not cross-machine generalization.

Engine	Layer	primary req/s	post-restart req/s	ratio
PostgreSQL	<code>pg</code>	3687	3100.5	0.84
PostgreSQL	<code>prisma</code>	1186	1153	0.97
PostgreSQL	<code>mikroorm</code>	516	434.5	0.84
MySQL	<code>mysql2</code>	2382	2375.5	1.00
MySQL	<code>prisma</code>	634	618.5	0.98
MySQL	<code>mikroorm</code>	480	437.5	0.91

20 Utilization-controlled tail latency

Tables S21 and S22 give the coordinated-omission-corrected p99 on the deep fetch when each layer is offered 50/70/85/95% of *its own* saturating throughput (its primary closed-loop deep-fetch rate), replicated five times, on both engines. At matched utilization the tails are small and similar across layers—for example on PostgreSQL every layer holds p99 near 2–4 ms at 50%—so the wide high-load p99 ladder in the main text is a capacity-and-queueing effect, not a difference in latency at a matched offered-load fraction; the tail rises only as a layer nears its own capacity. This is the capacity-normalized latency companion to the high-load primary p99. The 50/70% cells, which carry that comparison, are stable across runs; the 85/95% cells are increasingly noisy (p99 varies several-fold between runs as a layer approaches saturation), so those columns are read as trend, not precise values.

Table S21: Utilization-controlled open-loop tail on the deep fetch (PostgreSQL): coordinated-omission-corrected p99 (ms, median of five runs) when each layer is offered a fixed fraction of *its own* saturating throughput. At matched utilization the tails are small and similar across layers—the large high-load p99 gap is a capacity/queueing effect, not a sub-saturation latency difference; the tail rises only as each layer approaches its own capacity.

Layer	50%	70%	85%	95%
pg	3.9	3.1	4	5.9
knex	2.4	3.2	3.5	54
drizzle	3.8	8.8	30.9	7.4
prisma	4.2	7.8	18.8	77.5
sequelize	2.5	4.3	7.7	11.8
typeorm	3.1	3.7	7.9	37.1
objection	3.1	3.1	6.2	284.6
mikroorm	3.9	4.9	5.5	7.5

21 Longer-window tail sensitivity

The primary p99 is measured over a 12 s window, which yields only ≈ 60 requests above the p99 threshold in the slowest deep-fetch cells. To check that this per-run tail count is sufficient, Table S23 re-measures the deep fetch over a 60 s window (about five times the requests, ≈ 300 tail observations per run) at the same 50-connection operating point, ten repeated runs, on both engines. The p99 ranking is preserved exactly (Spearman $\rho = 1.00$ per engine), each cell’s p99 barely moves from the 12 s value, the p97.5 ordering matches the p99 ordering ($\rho = 1.00$), and the run-to-run p99 coefficient

Table S22: Utilization-controlled open-loop tail on the deep fetch (MySQL): coordinated-omission-corrected p99 (ms, median of five runs) when each layer is offered a fixed fraction of *its own* saturating throughput. At matched utilization the tails are small and similar across layers—the large high-load p99 gap is a capacity/queueing effect, not a sub-saturation latency difference; the tail rises only as each layer approaches its own capacity.

Layer	50%	70%	85%	95%
<code>mysql2</code>	4.2	7.7	20.9	39.7
<code>knex</code>	1.9	2.7	13.7	4.1
<code>drizzle</code>	3.7	6.1	14.9	17.6
<code>prisma</code>	5.2	5.6	11.9	51.4
<code>sequelize</code>	2.5	3.2	6.8	21
<code>typeorm</code>	2.8	3.6	3	36.5
<code>objection</code>	2.9	3.2	3.7	32.9
<code>mikroorm</code>	4	4.2	8.7	31.2

of variation stays small (maximum 6.0% on PostgreSQL, 3.8% on MySQL). The shorter window is therefore adequate for the tail *ranking*, and the tail conclusions do not rest on the exact per-run tail count.

22 Multi-worker scaling

Table S24 scales a representative subset (three layers per engine) to 1, 2, and 4 Express workers (a shared-port `node` cluster), median of three runs. Because each process uses about one core and gains nothing from more, aggregate throughput rises roughly in proportion to workers over this range: on PostgreSQL the native driver leads at one worker (`pg` 3,299 to Prisma’s 1,117); its aggregate throughput then plateaus by two workers (4,055, then 3,961 at four) — whether from database contention, shared-host CPU competition, or the fixed-connection generator is not instrumented here — while Prisma, low per process, rises 1,117 to 2,275 to 4,449 and reaches native-like aggregate throughput at four workers. This is a bounded check, not a demonstration of general scalability: it stops at four workers, measures no database-side saturation (DB CPU, connections, or wait events), and on MySQL horizontal scaling *widens* the gap — the native `mysql2` keeps scaling to 8,883 req/s at four workers while Prisma reaches only 2,427 (about `mysql2` at one worker). Within these limits a layer’s low per-process throughput may be addressed by additional application processes until another bottleneck is reached, at a proportional cost in cores — a partial test of the single-host resource-competition concern.

Table S23: Longer-window tail sensitivity on the deep/nested fetch. Each cell’s primary p99 is measured over a 12 s window (≈ 60 tail observations per run in the slowest cells); the re-measurement uses a 60 s window (≈ 300 tail observations) at the same 50-connection operating point, 10 repeated runs. The p99 ranking is preserved exactly (Spearman $\rho = 1.00$ on both engines), the per-cell p99 barely moves, the p97.5 ordering matches the p99 ordering ($\rho = 1.00$), and the run-to-run p99 CV stays small (max 6.0% PostgreSQL, 3.8% MySQL), so the 12 s window is adequate for the tail ranking despite the modest per-run tail count.

Layer	p99 12 s (ms)	p99 60 s (ms)	p97.5 60 s (ms)	p99 CV (%)	rps 60 s
<i>PostgreSQL</i>					
pg-tuned	18	18	17	2.3	3785.5
pg	20	20.5	18.5	6.0	3687.5
knex	27	28	26	2.9	2460
drizzle	35	35	33	2.3	1855.5
typeorm	50	51	49	3.7	1222
prisma	51	52.5	50	3.1	1176.5
objection	57	59	56.5	2.9	1026
sequelize	60	60	57	1.5	994
mikroorm	116	112.5	109	2.8	519.5
<i>MySQL</i>					
mysql2-tuned	25	25.5	24	3.3	2409.5
mysql2	28	28	25.5	2.6	2443.5
knex	30	29.5	28	3.3	2048
drizzle	47	48	45	3.2	1297
typeorm	57	58.5	56	1.8	1015.5
sequelize	67	66	63	2.2	889
objection	69	70	67.5	2.1	845.5
prisma	93	95	92	3.2	630
mikroorm	120	123.5	119	3.8	478.5

Table S24: Multi-worker deep-fetch throughput (req/s, median of three runs) as each layer is scaled to 1, 2, and 4 Express workers (node cluster, shared port). Because each process uses about one core, aggregate throughput rises roughly in proportion to workers over this range: Prisma, low at one worker (PostgreSQL 1,117), rises to 2,275 at two and 4,449 at four workers, recovering native-like aggregate throughput at four times the processes. This bounded 1-to-4-worker check measures no database-side saturation and stops before any layer’s knee, so a layer’s low per-process throughput may be addressed by additional application processes until another bottleneck is reached, at a proportional cost in cores.

Layer	1 worker	2 workers	4 workers
<i>PostgreSQL</i>			
pg	3299	4055	3961
prisma	1117	2275	4449
mikroorm	435	898	1773
<i>MySQL</i>			
mysql2	2327	4698	8883
prisma	591	1167	2427
mikroorm	418	851	1624

23 Pool-size frontier on MySQL

Table S25 is the MySQL companion to the PostgreSQL pool sweep (Supplement Table S7): deep-fetch throughput as the connection pool varies from 1 to 50, per layer. Each layer saturates at a pool well below 50 and the ordering is preserved at every pool size, so the fixed pool of ten does not appear to drive the ranking on either engine.

Table S25: Deep-fetch throughput (req/s, MySQL, 50 connections, median of 3) as the connection pool varies from 1 to 50, per access layer. The primary campaign fixes the pool at 10; each layer saturates at a pool well below 50 and the ordering is preserved at every pool size, so the fixed-pool choice does not appear to drive the ranking.

Layer	1	2	5	10	20	50
mysql2	1545	2156	2352	2446	2506	2494
knex	1255	1926	1984	2081	2131	2120
prisma	464	572	631	639	626	648
mikroorm	416	469	486	418	492	429

24 Mixed read/write workload

Table S26 interleaves the deep-fetch read with single-row inserts at 90:10 and 70:30 read:write, both engines, with a physical write reset per run. The layer ordering matches the read ordering and is insensitive to the mix, because the deep-fetch read dominates the cost; exploratory.

Table S26: Mixed read/write workload: throughput (req/s) and p99 (ms) under an autocannon request mix interleaving the deep-fetch read with single-row inserts, at 90:10 and 70:30 read:write, both engines, median of five runs with a physical write reset per run. The layer ordering matches the read ordering and is insensitive to the mix (the deep-fetch read dominates); exploratory.

Engine	Layer	read:write	req/s	p99
PG	pg	90:10	4236	19
PG	pg	70:30	4347	18
PG	knex	90:10	3024	26
PG	knex	70:30	2993	26
PG	prisma	90:10	1504	55
PG	prisma	70:30	1409	58
PG	mikroorm	90:10	676	103
PG	mikroorm	70:30	705	99
MySQL	mysql2	90:10	2283	44
MySQL	mysql2	70:30	2505	47
MySQL	knex	90:10	2275	37
MySQL	knex	70:30	2232	40
MySQL	prisma	90:10	718	114
MySQL	prisma	70:30	729	103
MySQL	mikroorm	90:10	626	109
MySQL	mikroorm	70:30	639	105

25 Concurrency sweep

Figure S2 sweeps the deep/nested fetch from 1 to 256 client connections for a representative layer of each tier, both engines. The 50-connection operating point used throughout sits above the knee for every layer on this pattern; the ordering is load-robust above about 32 connections. This is the evidence behind the choice of operating point.




Figure S2: Throughput of each access layer on the deep/nested fetch as concurrency increases (postgres). The native driver leads at every concurrency level, with the query builder and Drizzle next and the data-mapper ORMs (Prisma among them) below; the ordering is stable above about 32 connections and compresses only at very low concurrency.

26 Per-layer cross-engine ranks

Table S27 gives each portable layer’s throughput rank on PostgreSQL versus MySQL for the deep fetch, aggregation, and the insert, showing the concrete rank reversals behind the main text’s engine-dependence finding. The deep-fetch orderings largely agree, but aggregation and the insert reverse at the top: on aggregation **drizzle** (PostgreSQL rank 1) and **knex** (MySQL rank 1) trade the lead, and on the insert **knex** drops from rank 1 to 3 while **drizzle** rises to 1, with Prisma falling from mid-pack (rank 4) to last (rank 7) on both. The rank correlations these reversals produce ($\rho = 0.86$ deep fetch, 0.64 aggregation, 0.68 insert; coarse at $n = 7$) are reported only as descriptive summaries.

Because MikroORM is the largest single outlier in several spreads, we recompute the cross-engine agreement without it: over the remaining six portable layers the read ordering stays similar (deep-fetch $\rho = 0.77$, $\tau = 0.60$; range $\rho = 0.83$, $\tau = 0.73$) and the deep-fetch native-relative spread stays substantial at $3.72\times$ (PostgreSQL) and $3.76\times$ (MySQL), so the RQ2 and RQ3 findings do not depend on that one layer.

Table S27: Per-layer throughput rank (1 = fastest) on PostgreSQL versus MySQL for the seven portable layers, showing the concrete rank reversals behind the engine-dependence finding. The deep/nested-fetch orderings largely agree; aggregation and the insert reverse at the *top*, not among near-ties: on aggregation **drizzle** (PostgreSQL rank 1) and **knex** (MySQL rank 1) trade the lead (a three-place move each), and on the insert **knex** drops from rank 1 to 3 while **drizzle** rises from 2 to 1; Prisma falls from mid-pack (rank 4) to last (rank 7) on both. These reversals, not the coarse ($n = 7$) coefficients they produce (deep fetch $\rho = 0.86$, $\tau = 0.71$; aggregation $\rho = 0.64$; insert $\rho = 0.68$, $\tau = 0.52$), are the engine-dependence result. Excluding MikroORM (the largest single outlier) leaves the read ordering similar (deep-fetch $\rho = 0.77$, $\tau = 0.60$) and the deep-fetch native-relative spread at $3.72\times$ (PostgreSQL) and $3.76\times$ (MySQL), so no single layer drives these findings.

Layer	Deep fetch		Aggregation		Insert	
	PG	MySQL	PG	MySQL	PG	MySQL
knex	1	1	4	1	1	3
drizzle	2	2	1	4	2	1
prisma	4	6	6	7	4	7
sequelize	6	4	5	5	5	4
typeorm	3	3	2	2	6	5
objection	5	5	7	6	3	2
mikroorm	7	7	3	3	7	6

27 Layer-by-engine interaction magnitude

The blocked permutation test reports *that* the layer ordering interacts with the engine; Table S28 reports *how large* the interaction is, as the PostgreSQL $\div$ MySQL throughput ratio per layer and pattern with a 95% paired-bootstrap interval. Every layer is faster on PostgreSQL, so the interaction is in the spread of the advantage: the reads stay within ≈ 1.0 – 1.9 while the insert scatters from 1.6 to 3.2, the quantitative counterpart of the low insert rank correlation.

Table S28: Layer $\times$ engine interaction magnitude: the PostgreSQL $\div$ MySQL throughput ratio (geometric mean of the paired per-replicate ratios, with a within-campaign 95% paired-bootstrap interval) for each portable layer and pattern. A ratio > 1 means the layer runs faster on PostgreSQL, < 1 faster on MySQL. Every layer is faster on PostgreSQL here, so the interaction is in the *spread* of the advantage, not its sign: point read and range scan hold a tight band (≈ 1.0 – 1.6), the deep fetch and aggregation reach ≈ 1.9 , and the insert scatters widest (1.6–3.2), reordering the layers across engines — the interaction the blocked permutation test detects and the rank reversals of Table S27 make concrete.

Layer	Point read	Range scan	Deep fetch	Aggregation	Insert
knex	1.31 [1.28,1.34]	1.07 [1.06,1.09]	1.23 [1.21,1.25]	1.46 [1.44,1.48]	3.05 [2.89,3.23]
drizzle	1.44 [1.42,1.46]	1.28 [1.27,1.30]	1.41 [1.40,1.42]	1.80 [1.77,1.83]	2.62 [2.46,2.79]
prisma	1.57 [1.55,1.59]	1.59 [1.57,1.62]	1.89 [1.85,1.92]	1.90 [1.87,1.92]	3.23 [3.14,3.32]
sequelize	1.28 [1.26,1.29]	1.17 [1.15,1.19]	1.12 [1.10,1.13]	1.49 [1.48,1.51]	1.82 [1.73,1.91]
typeorm	1.38 [1.36,1.41]	1.15 [1.12,1.17]	1.19 [1.17,1.21]	1.60 [1.57,1.63]	2.05 [1.97,2.13]
objection	1.22 [1.20,1.24]	1.04 [1.03,1.05]	1.22 [1.21,1.24]	1.37 [1.35,1.40]	2.38 [2.28,2.50]
mikroorm	1.19 [1.17,1.21]	1.06 [1.04,1.07]	1.06 [1.04,1.08]	1.51 [1.48,1.54]	1.60 [1.56,1.64]

28 A reporting checklist for access-layer benchmarks

Distilled from the design decisions and the pitfalls this study confronted, the following items are the ones an access-layer benchmark should report so that its claims can be scoped and reproduced. It is a reporting aid, not a standard.

1. **Correctness before timing.** Assert byte-identical responses across layers (final serialized bodies), or state what differs; several defects in this study were caught only this way.
2. **Experimental unit.** Treat the freshly booted run, not the individual request, as the unit; requests within a run are not independent.
3. **Same-host repetition is not environmental replication.** Report how many runs, on how many hosts, over what period; do not call same-host repeats “independent.”

4. **Documentation-selected treatment.** State the exact API, its documentation page and version, the declined alternative, and whether logging, hooks, validation, and identity maps are on.
5. **Latency construct.** Distinguish high-load closed-loop response time (which tracks throughput and queueing) from utilization-controlled or open-loop latency; report the operating point.
6. **Same-SQL / shared-plan control.** Report it as a compound intervention (it changes query formulation, protocol, preparation, round-trips, and mapping together), not a mechanism decomposition.
7. **Resource accounting.** Report application-tier and database CPU separately; throughput per CPU-second, not throughput alone.
8. **Write-state control.** Reset the mutable state between runs and state the durability regime (default versus relaxed).
9. **Provenance.** Pin engine images by digest and dependencies by lock-file; archive raw per-cell records, seeds, and a table-to-record manifest.
10. **Scope.** State the host topology, workload, versions, and dates, and confine claims to them.

29 Study factors

Table [S29](#) records which factors the benchmark holds constant, which it varies, and which are uncontrolled on the single virtualized host; the uncontrolled factors are why results are reported for this configuration rather than as a general library ranking.

30 Paired significance of the p99 ladder

Table [S30](#) is the p99 counterpart of the main-text throughput significance table (Table 6): the same paired permutation, Wilcoxon, geometric-mean-ratio, paired-bootstrap-CI, and dominance analysis applied to each adjacent layer’s per-replicate deep-fetch p99 on both engines, delivering the per-cell p99 inference the outcomes table (Table 3) lists as a primary analysis. Every adjacent pair separates on both engines except two near-ties at ratio 1.03 — the TypeORM–Prisma step on PostgreSQL (permutation $p = 0.05$) and the Sequelize–Objection step on MySQL ($p = 0.08$); the p99 ordering otherwise mirrors the throughput ranking. Because the high-load tail tracks capacity, these are resolved capacity-and-queueing differences, not differences in sub-saturation latency.

Table S29: Factors in the benchmark: what is held constant, what varies (the treatments), and what is uncontrolled. The same-SQL control additionally holds the generated SQL and row mapping constant; because it changes query formulation, API, protocol, statement preparation, round-trip structure, and mapping *together*, it bounds the residual difference but does not isolate any single factor. The uncontrolled factors are inherent to a single virtualized host and are the reason results are reported for this configuration, not as general library performance.

Status	Factor	Detail
Held constant	schema, indexes, seed data	identical DDL + deterministic seed, both engines
	per-replicate request stream	paired PRNG stream, seeded on (endpoint, replicate)
	connection pool	fixed at ten for every layer
	response bytes	byte-identical JSON across layers (verified)
	host, OS, engine versions	one host; pinned July-2026 versions (Supplement Table S11)
	warm-up, run duration, cell order	15s warm-up, 12s runs, randomized order
Varies (treatments)	access layer	9 documentation-selected + 2 tuned implementations
	database engine	PostgreSQL, MySQL
	access pattern	point read, range scan, deep fetch, aggregation, insert
Uncontrolled	hypervisor neighbours, CPU steal	single virtualized host
	thermal / frequency scaling	not pinned; bounded by the drift and post-restart checks
	database cache / kernel history	warm by design; randomized order + fresh boot mitigate

31 Benchmarking-pitfalls checklist

The four confounds summarized in the main text’s Discussion, in full. Each is a generic failure mode of the genre, easy to introduce by accident and hard to detect: the correctness ones without an automated cross-check, the performance ones without profiling every cell.

- **Aggregation fan-out.** A naive join between `posts` and `comments` evaluated before the `SUM` multiplies each post’s view count by its number of matching comment rows, inflating the aggregate — a correctness bug invisible unless every layer’s output is checked against a reference, not a performance one.
- **OFFSET vs. keyset pagination.** An early, `OFFSET`-based design for the range-scan endpoint collapsed on MySQL at realistic offsets, which would have been misread as an access-layer weakness rather than what it is: a pagination-strategy cost that InnoDB pays more dearly for than PostgreSQL. Rewriting every adapter around keyset pagination (`WHERE id < cursor ORDER BY id DESC LIMIT n`) removed the confound.
- **Write-workload contamination.** The insert endpoint grows `posts` on every call; because engines and layers commit at different rates, later cells in a run would scan a differently sized table depending on run order rather than on layer identity, unless benchmark-inserted rows are deleted before every cell.

Table S30: Paired comparison of adjacently ranked access layers on the deep/nested fetch by response-time **p99** — the p99 counterpart to the throughput significance table (Table 6 in the main text), over the 25 repeated runs on each engine. Layers are ranked by median p99 (lower is better) and each adjacent pair is compared on its per-replicate p99 ratios: median p99 (ms) with a within-campaign bootstrap 95% CI, the geometric-mean paired ratio B/A (how many times the slower-tail layer’s p99 exceeds the faster’s) with a paired bootstrap 95% CI, paired dominance (fraction of replicates in which B’s p99 exceeds A’s), and the two-sided paired sign-flip permutation p (20000 permutations; 5.0×10^{-5} is the resolution floor $1/(B+1)$, and the Wilcoxon signed-rank test agrees, replication package). This delivers the per-cell p99 inference the outcomes table (Table 3 in the main text) lists as a primary analysis; because the high-load tail tracks capacity, these are resolved capacity-and-queueing differences, not sub-saturation latency.

Comparison (A < B, p99)	med. A p99 [95% CI]	med. B p99 [95% CI]	ratio B/A [95% CI]	dom.	perm. p
<i>PostgreSQL</i>					
pg < knex	20 [18–21]	27 [26–28]	1.40 [1.34–1.47]	1.00	5.0×10^{-5}
knex < drizzle	27 [26–28]	35 [32–35]	1.26 [1.21–1.31]	0.96	5.0×10^{-5}
drizzle < typeorm	35 [32–35]	50 [49–52]	1.49 [1.43–1.55]	1.00	5.0×10^{-5}
typeorm < prisma	50 [49–52]	51 [51–53]	1.03 [1.00–1.05]	0.62	0.0538
prisma < objection	51 [51–53]	57 [56–58]	1.10 [1.07–1.13]	0.92	5.0×10^{-5}
objection < sequelize	57 [56–58]	60 [59–61]	1.05 [1.03–1.08]	0.86	0.0001
sequelize < mikroorm	60 [59–61]	116 [113–119]	1.92 [1.87–1.98]	1.00	5.0×10^{-5}
<i>MySQL</i>					
mysql2 < knex	28 [28–29]	30 [29–31]	1.06 [1.03–1.08]	0.80	0.0003
knex < drizzle	30 [29–31]	47 [46–48]	1.57 [1.53–1.62]	1.00	5.0×10^{-5}
drizzle < typeorm	47 [46–48]	57 [56–59]	1.22 [1.20–1.25]	1.00	5.0×10^{-5}
typeorm < sequelize	57 [56–59]	67 [66–68]	1.17 [1.14–1.19]	0.96	5.0×10^{-5}
sequelize < objection	67 [66–68]	69 [67–71]	1.03 [1.00–1.05]	0.66	0.0799
objection < prisma	69 [67–71]	93 [89–95]	1.34 [1.30–1.39]	1.00	5.0×10^{-5}
prisma < mikroorm	93 [89–95]	120 [118–128]	1.34 [1.29–1.38]	1.00	5.0×10^{-5}

- **Query-formulation confounds.** A first attempt at the aggregation endpoint pre-aggregated *all* of `comments` before filtering to one author, so every request re-scanned the whole table regardless of which author was requested — a query-plan artifact masquerading as an access-layer effect, fixed by rewriting the query as a correlated subquery touching only the target author’s rows.

32 Statistical methods

This section details the tests summarized in the main text’s Analysis subsection. In every test the unit is the *run-level* statistic (one throughput or p99 value per repeated run, 25 per cell), never the individual request; all confidence intervals are percentile bootstraps resampling replicate indices (preserving the pairing) from a fixed seed (`mulberry32(0x57a75)`), so every resample is reproducible.

To test whether adjacently ranked layers are distinguishable, we compare each adjacent pair on the deep/nested fetch on their per-replicate throughput ratios with a two-sided paired sign-flip permutation test (20,000 permutations, so the smallest attainable p is $1/(20,001)$), cross-checked with the Wilcoxon signed-rank test, and report the geometric-mean paired ratio with a paired bootstrap 95% CI and the paired dominance (the fraction of replicates the faster layer wins). Because p99 is reported at millisecond resolution and contains ties, the permutation test keeps zero log-ratios and the Wilcoxon test drops them with the standard tie correction. Because the seven adjacent pairs follow the *observed* ranking rather than a preregistered one, their significances are a secondary, descriptive robustness check, and we apply a Bonferroni correction as a conservatism check rather than claiming a confirmatory family.

The layer×engine interaction is tested per pattern with a blocked permutation test: for each layer and replicate we form the paired PostgreSQL-minus-MySQL log-throughput difference $D_{\ell,i}$ and, under the null of no interaction, permute the layer labels of $\{D_{\ell,i}\}$ *within* each replicate i (20,000 permutations); the statistic is the between-layer F of a randomized-block ANOVA on D with replicate as the block. Rank agreement across engines (Spearman, Kendall) is reported descriptively over the seven evaluated layers — systems under study, not a sample from a population — so we attach no population confidence intervals to those correlations.

For the pg-Prisma pair we additionally report a paired two-one-sided-tests (TOST) equivalence result against a $\pm 5\%$ margin on the log-ratio. That margin is an author-selected smallest effect of practical interest for *this* study, not a general engineering threshold: at modest deployment scale a sub-5% throughput difference would not change a capacity-planning decision (for example, how many application instances to provision for a target load). We

set it from that engineering consequence rather than from measurement noise; it happens also to sit within the run-to-run and day-to-day variation typical of a virtualized host (this study’s own run-to-run CV reaches 15.9% on the noisiest MySQL-insert cell, within 4.2% on the reads), but that coincidence is not the justification. At very large scale even 5% can correspond to many machines, so the margin should be reset per deployment. We chose it before the analysis but did not preregister it, and the equivalence conclusion is not sensitive to the exact value within a few percent.

33 Measurement details

This section records the measurement specifics summarized in the main text’s Measurement procedure.

Warm-up length The fifteen-second warm-up follows a separate cold-boot measurement for a fast, a middling, and the slowest PostgreSQL layer: every layer reached and sustained at least 95% of steady-state throughput within 12 s (Prisma by 5 s, native by 9 s, MikroORM by 12 s), absorbing JIT compilation, pool fill, and plan caching.

Order invariance and completeness Cell order was reshuffled independently within each replicate; a forward-versus-reversed order check on the write endpoint (5 replicates per direction) found no detectable history effect on the read cells (reversed-to-forward ratio 0.977–1.280, median 1.012; the wide tail is confined to the noisy MySQL inserts, where five replicates cannot separate order from write noise). Of the 450 dedicated write-endpoint boots, 449 completed normally; the exception was one Drizzle/MySQL write replicate whose server boot did not clear the post-reset health check, leaving that cell with 24 of 25 runs (the runner records and skips such failures rather than retrying silently).

Tail estimation `autocannon` records a per-run latency histogram at millisecond resolution; even the slowest deep-fetch cells (≈ 480 –516 req/s over the 12 s window) complete $\approx 5,800$ –6,200 requests per run, so each run-level p99 is estimated from roughly 60 requests in its upper 1% tail, and the reported p99 is the median of the 25 run-level values.

Resource and framework sampling Server CPU and peak resident memory were sampled from `/proc` over the full process tree (parent plus descendants, so an in-process engine or spawned helper cannot escape measurement); a `perf_hooks` GC observer (event count, total pause time) and the 200 ms connection-pool sampler ran inside the server, polled through a

/stats endpoint after each run. A fixed-payload /baseline endpoint returning a seed-realistic thread-shaped payload measured the shared Express-plus-serialization floor: 10,909 (PostgreSQL server) and 10,921 (MySQL server) req/s, roughly $3\times$ the fastest deep-fetch cell, so the framework floor leaves between-layer spreads visible rather than compressed.

Insert configuration Each insert is an autocommit single-statement INSERT under each engine’s default isolation level (PostgreSQL Read Committed, MySQL Repeatable Read); the transactional variant wraps a post and its five comments in one transaction. The default durability settings are PostgreSQL fsync=on, synchronous_commit=on, full_page_writes=on and MySQL innodb_flush_log_at_trx_sync_binlog=1, log_bin=ON.

34 Artifact reproducibility summary

Table S31 summarizes what is needed to reproduce the study and which results are regenerated automatically; REPRODUCE.md in the replication package is the single runnable entry point.

35 Construct validity of the documentation-selected treatment

The *documentation-selected* rule operationalizes the eager-loading strategy a developer who follows a library’s official documentation would adopt on the deep fetch. Its construct validity rests on the selected API being the library’s *canonical* relations mechanism — the one its official documentation presents first for the pinned version (the selection rule; per-layer documentation page and version in METHODOLOGY.md and Table S16) — rather than an obscure or deprecated option. Table S33 records this per layer. Drizzle is the one borderline case: it exposes both a SQL-style query builder (the selected join) and a higher-level relational-query API that its documentation also presents, so we selected the builder join — comparable with the query-builder tier — and record the relational-query API as the documented alternative. Where the documented alternative is a byte-identical drop-in, the loading-strategy sensitivity check (Table S18) adds an empirical signal: for the join-default ORMs (Sequelize, MikroORM) the documentation-selected default is the *faster* of the two strategies, so a layer’s deep-fetch deficit does not appear to be an artifact of a poor default; for Objection the documented default (batched select-in) is the slower option on MySQL, which we disclose rather than correct; the Drizzle, Prisma, and TypeORM alternatives were not evaluated (TypeORM’s errors on the ordered nested fetch), itself a portability caveat. The construct is therefore a defined, reproducible practitioner persona — the documentation-following developer — and is *not* a claim about

Table S31: Artifact reproducibility summary: version and identifier, requirements, how to run, time and resources, and which results are reproduced automatically.

Item	Detail
Artifact & version	express-db-access-performance , release v1.12.8 (git tag v1.12.8); the archive is a git archive of the tagged commit and contains every tracked file, including the raw data and the table generators.
Persistent identifier	Zenodo 10.5281/zenodo.21497051 (concept 10.5281/zenodo.21313858).
Software	Node.js 24.18.0, PostgreSQL 18.4, MySQL 9.7.1, npm . Reference path: conda user-space (conda-forge); the Docker path is a convenience and pins older engines (PostgreSQL 16 / MySQL 8.4), so it only approximates the published numbers.
Hardware & disk	Single Linux host (full host and version capture in Table S11); ≈ 1 GB for the seeded databases.
Deterministic seed	2,000 authors / 100,000 posts / 1,000,000 comments from a fixed PRNG.
Smoke test (≈ 5 –10 min)	npm ci ; npm run migrate && npm run seed ; npm run bench:quick ; node bench/verify.mjs (read byte-equivalence, must print ALL BYTE-IDENTICAL) and node bench/verify-writes.mjs (write-state correctness, must print ALL WRITES SEMANTICALLY CORRECT).
Full campaign	Primary matrix ($\approx$ hours): INDEP=1 REPEATS=25 DURATION=12 WARMUP=15 node bench/runner.mjs $\rightarrow$ results/raw.json ; all secondary experiments add up to ≈ 25 h wall-clock.
Regenerate from raw	$\approx$ minutes, <i>no database</i> : the committed generators rebuild every table and figure from results/*.json , then npm run sync:tables and make .
Automatically reproduced	<i>Every</i> table and figure, from the archived raw data; results/checksums.sha256 verifies the 35 raw-data files. Reproducing the raw data itself requires the full campaign on the reference engines.
Entry point	REPRODUCE.md (smoke test, full matrix, clean-room reproduction from the tarball); the per-table data-and-generator map is in MANIFEST.md .

Table S32: Scope of every experiment: the access pattern(s), engine(s), number of access-layer implementations, and repeated runs each covers. The *primary comparison* (throughput and closed-loop p99) is the full five-pattern, two-engine matrix; every other experiment is a limited condition, most restricted to the deep/nested fetch, and each conclusion is scoped accordingly. “Layers” counts implementations (the nine access layers plus, where present, the two tuned native baselines); a representative subset is used where a caption says so. Engine entries name where the experiment was run (pg is PostgreSQL-only, `mysql12` MySQL-only).

Experiment	Pattern(s)	Engines	Layers	Runs
Primary matrix: throughput & closed-loop p99 (main paper)	all five	PG, MySQL	11	25 independent
Same-SQL raw-path control (Table S14)	deep fetch	PG, MySQL	9	median of 10
Open-loop CO-corrected tail (Tables S6, S17)	deep fetch	PG, MySQL	5	rate sweep
Utilization-controlled tail (Tables S21, S22)	deep fetch	PG, MySQL	8	median of 5
Equal-CPU, 1/2/4 cores (Table S4), exploratory	deep fetch	PG	3	median of 3
Node-cluster multi-worker (Table S24)	deep fetch	PG, MySQL	3	median of 3
Pool-size sweep, 1–50 (Tables S7, S25)	deep fetch	PG, MySQL	4	median of 3
Mixed read/write, 90:10 & 70:30 (Table S26)	deep fetch, insert	PG, MySQL	5	median of 5
Relaxed vs default durability (Table S3)	insert	PG, MySQL	11	25 def. + 10 relaxed
Wait-event flush decomposition (Table S19)	insert	MySQL	3	4000 inserts/layer
Transactional multi-statement write (Table S8)	POST /threads	PG	5	median of 5
Longer-window (60 s) tail (Table S23)	deep fetch	PG, MySQL	11	10 runs
Post-restart cold cache (Table S20)	deep fetch	PG, MySQL	3	vs primary median
Alternative eager-loading (Table S18)	deep fetch	PG, MySQL	3	median of 5

performance-tuned expert usage or measured production frequency, which this study does not observe; RQ1 is scoped accordingly (main text, Threats to Validity).

Table S33: Construct-validity corroboration for the documentation-selected deep-fetch treatment: the selected eager-loading API, whether it is the library’s canonical (documentation-primary) relations mechanism, and the empirical loading-strategy signal where the documented alternative is a byte-identical drop-in (Table S18).

Layer	Selected API	Canonical relations mechanism?	Loading-strategy signal
<code>pg, mysql2</code> (+tuned)	hand-written JOIN	Yes — the reference way to fetch a graph with a driver; no higher-level relations API applies	n/a (hand-written SQL)
<code>knex</code>	<code>builder .join()</code>	Yes — the query builder’s documented join	n/a (hand-written SQL)
<code>drizzle</code>	<code>builder join</code>	Documented; the relational-query API is the alternative	alternative (relational-query API) not evaluated
<code>prisma</code>	<code>include</code>	Yes — Prisma’s primary relation-query API	alternative (join strategy) not evaluated
<code>sequelize</code>	<code>include (join)</code>	Yes — Sequelize’s canonical eager loading	default <i>faster</i> than select-in
<code>typeorm</code>	<code>relations</code>	Yes — TypeORM’s documented relations option	alternative errors on the ordered nested fetch (not evaluated)
<code>objection</code>	<code>withGraphFetched</code>	Yes — Objection’s flagship eager-loading	default (select-in) slower than join on MySQL (disclosed)
<code>mikroorm</code>	<code>populate (join)</code>	Yes — MikroORM’s canonical eager loading	default <i>faster</i> than select-in

36 Results roadmap

Table S34 classifies every result in the paper by its analysis role — which outcomes carry the confirmatory, paired inference (*primary*), and which are reported descriptively as controls or sensitivity analyses (*secondary* and *exploratory*) that bound or contextualize the primary results rather than adding confirmatory hypotheses. It is the navigational companion to the reproduction map in `experiments/MANIFEST.md`, which ties each table to its generator and input data. The two co-primary deep-fetch regimes it lists under RQ1 — documentation-primary and performance-conscious — are now measured and reported in the main text’s Results section, no longer relegated to the supplement.

Table S34: Outcomes and their analysis role. The primary outcomes carry the paired inference and the research-question claims; the secondary and exploratory outcomes are reported descriptively as controls and sensitivity analyses that bound or contextualize the primary results, and are not treated as additional confirmatory hypotheses.

Role	Outcome	Analysis
Primary	Throughput (req/s), 5 patterns $\times$ 2 engines, documentation-selected layers, 50-connection operating point	Paired permutation, Wilcoxon, bootstrap median CI; spread & rank
Primary	Closed-loop p99 tail latency, same cells	Paired permutation, Wilcoxon, per-cell bootstrap CI on p99
Secondary	Same-SQL (raw-path) contrast	Bounds the residual (compound) same-SQL difference
Secondary	Sub-saturation open-loop tail (coordinated-omission corrected)	Lower-load latency companion to the high-load primary p99
Secondary	Layer $\times$ engine interaction (per pattern)	PG $\div$ MySQL effect ratios (S28), blocked permutation (F on D)
Secondary	Combined app+db CPU efficiency, equal-CPU control	End-to-end compute view; operational confirmation of the CPU trade
Explor.	Aggregation fan-out; OFFSET vs. keyset; pool-size sweep; durability & isolation sensitivity; RSS	Descriptive; pitfall illustration and robustness checks

37 Per-pattern capacity and utilization at the operating point

The operating-point protocol requires that capacity be identified so the fixed operating point can be read against it (main text, Study Design). The deep-fetch concurrency sweep (Section 25, Figure S2) does this in per-layer detail for one pattern; review point 6.5 asks whether the other four patterns, reported only at the fixed 50-connection demand, sit at comparable utilization. To answer it we swept every access layer over a connection ladder (1, 4, 8, 16, 32, 50, 100, 200) on *all five* patterns and both engines, medianed over two runs per point after a warm-up, and for each layer computed its saturating throughput (the ladder maximum), the utilization at 50 connections (throughput at 50 connections divided by that maximum), and the knee (the smallest connection count reaching 95% of the maximum).

Table S35 reports the result. The knee lies at or below 50 connections for every pattern on both engines, so the fixed 50-connection operating point places all five patterns at high utilization of their own capacity: cross-pattern comparisons of p99 and spread at that point are therefore made at comparable (high) relative utilization, not at unknown or widely different fractions of capacity. The binding constraint is the ten-connection pool shared by every pattern (main text, Controls), which is why the knee clusters well below 50 connections regardless of the pattern’s absolute throughput. This extends the capacity-identification stage of the protocol from the deep fetch to all five patterns; the equal-utilization open-loop tail (Section 20) and the exploratory equal-compute check (Section 4), which each need a per-layer saturating-throughput *target*, remain demonstrated on the deep fetch, the

pattern on which the layer spread is largest.

Table S35: Per-pattern capacity and utilization at the fixed 50-connection operating point, from a concurrency sweep (connection ladder 1–200) of every access layer on all five patterns and both engines. *Knee* is the median (across layers) smallest connection count reaching 95% of a layer’s own saturating throughput; *util. @50* is the median (and, in parentheses, minimum) across layers of the ratio (throughput at 50 connections)/(saturating throughput). The knee sits at or below 50 connections for every pattern on both engines, and the median utilization at 50 connections is 97–100% across all five patterns at high utilization of their own capacity. The capacity-identification stage of the protocol is thus demonstrated for all five patterns, not only the deep fetch; the equal-utilization and equal-compute conditions remain demonstrated on the deep fetch.

Pattern	PostgreSQL		MySQL	
	knee	util. @50 (min)	knee	util. @50 (min)
Point read	32	98% (92%)	32	100% (89%)
Range scan	16	100% (96%)	16	99% (89%)
Deep fetch	8	99% (88%)	8	97% (90%)
Aggregation	32	99% (95%)	8	98% (94%)
Insert	32	100% (97%)	16	98% (82%)

38 Paired significance of the deep-fetch ranking

The main text foregrounds paired *effect sizes* for the deep-fetch ranking — geometric-mean per-replicate throughput ratios and the fraction of replicates in which the faster layer wins — because they, not the significance tests, carry the ordering. Table S36 reports the post-hoc adjacent-pair tests as a descriptive companion: since the ranking itself selects which pairs are compared, the inference is conditional, so the permutation and Wilcoxon results are not treated as confirmatory. The bootstrap intervals are within-campaign (main text, Study Design), not general over deployment environments.

39 Application-tier CPU trade-off

Figure S3 plots each layer’s application-tier CPU against its throughput on the deep fetch (the trade-off discussed in the main text’s Resource-footprint paragraph): every layer draws about one application core, and the native driver’s throughput lead comes from issuing lean SQL that pushes work onto the database tier rather than from spending more application CPU.

Table S36: Paired comparison of adjacently ranked access layers on the deep/nested fetch (MySQL), over the 25 repeated runs. Because every layer runs the identical per-replicate request stream, layers are compared on their per-replicate throughput ratios: median throughput (req/s; the per-cell bootstrap CI is in the pattern tables), the geometric-mean paired ratio A/B with a within-campaign paired bootstrap 95% CI, paired dominance (fraction of replicates in which A exceeds B), and the two-sided paired sign-flip permutation p (20000 permutations; a value of 5.0×10^{-5} is the resolution floor $1/(B+1) = 1/20,001$, i.e. the observed statistic exceeded every permutation; the Wilcoxon signed-rank test agrees, replication package).

Comparison (A > B)	med. A	med. B	ratio A/B [95% CI]	dom.	perm. p
mysql2 > knex	2382	2007	1.19 [1.18–1.20]	1.00	5.0×10^{-5}
knex > drizzle	2007	1318	1.53 [1.51–1.55]	1.00	5.0×10^{-5}
drizzle > typeorm	1318	1017	1.30 [1.28–1.32]	1.00	5.0×10^{-5}
typeorm > sequelize	1017	886	1.15 [1.13–1.16]	1.00	5.0×10^{-5}
sequelize > objection	886	834	1.05 [1.03–1.07]	0.88	0.0001
objection > prisma	834	634	1.33 [1.31–1.36]	1.00	5.0×10^{-5}
prisma > mikroorm	634	480	1.31 [1.29–1.33]	1.00	5.0×10^{-5}




Figure S3: Throughput versus application-side CPU on the deep/nested fetch (PostgreSQL, ten-connection pool, medians of 25 replicates). Every layer sits at about 100–110% of one application core; the native driver leads on the throughput axis, and the slower ORMs are low on throughput at the same CPU. CPU is the server process only; database-side CPU is excluded for all layers alike.

40 Retrospective application of the protocol to prior benchmarks

Table S37 applies the protocol retrospectively to the eight prior benchmarks located in the sources searched. The exercise is purely analytical: we re-ran nothing and dispute none of their reported numbers. For each study we name the protocol stages that are absent and state, in one line, why the study’s headline conclusion is not interpretable as a general access-layer claim without them. The point is not that these studies erred, but that naming the missing stages localizes exactly what each study’s measurements can and cannot support — and, for the one engine-only study, why the access-layer unit of comparison is absent by design.

41 Randomized differential semantic-equivalence check

The fixed 12-probe gate (`bench/verify.mjs`) establishes finite pre-measurement conformance; the main text (Study Design) also reports a broader, property-based level. Because the benchmark dataset is deterministic and finite, testing far more inputs is inexpensive relative to the measurement campaign. A second gate (`bench/verify-property.mjs`) generates a fixed-seed sample of one thousand random post IDs and one thousand random author IDs across the full key range, together with an explicit edge set (nonexistent, boundary, negative, and $2^{31}-1$ IDs; empty and boundary keyset pages), and asserts that every adapter is byte-identical to the native-driver oracle on all five read methods. Table S38 summarizes the coverage: 3,800 distinct inputs per adapter, 30,400 adapter-versus-baseline comparisons per engine, 60,800 across both engines, with zero divergences. This does not prove equivalence on every possible input — that remains infeasible for arbitrary code — but it broadens the finite check by more than two orders of magnitude over the 12-probe gate and exercises the boundary and error-path inputs most likely to reveal an input-dependent divergence. The sample size is a script parameter, so broadening the randomized check further is a re-run at larger N against the deterministic dataset, not a redesign.

42 The five access patterns

Table S39 lists the five HTTP endpoints of the benchmark workload, the access pattern each realizes, and what each is intended to stress. The main text (Study Design) describes the patterns in prose; this table is the compact reference.

Table S37: Retrospective application of the comparability protocol to the eight prior benchmarks located in the sources searched. The exercise is analytical only: nothing was re-run and no reported figure is disputed. For each study we name the protocol stages that are absent and state why its headline conclusion is not interpretable as a general access-layer claim without them.

Study	Stages absent	Conclusion not interpretable without them
Colley et al. (2018)	1 (oracle), 3 (same-SQL control), 4 (capacity), 5 (demand/utilization)	Its cross-language ORM latency ranking cannot be read as a general access-layer ordering: with no capacity sweep (4), each latency sits at an unlocated operating point.
Procedia CS (prisma_rawsql_2025)	1 (oracle), 3 (cross-SQL control), 4 (capacity), 5 (demand/utilization)	“Optimized raw SQL beats Prisma” does not generalize to ORMs-versus-raw broadly, because only one library family (Prisma) is sampled and no capacity sweep (4) locates where the compared operating point sits relative to saturation.
JCCT (orm_compare_jccs2025)	1 (oracle), 3 (same-SQL control), 4 (capacity), 5 (demand/utilization)	Its three-ORM response-time comparison is not interpretable as a library ordering without a same-SQL control (stage 3), since nothing separates each ORM’s default loading strategy from the library machinery itself.
JCSI (zhadko2025orm)	1 (oracle), 3 (same-SQL control), 4 (capacity), 5 (client-observed tail)	The concurrency-dependent Prisma ranking flip cannot be read as client-facing: internal query stages, not client-observed request tails (5), were timed, and saturation was never located (4).
Salunke & Ouda (2024)	Not applicable: no access-layer treatment, so stages 2 and 3 have no referent	They compare the PostgreSQL and MySQL engines head-to-head, not application access layers, so the protocol’s unit of comparison is absent by design; their (valid) engine result yields no access-layer conclusion.
Drizzle (vendor)	1 (oracle), 3 (same-SQL control), 4 (capacity), 5 (matched-utilization)	Without a capacity sweep (stage 4) to locate the fixed k6 load point against each target’s saturation throughput, the Drizzle suite’s req/s and P95 ranking of the nine targets is not interpretable as a general access-layer claim.
Prisma (vendor)	1 (oracle), 3 (same-SQL control), 4 (capacity), 5 (tail/utilization)	Because the Prisma vendor benchmark reports only median query latency over 500 iterations—no latency percentile and no throughput—its “lowest median latency” finding is not interpretable as a statement about behavior under load: no tail at equal demand or matched utilization (stage 5) and no saturating-capacity identification (stage 4) are reported.
IMDBench / Gel Data (vendor)	1 (oracle), 3 (same-SQL control), 4 (capacity), 5 (matched-utilization)	Its throughput/latency ordering over three CRUD operations is uninterpretable as a general access-layer ranking because no capacity sweep (4) fixes the operating point across layers.

Table S38: Randomized differential semantic-equivalence coverage (`bench/verify-property.mjs`). A fixed-seed generator (mulberry32, seed 99005011) draws inputs across the full key range; the native driver is the oracle and every adapter response must be byte-identical to it, with any thrown error captured as a comparable value so an input-dependent crash also counts as a divergence. The gate is re-run against the seeded database (2,000 authors, 100,000 posts, 1,000,000 comments), not derived from `results/raw.json`.

Dimension	Coverage
Read methods checked	<code>getPost</code> ; deep fetch (<code>getThread</code>); same-SQL deep fetch (<code>getThreadRaw</code>); aggregation (<code>authorSummary</code>); keyset range scan (<code>listPosts</code>)
Seeded-random inputs	1,000 post IDs + 1,000 author IDs, uniform over the full range
Edge cases	9 post IDs (0, -1, 1, 2, 99,999, 100,000, 100,001, 101,000, $2^{31}-1$); 7 author IDs (0, -1, 1, 1,999, 2,000, 2,001, 3,000); 8 keyset pages (empty / boundary; limit 0, 1, 100)
Distinct inputs / adapter	3,800 (after de-duplication)
Adapters vs. baseline	8 per engine
Comparisons	30,400 per engine; 60,800 across PostgreSQL and MySQL
Divergences	0 (ALL BYTE-IDENTICAL)

Table S39: The five access patterns and what each stresses.

Endpoint	Pattern	Stresses
GET <code>/posts/:id</code>	point read (PK)	per-query overhead
GET <code>/posts?before=</code>	keyset range scan	index seek + hydration
GET <code>/posts/:id/thread</code>	deep/nested fetch	join strategy, N+1 avoidance
GET <code>/authors/:id/summary</code>	aggregation	grouped counts / raw SQL
POST <code>/posts</code>	insert	write path

43 Protocol compliance levels and their coverage

The main text (Study Design) names five protocol compliance levels; Table S40 records precisely which workload cells satisfy each in this case study. The mandatory validity core and capacity characterization hold across all five patterns on both engines; the three extensions — standardized-strategy, utilization-controlled, and resource-normalized — are demonstrated on the deep fetch, the most consequential pattern, as a representative-point instantiation rather than a claim of full-matrix coverage. Stating the levels and their coverage explicitly is what makes the protocol reusable: a later study can report its own compliance level per cell against the same ladder.

Table S40: Protocol compliance levels and which workload cells satisfy each in this case study. The mandatory validity core and capacity characterization are exercised across the *whole* matrix; the three extensions are demonstrated on the deep fetch — the most consequential pattern (main text, Results) — as a representative-point instantiation, not a claim of full-matrix coverage. This is what the protocol box means by recommended stages that “may be scoped to a representative workload point.”

Compliance level	Requirement it adds	Cells satisfying it here	Evidence
Mandatory validity core	Correctness oracle (byte-identical outputs + correct write state) and a predeclared treatment rule, on <i>every</i> timed cell	All five patterns × two engines (all 90 cells)	Study Design; Table S38; <code>verify-writes.mjs</code>
Capacity characterization	Locate saturating throughput by a concurrency sweep at each workload point	All five patterns × two engines	Table S35
Standardized-strategy extension	Add a same-SQL standardized contrast against a common-strategy baseline	Deep fetch × two engines	Main-text same-SQL table; Table S14
Utilization-controlled extension	Measure the tail at matched fractions of each treatment’s own capacity	Deep fetch × two engines	Main-text tail-regime table; Tables S21–S22
Resource-normalized extension	Compare under an equal compute budget	Deep fetch × two engines	Tables S4, S5

44 Predefined contrasts against the native-driver reference

Selecting adjacent pairs *after* seeing the ranking makes those tests conditional on the observed ordering, so their p -values add little (main text, Mea-

surement stability; the adjacent-pair table is Table S36). A predeclared contrast set avoids this. Table S41 fixes the reference (the native driver) and the contrasts (each portable layer against it) in advance, independently of the ranking, and reports for each layer on the deep fetch the geometric-mean paired ratio of native-to-layer throughput, a within-campaign 95% paired-bootstrap interval, and the paired dominance. Every portable layer trails the native driver in all 25 replicates on both engines (dominance 1.00), and the ratios range from $1.48\times$ (**knex**, PostgreSQL) to $7.16\times$ (**mikroorm**, PostgreSQL) — the same spread the main text reports, now as predeclared contrasts rather than post-hoc pairwise tests. The contrasts are within-campaign, not deployment-general.

Table S41: Predefined contrasts against the native-driver reference on the deep fetch, fixed in advance rather than selected after the ranking (reviewer 8.1). For each portable layer: the geometric-mean paired ratio of *native-driver to layer* throughput (how many times faster the native driver is, per replicate), a within-campaign 95% paired-bootstrap interval, and the paired *dominance* (fraction of the 25 replicates in which the native driver is faster). Because the reference and the contrast set are predeclared, these are not conditional on the observed ordering, unlike the adjacent-pair tests. Values are within-campaign, not deployment-general; the reference is **pg** on PostgreSQL and **mysql2** on MySQL.

Layer	PostgreSQL		MySQL	
	native/layer [95% CI]	dom.	native/layer [95% CI]	dom.
knex	1.48 [1.46,1.50]	1.00	1.19 [1.18,1.20]	1.00
drizzle	1.97 [1.93,2.00]	1.00	1.82 [1.79,1.84]	1.00
prisma	3.09 [3.04,3.14]	1.00	3.80 [3.74,3.87]	1.00
sequelize	3.71 [3.65,3.76]	1.00	2.71 [2.67,2.76]	1.00
typeorm	3.04 [2.99,3.08]	1.00	2.37 [2.33,2.40]	1.00
objection	3.56 [3.50,3.63]	1.00	2.85 [2.80,2.90]	1.00
mikroorm	7.16 [7.02,7.31]	1.00	4.97 [4.92,5.02]	1.00

45 Run-to-run p99 spread

Each cell’s reported p99 is the median of 25 run-level p99 values, and each run-level p99 is itself estimated from roughly the top 1% of that run’s requests, so a single run’s p99 is noisy. Figure S4 plots all 25 per-run p99 values behind each deep-fetch cell (PostgreSQL) with the median marked, so the run-to-run stability of the estimate is visible rather than summarized by a single number: the spread is tight for the fast layers and wider for the slow ones, consistent with the coarser p99 estimate at higher latency, and the medians preserve the reported ordering. The load generator reports latency

percentiles rather than raw per-request latencies, so a full HDR-histogram reconstruction is left to future work; the median-of-run-level-p99 estimator, its within-campaign bootstrap interval, and the longer-window (60 s) sensitivity check (Section 21) are the mitigations used here.




Figure S4: The 25 per-run p99 values behind each reported deep-fetch p99 on PostgreSQL (red bar = median, the value tabulated in the main text). Each run’s p99 is estimated from roughly the top 1% of that run’s requests, so a single run’s p99 is noisy; the paper reports the *median* of these 25 run-level p99 values with a within-campaign bootstrap interval rather than pooling dependent requests. The plot shows the run-to-run dispersion this median summarizes: it is tight for the fast layers and widens for the slower ones, consistent with the coarser p99 estimate at higher latency. A full per-request HDR recorder was not used (the load generator reports latency percentiles, not raw request latencies), so an HDR reconstruction is left to future work.

46 The same-SQL control: components of the compound intervention

The same-SQL (raw-path) control (main text, RQ1) is a *compound* intervention: replacing each layer’s documentation-selected relational-loading path with the same hand-written SQL through its raw facility changes several mechanisms at once. Table S42 enumerates them. The residual spread it leaves is a standardized-SQL contrast — not a bound on the combined effect, and not attributable to any single component, in particular not to SQL formulation or query count.

Table S42: The *same-SQL control* is a compound intervention: switching each layer from its documentation-selected path to the same hand-written SQL through its raw facility moves every component below at once — equalizing the SQL text, query strategy, round-trip count, and result mapping, while the raw-SQL facility, statement preparation, and disclosed wire protocol still differ across layers. The residual difference it leaves is a standardized-SQL contrast, not a bound on their combined effect, and cannot be attributed to any single component — in particular, not to SQL formulation or query count.

Component	Documentation-selected path	Same-SQL (raw-path) control
Loading API	the ORM’s eager-load API (<code>include</code> , <code>populate</code> , ...) or a builder join	the layer’s raw-SQL facility
Query strategy	per layer: single join, batched select-in, or multi-query	one hand-written two-statement join, identical across layers
SQL formulation	per-ORM generated statement text	byte-identical statement text (server-side confirmed)
Round-trip count	one to four statements	two statements
Statement preparation	ORM-managed	the raw facility’s default
Wire protocol	per driver (extended/simple, text/binary)	per driver, unchanged and disclosed
Result hydration	the ORM’s object-graph mapping	shared canonical constructors, identical across layers

47 Deep-fetch tail: equal closed-loop demand and matched utilization

The high-load p99 under equal closed-loop demand (main text, Tail latency) and the p99 at matched fractions of each layer’s own capacity read the deep-fetch tail two ways. Table S43 shows both per layer, so the large high-load gap (pg 20 to mikroorm 116 ms on PostgreSQL) is seen to be a capacity-and-queueing effect that largely dissolves to a 2–5 ms band at matched utilization, rather than a sub-saturation latency difference.

Table S43: The deep-fetch tail read two ways, per layer and engine. *Equal demand* is the high-load response-time p99 (ms) at the fixed 50-connection closed-loop point — the headline deep-fetch p99, which at this saturating load largely tracks capacity and connection-acquisition queueing. *50% util.* is the coordinated-omission-corrected p99 when each layer is instead offered half of *its own* saturating throughput (Supplement Tables S21–S22). The wide equal-demand ladder (pg 20 to mikroorm 116 ms on PostgreSQL) collapses to a 2–5 ms band at matched utilization, so the high-load gap is a capacity-and-queueing effect, *not* a sub-saturation latency difference; the matched-utilization reading is the more informative one.

Layer	Category	PostgreSQL p99 (ms)		MySQL p99 (ms)	
		equal demand	50% util.	equal demand	50% util.
pg	native-driver	20	3.9	–	–
mysql2	native-driver	–	–	28	4.2
knex	query-builder	27	2.4	30	1.9
drizzle	orm	35	3.8	47	3.7
prisma	orm	51	4.2	93	5.2
sequelize	orm	60	2.5	67	2.5
typeorm	orm	50	3.1	57	2.8
objection	orm	57	3.1	69	2.9
mikroorm	orm	116	3.9	120	4