Online supplement: A Reproducible Benchmark of Relational Database Access-Layer Performance in Express.js: A Configuration-Specific Comparison on PostgreSQL and MySQL

Mateusz Miotk
Faculty of Mathematics, Physics and Informatics, University of Gdańsk
mateusz.miotk@ug.edu.pl

This supplement carries measurement detail referenced from the main text. All tables are generated by the replication package (https://github.com/mmiotk/express-db-access-performance; archived at Zenodo as release v1.6.1, https://doi.org/10.5281/zenodo.21428215) from the same raw data as the main-text tables.

1 Dispersion on MySQL

The main text reports the coefficient-of-variation table for PostgreSQL and the per-engine maxima. Table S1 gives the full per-layer, per-pattern CV on MySQL across the 25 repeated runs (maximum 15.9%, `drizzle`/insert; the insert cells are the noisiest in the study). Figure S1 plots the raw per-replicate insert throughput behind these coefficients: several cells are visibly bimodal or heavy-tailed within a cell, structure a single CV cannot convey, which is why the insert rank correlation and dispersion are reported conservatively.

2 Idiomatic round-trip counts

Table S2 reports, for each layer, the number of SQL statements the idiomatic deep fetch issues per request, captured from server-side statement logs. The counts motivate the main text's observation that round-trip count alone does not determine throughput: the two layers issuing four statements (Prisma, Objection) sit at opposite ends of the throughput ladder.

Table S1: Coefficient of variation (%) of throughput across the 25 repeated runs, per access layer and pattern on MySQL. Low values indicate stable measurements.

Layer	**point_read**	**range_scan**	**deep_fetch**	**aggregation**	**write**
mysql2	2.4	2.6	2.4	2.5	10.5
knex	3.7	2.7	2.7	1.9	13.2
drizzle	2.9	2.2	2.0	3.5	15.9
prisma	1.9	2.6	3.3	1.9	7.1
sequelize	2.1	1.8	2.5	1.6	12.5
typeorm	3.6	4.0	3.2	1.8	8.5
objection	2.5	1.9	3.8	2.9	11.7
mikroorm	3.0	4.0	2.2	2.1	5.6

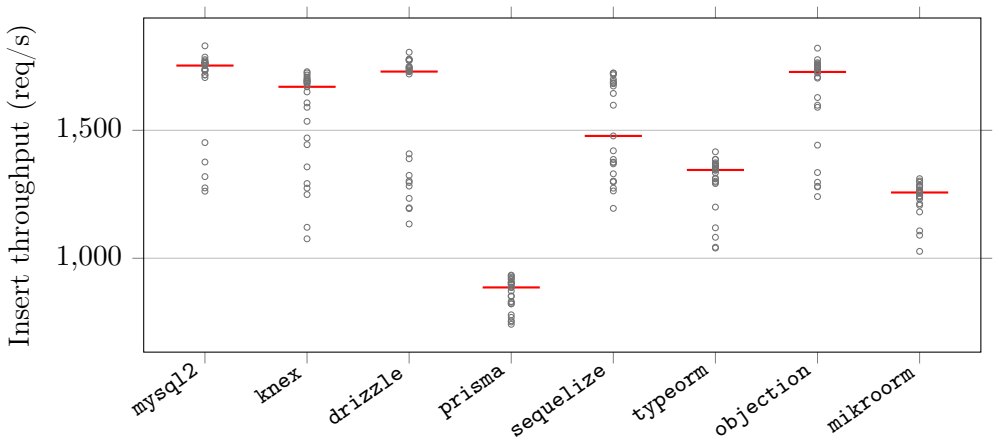

Figure S1: Raw per-replicate insert throughput on MySQL for the eight idiomatic layers (25 independent runs each, red bar = median). The insert cells are the noisiest in the study; several are visibly bimodal or heavy-tailed within a cell, structure a coefficient of variation cannot convey, which is why the insert rank correlation and dispersion are reported conservatively.

Table S2: Number of SQL round trips each access layer issues per request on PostgreSQL, captured from server-side statement logging. The native drivers, Knex, and Drizzle issue a two-query deep fetch; the data-mapper ORMs' eager-loading strategies choose their own counts. Full generated SQL and `EXPLAIN` plans are in the replication package.

Layer	Point read	Range scan	Deep fetch	Aggregation	Insert
`pg`	1	1	2	1	1
`knex`	1	1	2	1	1
`drizzle`	1	1	2	1	1
`prisma`	1	1	4	1	1
`sequelize`	1	1	1	1	1
`typeorm`	1	1	2	1	2
`objection`	1	1	4	1	1
`mikroorm`	1	1	1	1	1

3 Write throughput under default versus relaxed durability

Table S3 gives the per-layer insert throughput under each engine's default durability configuration (the paper's primary write regime) and under the symmetric relaxed regime (secondary, 10 replicates), with the relaxed÷default ratio quoted in the main text.

4 Equal-CPU-budget control

Table S4 gives the deep-fetch throughput with the server confined to 1, 2, or 4 cores (database and load generator pinned to disjoint cores); all nine values are quoted in the main text.

5 Application-tier versus end-to-end CPU efficiency

Table S5 reports, per layer on the deep/nested fetch, the median application-process-tree CPU and database-server CPU, the combined core-count, and throughput per application-CPU-second and per combined (application+database) CPU-second. The application-only figure understates the compute of layers that delegate more work to the database, but on the current versions the application-only and combined figures largely agree, because no layer draws more than about one application core: `mysql2` leads application-tier efficiency by about 3.8× over Prisma (2,269 vs. 604 req per app-CPU-s) and by 3.3× end-to-end (1,254 vs. 384 req per combined-CPU-s). There is no CPU-for-throughput trade of the kind an earlier version of this study reported.

Table S3: Insert throughput (req/s) under the default durability configuration (PostgreSQL `fsync=on`, `synchronous_commit=on`; MySQL `innodb_flush_log_at_trx_commit=1`, `sync_binlog=1`; median of 25 replicates) and under the relaxed regime (all of the above off; median of 10), with the relaxed÷default ratio. At 50 concurrent connections, group commit amortizes the per-commit flushes: relaxing durability buys up to 51% on PostgreSQL (the write-bound MikroORM; less for the others) and 2.1–2.4× on MySQL.

	PostgreSQL			MySQL		
Layer	default	relaxed	ratio	default	relaxed	ratio
`pg`	6402	7228	1.13	–	–	–
`mysql2`	–	–	–	1753	4114	2.35
`pg-tuned`	6413	7281	1.14	–	–	–
`mysql2-tuned`	–	–	–	1750	3980	2.27
`knex`	4841	5425	1.12	1670	3528	2.11
`drizzle`	4085	4557	1.12	1730	3294	1.90
`prisma`	2774	2833	1.02	886	1014	1.14
`sequelize`	2732	2727	1.00	1478	2341	1.58
`typeorm`	2648	2819	1.06	1345	1675	1.25
`objection`	3907	4390	1.12	1728	3295	1.91
`mikroorm`	1979	2981	1.51	1257	3009	2.39

Table S4: Deep-fetch throughput (req/s, PostgreSQL, median of 3) with the server process confined by `taskset` to 1, 2, or 4 cores (database and load generator pinned to disjoint cores). `pg` reaches full throughput on a single core; Prisma requires roughly four to approach it, quantifying the CPU subsidy behind its near-native throughput.

Layer	1 core	2 cores	4 cores
`pg`	3689	3688	3687
`prisma`	1079	1052	1219
`mikroorm`	446	495	522

Table S5: Application-tier versus end-to-end CPU efficiency on the deep/nested fetch (MySQL): median CPU (% of one logical core) of the Node.js process tree and of the database server, the combined core-count, and throughput per application-CPU-second and per combined (app+database) CPU-second. The application-only figure understates the compute of layers that delegate more work to the database; the combined figure is the end-to-end view. Load-generator CPU is excluded.

Layer	app CPU%	db CPU%	comb. cores	req/app-CPU-s	req/comb-CPU-s
mysql2	105	85	1.90	2269	1254
knex	100	75	1.75	2007	1147
drizzle	105	50	1.55	1255	850
prisma	105	60	1.65	604	384
sequelize	100	30	1.30	886	682
typeorm	100	70	1.70	1017	598
objection	100	45	1.45	834	575
mikroorm	110	25	1.35	436	356

6 Coordinated-omission-corrected open-loop sweep

Table S6 gives the full constant-arrival sweep (PostgreSQL, five layers, offered rates 250–4,000 requests/s) behind the open-loop paragraph of the main text: achieved throughput, corrected latency percentiles measured from each request's *intended* send time, and timeout counts (10 s cap, clipped into the distribution).

Table S6: Open-loop tail latency, corrected for coordinated omission: p99 (ms) on the deep/nested fetch (PostgreSQL) under a constant offered request rate, with every latency measured from the request's intended start time and timed-out requests included at their clipped value. †marks cells that did not sustain the offered rate (achieved<95% or any timeouts).

Layer	250/s	500/s	1000/s	2000/s	4000/s
`pg`	2	2	2	3	1483†
`prisma`	3	3	13	10721†	10947†
`knex`	2	2	2	89	8812†
`sequelize`	3	3	183	10745†	10940†
`mikroorm`	4	346	11283†	11677†	11935†

7 Pool-size sensitivity

The primary campaign fixes every layer's connection pool at 10 to isolate the access layer from pool tuning. Table S7 varies the pool from 1 to 50

for a representative layer of each tier on the deep fetch (PostgreSQL, 50 connections). The deep-fetch ordering (`pg` > Knex > Prisma > MikroORM) is preserved at every pool size, and each layer's throughput saturates at a pool well below 50, so the fixed-pool choice does not drive the ranking; a very small pool (1–2) is the one regime that compresses the differences, because every layer then queues on connection acquisition.

Table S7: Deep-fetch throughput (req/s, PostgreSQL, 50 connections, median of 3) as the connection pool varies from 1 to 50, per access layer. The primary campaign fixes the pool at 10; each layer saturates at a pool well below 50 and the ordering is preserved at every pool size, so the fixed-pool choice does not drive the ranking.

Layer	1	2	5	10	20	50
`pg`	905	1762	3224	3667	3798	3768
`knex`	741	1413	2392	2532	2683	2636
`prisma`	870	1072	1183	1234	1207	1105
`mikroorm`	494	487	507	527	527	522

8 Transactional multi-statement write

Beyond the single-row insert of the primary matrix, Table S8 reports a transactional write, `POST /threads`, which inserts a post and five comments in one transaction through each layer's transaction facility, for a representative layer of each tier on PostgreSQL under default durability. The full layer × engine matrix for this pattern is left to future work. Because the transaction commits once, its per-commit fsync dominates and the layers cluster more tightly than the widest read spread (3.7×, `pg` 2,718 down to Prisma 733), though the ordering does not simply track the single-row insert: Prisma, mid-pack on the single-row PostgreSQL insert, falls to the bottom of the transactional subset.

9 Dispersion on PostgreSQL

Table S9 gives the full per-layer, per-pattern coefficient of variation of throughput on PostgreSQL across the 25 repeated runs (maximum 7.6%, `pg`/insert), the companion to the MySQL table in Section S1.

10 Resource footprint

Table S10 gives the per-layer application-process CPU (median % of one logical core) and peak resident memory on the deep/nested fetch (MySQL),

Table S8: Transactional write throughput (req/s) and tail latency (p99, ms) for `POST /threads`, which inserts a post and five comments in a single transaction, on PostgreSQL under default durability (50 connections, median of 5). A representative layer of each tier is shown; the full matrix is left to future work. Between-layer spread is 3.7× (`pg` 2718 down to Prisma 733, now the slowest), well below the 7.15× read spread, and the ordering tracks the single-row insert, so the write conclusions extend to a multi-statement transactional write.

Layer	req/s	p99 (ms)
`pg`	2718	22
`knex`	2322	27
`prisma`	733	80
`typeorm`	1935	32
`mikroorm`	856	68

Table S9: Coefficient of variation (%) of throughput across the 25 repeated runs, per access layer and pattern on PostgreSQL. Low values indicate stable measurements.

Layer	**point_read**	**range_scan**	**deep_fetch**	**aggregation**	**write**
pg	4.0	2.3	3.4	2.0	5.3
knex	4.1	3.0	2.7	3.6	7.3
drizzle	2.8	2.5	2.0	3.4	4.0
prisma	3.0	2.5	3.0	2.7	3.8
sequelize	2.6	3.0	2.2	2.2	2.6
typeorm	2.3	2.1	2.6	3.9	3.6
objection	2.8	2.4	2.6	3.3	5.0
mikroorm	2.4	3.8	4.0	4.2	2.5

the detail behind the CPU trade-off discussed in the main text (all layers ≈ 100–110% of one core; peak RSS 215–356 MB).

Table S10: Server-process resource footprint on the deep/nested fetch (MySQL): median CPU utilization (% of one core) and peak resident memory (MB) of the Node.js process under load. The load generator and database run in separate processes.

Layer	Category	CPU (%)	RSS (MB)
mysql2	native-driver	105	294
knex	query-builder	100	215
drizzle	orm-lightweight	105	297
prisma	orm	105	334
sequelize	orm	100	262
typeorm	orm	100	225
objection	orm	100	221
mikroorm	orm	110	356

11 Pinned versions

Table S11 lists the pinned versions of every evaluated access layer and tool, recorded from the lockfile on the measurement date.

Table S11: Pinned versions of the evaluated access layers and tooling.

Layer / tool	Version	Layer / tool	Version
`pg`	8.22.0	`sequelize`	6.37.8
`mysql2`	3.23.0	`typeorm`	1.1.0
`knex`	3.3.0	`objection`	3.1.5
`drizzle-orm`	0.45.2	`@mikro-orm/core`	7.1.6
`@prisma/client`	7.8.0	`express`	5.2.1
`autocannon`	8.0.0	Node.js	24.18.0
PostgreSQL	18.4	MySQL	9.7.1

Pinned to the latest stable release compatible with the harness adapter contract at the freeze date (2026-07-15); Sequelize's version-7 line is a prerelease, so the 6.x stable line is used. Prisma 7 is Rust-free and connects through an official JS driver adapter (`@prisma/adapter-pg` for PostgreSQL, `@prisma/adapter-mariadb` with the `mariadb` 3.5.3 driver for MySQL, as Prisma ships no `mysql2` adapter).

12 Per-pattern detail: the read patterns

The two cheapest read patterns show the smallest between-layer spread and are summarized in the main text; their full throughput and p99 tables are given here. Table S12 reports the point read (primary-key lookup) and Table S13 the keyset range scan, each by access layer and engine. The consequential patterns (deep/nested fetch, aggregation, insert) are tabulated in the main text.

Table S12: Point read: throughput (req/s, higher is better; median with a within-campaign 95% bootstrap interval over the 25 repeated runs — run-to-run variability within one host and campaign, not across hardware or days) and response-time p99 (ms, lower is better) by access layer and engine.

		PostgreSQL		MySQL	
Layer	Category	req/s [95% CI]	p99	req/s [95% CI]	p99
pg	native-driver	6835 [6597–6892]	9	–	–
mysql2	native-driver	–	–	4368 [4314–4420]	16
pg-tuned	native-tuned	6806 [6622–6964]	9	–	–
mysql2-tuned	native-tuned	–	–	4268 [4212–4331]	14
knex	query-builder	4942 [4794–5131]	13	3814 [3771–3870]	15
drizzle	orm-lightweight	4199 [4161–4254]	14	2928 [2866–2977]	21
prisma	orm	3255 [3222–3310]	20	2074 [2050–2095]	29
sequelize	orm	3518 [3483–3591]	17	2764 [2726–2798]	20
typeorm	orm	4185 [4171–4242]	15	3059 [3032–3097]	19
objection	orm	3977 [3943–4040]	15	3275 [3221–3310]	17
mikroorm	orm	2455 [2425–2484]	24	2067 [2044–2086]	27

13 Same-SQL control: full table

Table S14 gives the full same-SQL (raw-path) control on the deep fetch for both engines: throughput when every layer executes the identical control statement (verified byte-identical by server-side capture). The main text reports the resulting collapse toward parity and reads it as a bound on the residual same-SQL difference, a compound intervention that isolates no single factor.

14 Per-pattern detail: aggregation

Table S15 gives the full aggregation-pattern throughput and p99 by access layer and engine. RQ3 finds aggregation the least layer-sensitive pattern once every layer runs the same raw correlated-subquery plan; the main text reports the ranking and the leading figures.

Table S13: Keyset range scan: throughput (req/s, higher is better; median with a within-campaign 95% bootstrap interval over the 25 repeated runs — run-to-run variability within one host and campaign, not across hardware or days) and response-time p99 (ms, lower is better) by access layer and engine.

		PostgreSQL		MySQL	
Layer	Category	req/s [95% CI]	p99	req/s [95% CI]	p99
pg	native-driver	3811 [3747–3865]	18	–	–
mysql2	native-driver	–	–	3295 [3274–3341]	21
pg-tuned	native-tuned	3830 [3766–3872]	18	–	–
mysql2-tuned	native-tuned	–	–	3252 [3233–3271]	19
knex	query-builder	3220 [3198–3240]	18	2995 [2942–3022]	20
drizzle	orm-lightweight	2780 [2749–2822]	20	2165 [2145–2192]	29
prisma	orm	2443 [2428–2473]	26	1538 [1517–1556]	39
sequelize	orm	2478 [2424–2498]	23	2105 [2090–2117]	27
typeorm	orm	2564 [2537–2590]	24	2262 [2150–2298]	26
objection	orm	2686 [2659–2697]	22	2582 [2551–2597]	23
mikroorm	orm	905 [891–915]	64	856 [846–865]	66

Table S14: Bounding the residual same-SQL difference on the deep-fetch spread: throughput (req/s) of the idiomatic deep fetch versus the *same-SQL* control (median of 10 runs), in which every layer executes the identical two-statement SQL with identical row mapping through its raw-SQL facility. The ratio (idiomatic ÷ same-SQL) measures the gap between each layer's idiomatic relational-fetch path and the common raw-SQL path, a composite of query strategy, query count, protocol, the raw versus ORM API, and result marshalling and hydration changed together; the residual native-relative same-SQL spread (1.68× on PostgreSQL, 2.01× on MySQL) bounds that compound difference without isolating any single factor.

	PostgreSQL			MySQL		
Layer	idiom.	same-SQL	ratio	idiom.	same-SQL	ratio
`pg`	3687	3625	1.02	–	–	–
`mysql2`	–	–	–	2382	2427	0.98
`knex`	2487	2926	0.85	2007	2245	0.89
`drizzle`	1865	3237	0.58	1318	2008	0.66
`prisma`	1186	2155	0.55	634	1206	0.53
`sequelize`	991	2436	0.41	886	1903	0.47
`typeorm`	1204	3170	0.38	1017	2267	0.45
`objection`	1028	2831	0.36	834	2310	0.36
`mikroorm`	516	3239	0.16	480	2453	0.20

Table S15: Aggregation: throughput (req/s, higher is better; median with a within-campaign 95% bootstrap interval over the 25 repeated runs — run-to-run variability within one host and campaign, not across hardware or days) and response-time p99 (ms, lower is better) by access layer and engine.

		PostgreSQL		MySQL	
Layer	Category	req/s [95% CI]	p99	req/s [95% CI]	p99
pg	native-driver	6978 [6865–7038]	11	–	–
mysql2	native-driver	–	–	3994 [3941–4018]	18
pg-tuned	native-tuned	7538 [7493–7705]	10	–	–
mysql2-tuned	native-tuned	–	–	3958 [3933–3985]	15
knex	query-builder	5666 [5608–5710]	13	3842 [3815–3883]	19
drizzle	orm-lightweight	6280 [6156–6427]	12	3520 [3447–3541]	21
prisma	orm	4408 [4368–4476]	16	2328 [2301–2349]	30
sequelize	orm	4813 [4758–4870]	14	3216 [3203–3235]	20
typeorm	orm	6098 [5985–6179]	13	3767 [3731–3812]	18
objection	orm	3812 [3724–3851]	17	2763 [2743–2785]	25
mikroorm	orm	5745 [5637–5805]	13	3741 [3711–3804]	19

15 Idiomatic deep-fetch choices per access layer

Table S16 documents, for each access layer, the exact eager-loading API used on the deep/nested fetch, the number of SQL round trips it issues, the documented alternative loading strategy we did *not* use, and the query-logging state. Every layer uses its documented default (idiomatic) path; query logging is disabled and verified on all eleven adapters, and no lifecycle hooks or validation run on the read path. The six data-mapper and ORM layers split both ways on the join-versus-select-in default: Sequelize, TypeORM, and MikroORM default to a single join, while Prisma and Objection default to separate batched queries. The loading-strategy sensitivity check reported in the main text (Table S18) switches the join-default layers to select-in where that alternative is a byte-identical drop-in. Full per-adapter detail is in the replication package's `METHODOLOGY.md`.

16 Open-loop tail latency on MySQL

Table S17 is the MySQL companion to the PostgreSQL open-loop sweep (Supplement Table S6): coordinated-omission-corrected p99 on the deep/nested fetch at a constant offered rate. The sub-saturation tail is flat and small on both engines, and each layer collapses once the offered rate crosses its capacity, so the saturated-tail ordering holds on MySQL as on PostgreSQL.

Table S16: Idiomatic deep-fetch implementation per access layer: the eager-loading API actually used, the SQL round trips it issues (measured from server-side statement logging for the portable layers; by construction for the native variants), the documented alternative loading strategy we did *not* use, and the query-logging state. Query logging is disabled everywhere; no lifecycle hooks or validation run on the read path (details in the replication package's `METHODOLOGY.md`).

Layer	Deep-fetch API	Round-trips	Alternative not used	Logging
`pg`	hand-written JOIN ×2	2	n/a (hand-written SQL)	off
`mysql2`	hand-written JOIN ×2	2	n/a (hand-written SQL)	off
`pg-tuned`	named-prepared JOIN ×2	2	n/a (hand-written SQL)	off
`mysql2-tuned`	`execute()` JOIN ×2	2	n/a (hand-written SQL)	off
`knex`	builder `.join()` ×2	2	n/a (hand-written SQL)	off
`drizzle`	builder `.innerJoin()` ×2	2	relational query API (`db.query`)	off
`prisma`	`include:`	4	`relationLoadStrategy: 'join'`	off
`sequelize`	`include:` (`separate:false`)	1	`separate:true` (select-in)	off
`typeorm`	`relations:`	2	`relationLoadStrategy: 'query'`	off
`objection`	`.withGraphFetched()`	4	`.withGraphJoined()`	off
`mikroorm`	`populate:`	1	`strategy: 'select-in'`	off

Table S17: Open-loop tail latency on **MySQL**, corrected for coordinated omission: p99 (ms) on the deep/nested fetch under a constant offered request rate, every latency measured from the intended start time and timed-out requests clipped into the distribution. †marks cells that did not sustain the offered rate. As on PostgreSQL (Supplement Table S6), the sub-saturation tail is flat and small and each layer collapses once the offered rate crosses its capacity, so the saturated-tail ordering is engine-robust.

Layer	250/s	500/s	1000/s	2000/s	4000/s
`mysql2`	2	2	3	15	10856†
`prisma`	5	7	9518†	11723†	11900†
`knex`	2	2	2	243	10904†
`sequelize`	3	3	1628†	10797†	10962†
`mikroorm`	4	953†	11314†	11788†	11912†

17 Alternative eager-loading sensitivity

Table S18 reports the loading-strategy sensitivity check: for the data-mapper ORMs whose alternative strategy is a byte-identical drop-in, idiomatic deep-fetch throughput versus the alternative, on both engines. Both endpoints were verified to return byte-identical responses before timing.

Table S18: Alternative eager-loading sensitivity on the deep/nested fetch: idiomatic throughput (req/s) versus the alternative loading strategy, both engines, for the layers whose alternative is a byte-identical drop-in. Sequelize and MikroORM switch from their idiomatic single join to select-in; in both cases the idiomatic default is the *faster* strategy, so their deep-fetch deficit is not an artifact of a poor default. `TypeORM` is omitted, its alternative strategy not a byte-identical drop-in in the pinned versions (its query strategy errors on the ordered nested `findOne`, and Objection's `withGraphJoined` changes row handling), which is itself a portability caveat.

		PostgreSQL		MySQL	
Layer	Switch	idiom.	alt.	idiom.	alt.
`sequelize`	join→select-in	915	690	830	602
`mikroorm`	join→select-in	447	357	420	324
`objection`	select-in→join	–	–	820	1301

18 MySQL insert commit-flush mechanism

Table S19 gives the direct `performance_schema` evidence behind the MySQL insert bottleneck: sampling the redo-log and binary-log flush wait instruments around a fixed insert workload at default durability, the flush is a shared engine floor—similar per-insert across the native driver, query builder, and ORM despite their very different read performance—and relaxing durability removes it.

19 Post-restart environmental-robustness check

Table S20 reports the environmental-robustness check: after a full restart of both database engines, a representative deep-fetch subset re-measured eight days after the primary campaign. The layer ranking reproduces on both engines (`pg` > Prisma ≫ MikroORM on PostgreSQL; Prisma > `mysql2` ≫ MikroORM on MySQL); absolute throughput is 3–23% lower, the day-to-day host drift the paper's relative analysis is designed to absorb. Because PostgreSQL and MySQL cells are not measured simultaneously, this drift is non-uniform across engines (3% on `mysql2`, 13% on `pg`, 23% on MikroORM/PostgreSQL),

Table S19: MySQL insert commit-flush mechanism (`performance_schema` wait instruments, default durability `innodb_flush_log_at_trx_commit=1`, `sync_binlog=1`): per-insert wall time, the redo-log+binlog flush wait per insert, and the flush share of wall time, over 4000 inserts per layer. The flush wait is a shared engine floor – similar across layers despite their very different read performance – so it caps insert throughput regardless of the access layer. Relaxing durability (`mysql2`, `trx_commit=0`, `sync_binlog=0`) removes the flush and lifts throughput to 3238 req/s, confirming the floor.

Layer (MySQL)	req/s	per-insert (ms)	flush/insert (ms)	flush share
`mysql2`	1576	0.634	1.001	158%
`knex`	1722	0.581	1.009	174%
`mikroorm`	2584	0.387	0	0%
Relaxed durability (control):				
`mysql2`*	3238	0.309	0.01	3%

so the cross-engine interaction ratios could in principle carry some residual drift; they are far larger than 23% (reads 1.0–1.6, writes 2.4–7.7), so the engine-dependence conclusion holds.

Table S20: Environmental-robustness check (review 6.7): after a full restart of both database engines (cold buffer pools, fresh connections), deep-fetch throughput of a representative subset re-measured eight days after the primary campaign, against the primary median. The layer ranking reproduces on both engines; absolute throughput drifts with host conditions (larger on PostgreSQL), which is exactly the day-to-day variation the paper's *relative* analysis is designed to absorb. Single host, so this checks temporal robustness, not cross-machine generalization.

Engine	Layer	primary req/s	post-restart req/s	ratio
PostgreSQL	`pg`	3687	3100.5	0.84
PostgreSQL	`prisma`	1186	1153	0.97
PostgreSQL	`mikroorm`	516	434.5	0.84
MySQL	`mysql2`	2382	2375.5	1.00
MySQL	`prisma`	634	618.5	0.98
MySQL	`mikroorm`	480	437.5	0.91

20 Utilization-controlled tail latency

Tables S21 and S22 give the coordinated-omission-corrected p99 on the deep fetch when each layer is offered 50/70/85/95% of *its own* saturating through-

put (its primary closed-loop deep-fetch rate), replicated five times, on both engines. At matched utilization the tails are small and similar across layers—for example on PostgreSQL every layer holds p99 near 2–4 ms at 50%—so the wide saturated-p99 ladder in the main text is a capacity-and-queueing effect, not a difference in tail latency at matched utilization; the tail rises only as a layer nears its own capacity. This is the capacity-normalized latency companion to the saturated primary p99. The 50/70% cells, which carry that comparison, are stable across runs; the 85/95% cells are increasingly noisy (p99 varies several-fold between runs as a layer approaches saturation), so those columns are read as trend, not precise values.

Table S21: Utilization-controlled open-loop tail on the deep fetch (PostgreSQL): coordinated-omission-corrected p99 (ms, median of five runs) when each layer is offered a fixed fraction of *its own* saturating throughput. At matched utilization the tails are small and similar across layers—the large saturated p99 gap is a capacity/queueing effect, not intrinsic tail latency; the tail rises only as each layer approaches its own capacity.

Layer	50%	70%	85%	95%
`pg`	3.9	3.1	4	5.9
`knex`	2.4	3.2	3.5	54
`drizzle`	3.8	8.8	30.9	7.4
`prisma`	4.2	7.8	18.8	77.5
`sequelize`	2.5	4.3	7.7	11.8
`typeorm`	3.1	3.7	7.9	37.1
`objection`	3.1	3.1	6.2	284.6
`mikroorm`	3.9	4.9	5.5	7.5

21 Longer-window tail sensitivity

The primary p99 is measured over a 12 s window, which yields only ≈ 60 requests above the p99 threshold in the slowest deep-fetch cells. To check that this per-run tail count is sufficient, Table S23 re-measures the deep fetch over a 60 s window (about five times the requests, ≈ 300 tail observations per run) at the same 50-connection operating point, ten independent runs, on both engines. The p99 ranking is preserved exactly (Spearman $\rho = 1.00$ per engine), each cell's p99 barely moves from the 12 s value, the p97.5 ordering matches the p99 ordering ($\rho = 1.00$), and the run-to-run p99 coefficient of variation stays small (maximum 6.0% on PostgreSQL, 3.8% on MySQL). The shorter window is therefore adequate for the tail *ranking*, and the tail conclusions do not rest on the exact per-run tail count.

Table S22: Utilization-controlled open-loop tail on the deep fetch (MySQL): coordinated-omission-corrected p99 (ms, median of five runs) when each layer is offered a fixed fraction of *its own* saturating throughput. At matched utilization the tails are small and similar across layers—the large saturated p99 gap is a capacity/queueing effect, not intrinsic tail latency; the tail rises only as each layer approaches its own capacity.

Layer	50%	70%	85%	95%
`mysql2`	4.2	7.7	20.9	39.7
`knex`	1.9	2.7	13.7	4.1
`drizzle`	3.7	6.1	14.9	17.6
`prisma`	5.2	5.6	11.9	51.4
`sequelize`	2.5	3.2	6.8	21
`typeorm`	2.8	3.6	3	36.5
`objection`	2.9	3.2	3.7	32.9
`mikroorm`	4	4.2	8.7	31.2

22 Multi-worker scaling

Table S24 scales each of a representative subset to 1, 2, and 4 Express workers (a shared-port `node` cluster). Because every layer is single-core-bound per process, throughput scales roughly linearly with workers until the host saturates: on PostgreSQL the native driver leads at one worker (`pg` 3,299 to Prisma's 1,117) and saturates near two workers as it becomes database-bound (4,055, then 3,961 at four), while Prisma, low per process, rises 1,117 to 2,275 to 4,449 and reaches native-like aggregate throughput at four workers. A layer's low per-process throughput is therefore recoverable by horizontal scaling, at a proportional cost in cores—the direct test of the single-host resource-competition concern.

23 Pool-size frontier on MySQL

Table S25 is the MySQL companion to the PostgreSQL pool sweep (Supplement Table S7): deep-fetch throughput as the connection pool varies from 1 to 50, per layer. Each layer saturates at a pool well below 50 and the ordering is preserved at every pool size, so the fixed pool of ten does not drive the ranking on either engine.

24 Mixed read/write workload

Table S26 interleaves the deep-fetch read with single-row inserts at 90:10 and 70:30 read:write, both engines, with a physical write reset per run. The layer

Table S23: Longer-window tail sensitivity on the deep/nested fetch. Each cell's primary p99 is measured over a 12 s window (≈60 tail observations per run in the slowest cells); the re-measurement uses a 60 s window (≈300 tail observations) at the same 50-connection operating point, 10 independent runs. The p99 ranking is preserved exactly (Spearman $\rho = 1.00$ on both engines), the per-cell p99 barely moves, the p97.5 ordering matches the p99 ordering ($\rho = 1.00$), and the run-to-run p99 CV stays small (max 6.0% PostgreSQL, 3.8% MySQL), so the 12 s window is adequate for the tail ranking despite the modest per-run tail count.

Layer	p99 12 s (ms)	p99 60 s (ms)	p97.5 60 s (ms)	p99 CV (%)	rps 60 s
PostgreSQL					
`pg-tuned`	18	18	17	2.3	3785.5
`pg`	20	20.5	18.5	6.0	3687.5
`knex`	27	28	26	2.9	2460
`drizzle`	35	35	33	2.3	1855.5
`typeorm`	50	51	49	3.7	1222
`prisma`	51	52.5	50	3.1	1176.5
`objection`	57	59	56.5	2.9	1026
`sequelize`	60	60	57	1.5	994
`mikroorm`	116	112.5	109	2.8	519.5
MySQL					
`mysql2-tuned`	25	25.5	24	3.3	2409.5
`mysql2`	28	28	25.5	2.6	2443.5
`knex`	30	29.5	28	3.3	2048
`drizzle`	47	48	45	3.2	1297
`typeorm`	57	58.5	56	1.8	1015.5
`sequelize`	67	66	63	2.2	889
`objection`	69	70	67.5	2.1	845.5
`prisma`	93	95	92	3.2	630
`mikroorm`	120	123.5	119	3.8	478.5

Table S24: Multi-worker deep-fetch throughput (req/s, median of three runs) as each layer is scaled to 1, 2, and 4 Express workers (node cluster, shared port). Every layer is single-core-bound per process, so throughput scales roughly linearly with workers until the host saturates: Prisma, low at one worker (PostgreSQL 1,117), rises to 2,275 at two and 4,449 at four workers, recovering native-like aggregate throughput at four times the processes. A layer's low per-process throughput is therefore recoverable by horizontal scaling, at a proportional cost in cores.

Layer	1 worker	2 workers	4 workers
PostgreSQL			
`pg`	3299	4055	3961
`prisma`	1117	2275	4449
`mikroorm`	435	898	1773
MySQL			
`mysql2`	2327	4698	8883
`prisma`	591	1167	2427
`mikroorm`	418	851	1624

Table S25: Deep-fetch throughput (req/s, MySQL, 50 connections, median of 3) as the connection pool varies from 1 to 50, per access layer. The primary campaign fixes the pool at 10; each layer saturates at a pool well below 50 and the ordering is preserved at every pool size, so the fixed-pool choice does not drive the ranking.

Layer	1	2	5	10	20	50
`mysql2`	1545	2156	2352	2446	2506	2494
`knex`	1255	1926	1984	2081	2131	2120
`prisma`	464	572	631	639	626	648
`mikroorm`	416	469	486	418	492	429

ordering matches the read ordering and is insensitive to the mix, because the deep-fetch read dominates the cost; exploratory.

Table S26: Mixed read/write workload: throughput (req/s) and p99 (ms) under an autocannon request mix interleaving the deep-fetch read with single-row inserts, at 90:10 and 70:30 read:write, both engines, median of five runs with a physical write reset per run. The layer ordering matches the read ordering and is insensitive to the mix (the deep-fetch read dominates); exploratory.

Engine	Layer	read:write	req/s	p99
PG	`pg`	90:10	4236	19
PG	`pg`	70:30	4347	18
PG	`knex`	90:10	3024	26
PG	`knex`	70:30	2993	26
PG	`prisma`	90:10	1504	55
PG	`prisma`	70:30	1409	58
PG	`mikroorm`	90:10	676	103
PG	`mikroorm`	70:30	705	99
MySQL	`mysql2`	90:10	2283	44
MySQL	`mysql2`	70:30	2505	47
MySQL	`knex`	90:10	2275	37
MySQL	`knex`	70:30	2232	40
MySQL	`prisma`	90:10	718	114
MySQL	`prisma`	70:30	729	103
MySQL	`mikroorm`	90:10	626	109
MySQL	`mikroorm`	70:30	639	105

25 Concurrency sweep

Figure S2 sweeps the deep/nested fetch from 1 to 256 client connections for a representative layer of each tier, both engines. The 50-connection operating point used throughout sits above the knee for every layer; the ordering is load-robust above about 32 connections. This is the evidence behind the choice of operating point.

26 Per-layer cross-engine ranks

Table S27 gives each portable layer's throughput rank on PostgreSQL versus MySQL for the deep fetch and the insert, backing the rank correlations in the main text: the deep-fetch orderings largely agree ($\rho = 0.86$) while the

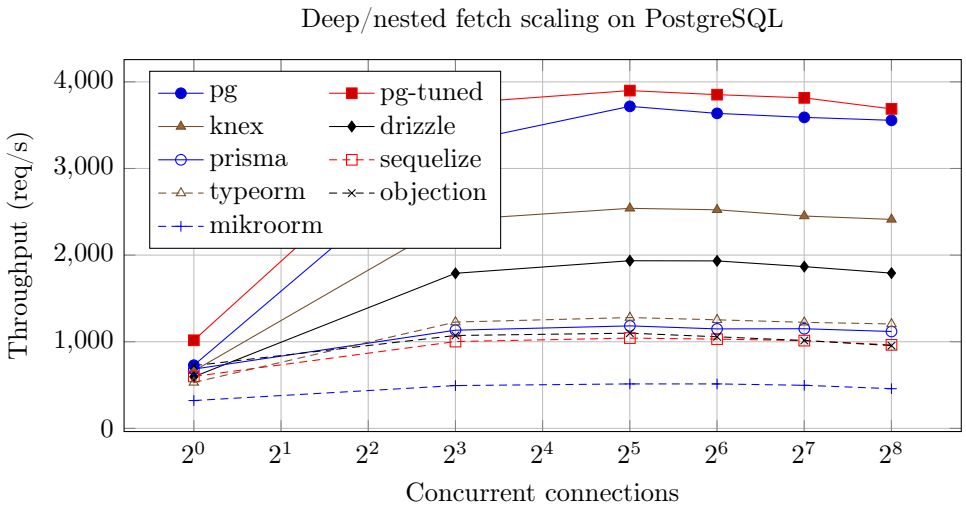

Figure S2: Throughput of each access layer on the deep/nested fetch as concurrency increases (postgres). The native driver leads at every concurrency level, with the query builder and lightweight ORM Drizzle next and the data-mapper ORMs (Prisma among them) below; the ordering is stable above about 32 connections and compresses only at very low concurrency.

insert orderings scramble more ($\rho = 0.68$), most visibly Prisma (rank 4 on PostgreSQL, 7 on MySQL).

Because MikroORM is the largest single outlier in several spreads, we recompute the cross-engine agreement without it: over the remaining six portable layers the read ordering stays similar (deep-fetch $\rho = 0.77$, $\tau = 0.60$; range $\rho = 0.83$, $\tau = 0.73$) and the deep-fetch native-relative spread stays substantial at 3.72× (PostgreSQL) and 3.76× (MySQL), so the RQ2 and RQ3 findings do not depend on that one layer.

27 Layer-by-engine interaction magnitude

The blocked permutation test reports *that* the layer ordering interacts with the engine; Table S28 reports *how large* the interaction is, as the PostgreSQL÷MySQL throughput ratio per layer and pattern with a 95% paired-bootstrap interval. Every layer is faster on PostgreSQL, so the interaction is in the spread of the advantage: the reads hold a narrow band across layers (≈1.0–1.6) while the insert scatters from 1.6 to 3.2, the quantitative counterpart of the low insert rank correlation.

Table S27: Per-layer throughput rank (1 = fastest) on PostgreSQL versus MySQL for the seven portable layers, backing the rank correlations. On the deep/nested fetch the two engines' orderings largely agree (Spearman $\rho = 0.86$, Kendall $\tau = 0.71$); on the insert they scramble more ($\rho = 0.68$, $\tau = 0.52$), most visibly Prisma, which is mid-pack on PostgreSQL (rank 4) but last on MySQL (rank 7). This is why aggregation and writes are reported as transferring only moderately across engines. Excluding MikroORM (the largest single outlier) leaves the read ordering similar (deep-fetch $\rho = 0.77$, $\tau = 0.60$) and the deep-fetch native-relative spread at 3.72× (PostgreSQL) and 3.76× (MySQL), so no single layer drives these findings.

	Deep fetch		Insert	
Layer	PG	MySQL	PG	MySQL
`knex`	1	1	1	3
`drizzle`	2	2	2	1
`prisma`	4	6	4	7
`sequelize`	6	4	5	4
`typeorm`	3	3	6	5
`objection`	5	5	3	2
`mikroorm`	7	7	7	6

Table S28: Layer × engine interaction magnitude: the PostgreSQL ÷ MySQL throughput ratio (geometric mean of the paired per-replicate ratios, with a 95% paired-bootstrap interval) for each portable layer and pattern. A ratio > 1 means the layer runs faster on PostgreSQL, < 1 faster on MySQL. Every layer is faster on PostgreSQL here, so the interaction is in the *spread* of the advantage, not its sign: the three read patterns hold a narrow band across layers (≈1.0–1.6, a near-parallel profile, small interaction), whereas aggregation and especially the insert scatter widely (insert 1.6–3.2), reordering the layers across engines — the interaction the blocked permutation test detects and the rank correlations of Table S27 summarize.

Layer	Point read	Range scan	Deep fetch	Aggregation	Insert
`knex`	1.31 [1.28,1.34]	1.07 [1.06,1.09]	1.23 [1.21,1.25]	1.46 [1.44,1.48]	3.05 [2.89,3.23]
`drizzle`	1.44 [1.42,1.46]	1.28 [1.27,1.30]	1.41 [1.40,1.42]	1.80 [1.77,1.83]	2.62 [2.46,2.79]
`prisma`	1.57 [1.55,1.59]	1.59 [1.57,1.62]	1.89 [1.85,1.92]	1.90 [1.87,1.92]	3.23 [3.14,3.32]
`sequelize`	1.28 [1.26,1.29]	1.17 [1.15,1.19]	1.12 [1.10,1.13]	1.49 [1.48,1.51]	1.82 [1.73,1.91]
`typeorm`	1.38 [1.36,1.41]	1.15 [1.12,1.17]	1.19 [1.17,1.21]	1.60 [1.57,1.63]	2.05 [1.97,2.13]
`objection`	1.22 [1.20,1.24]	1.04 [1.03,1.05]	1.22 [1.21,1.24]	1.37 [1.35,1.40]	2.38 [2.28,2.50]
`mikroorm`	1.19 [1.17,1.21]	1.06 [1.04,1.07]	1.06 [1.04,1.08]	1.51 [1.48,1.54]	1.60 [1.56,1.64]

28 A reporting checklist for access-layer benchmarks

Distilled from the design decisions and the pitfalls this study confronted, the following items are the ones an access-layer benchmark should report so that its claims can be scoped and reproduced. It is a reporting aid, not a standard.

1. **Correctness before timing.** Assert byte-identical responses across layers (final serialized bodies), or state what differs; several defects in this study were caught only this way.
2. **Experimental unit.** Treat the freshly booted run, not the individual request, as the unit; requests within a run are not independent.
3. **Same-host repetition is not environmental replication.** Report how many runs, on how many hosts, over what period; do not call same-host repeats "independent."
4. **Idiomatic treatment.** State the exact API, its documentation page and version, the declined alternative, and whether logging, hooks, validation, and identity maps are on.
5. **Latency construct.** Distinguish saturated closed-loop response time (which tracks throughput and queueing) from utilization-controlled or open-loop latency; report the operating point.
6. **Same-SQL / shared-plan control.** Report it as a compound intervention (it changes query formulation, protocol, preparation, round-trips, and mapping together), not a mechanism decomposition.
7. **Resource accounting.** Report application-tier and database CPU separately; throughput per CPU-second, not throughput alone.
8. **Write-state control.** Reset the mutable state between runs and state the durability regime (default versus relaxed).
9. **Provenance.** Pin engine images by digest and dependencies by lockfile; archive raw per-cell records, seeds, and a table-to-record manifest.
10. **Scope.** State the host topology, workload, versions, and dates, and confine claims to them.

29 Study factors

Table S29 records which factors the benchmark holds constant, which it varies, and which are uncontrolled on the single virtualized host; the uncontrolled factors are why results are reported for this configuration rather than as a general library ranking.

Table S29: Factors in the benchmark: what is held constant, what varies (the treatments), and what is uncontrolled. The same-SQL control additionally holds the generated SQL and row mapping constant; because it changes query formulation, API, protocol, statement preparation, round-trip structure, and mapping *together*, it bounds the residual difference but does not isolate any single factor. The uncontrolled factors are inherent to a single virtualized host and are the reason results are reported for this configuration, not as general library performance.

Status	Factor	Detail
Held constant	schema, indexes, seed data	identical DDL + deterministic seed, both engines
	per-replicate request stream	paired PRNG stream, seeded on (endpoint, replicate)
	connection pool	fixed at ten for every layer
	response bytes	byte-identical JSON across layers (verified)
	host, OS, engine versions	one host; pinned July-2026 versions (Supplement Table S11)
	warm-up, run duration, cell order	15 s warm-up, 12 s runs, randomized order
Varies (treatments)	access layer	9 idiomatic + 2 tuned implementations
	database engine	PostgreSQL, MySQL
	access pattern	point read, range scan, deep fetch, aggregation, insert
Uncontrolled	hypervisor neighbours, CPU steal	single virtualized host
	thermal / frequency scaling	not pinned; bounded by the drift and post-restart checks
	database cache / kernel history	warm by design; randomized order + fresh boot mitigate